



CIS 415

Operating Systems

Introduction

Prof. Allen D. Malony

Department of Computer and Information Science
Spring 2020



UNIVERSITY OF OREGON

Logistics

- Do not post links to a PDF version of the Operating Systems Concepts book online
 - OSC is not free and it is copyrighted material
 - It is a violation of copyright to provide a link to a PDF
 - UO can get in big trouble!
 - You can get in big trouble!
- Labs on Friday!
 - Begin working on Lab 1 assignment
- Project 1 to be posted by COB on Friday if not sooner

Objectives

- Operating systems
 - What is it?
 - What motivates it?
 - Bit of history
 - Overview
- Review of computer systems organization

What is an Operating System?

- A program that acts as an intermediary between a user of a computer and the computer hardware



What is an Operating System?

- A program that acts as an intermediary between a user of a computer and the computer hardware
- Operating system goals:
 - Execute user programs
 - Make solving user problems easier
 - Make the computer system convenient to use
 - Use the computer hardware in an efficient manner
- Ok, do I need an OS just to print something?

```
main()
{
    printf("This is not a printout ...\\n");
}
```

Without an OS

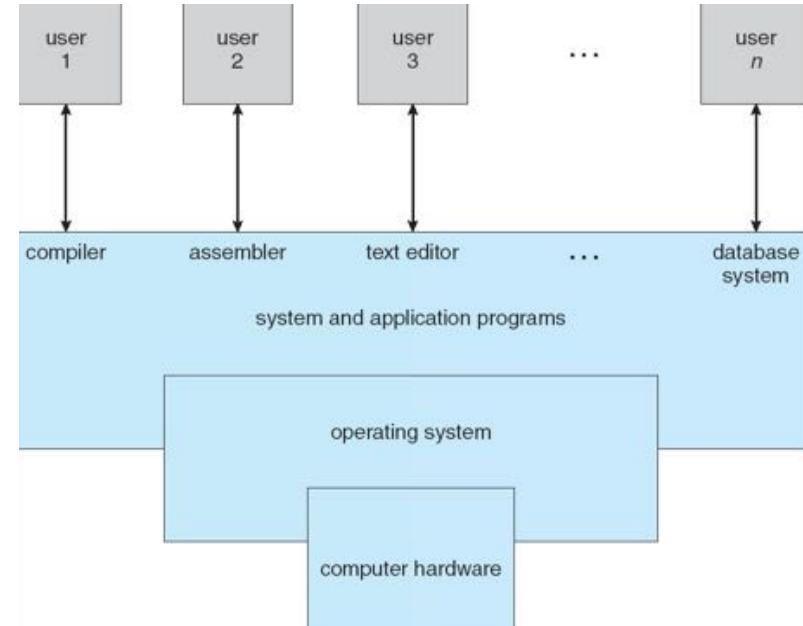
- Get printer manual
 - Figure out how to send messages to it
 - Write the program
 - Put the character string “This is not a printout ...” in a memory buffer
 - Do the stuff printer requires to send buffer to it
 - Go into endless loop
 - wait for someone to turn off computer eventually
 - Get hold of a computer
 - It has to be all yours!
 - Translate your program into machine code
 - Has to be specific to the computer you are using
 - Figure out how to get program into memory
 - Font panel switches
 - Burn a ROM
 - Punch cards
 - Somehow start program
 - Turn off computer when done
- YUCK!!!

With an OS

- Put program in file
 - Put character string “This is not a printout ...” in a memory buffer
 - Issue system call to send buffer to printer
- Compile the program
- Tell OS to run the program file
 - Turn on the printer first
- That’s everything!
- Uh, ok, but what’s a system call?

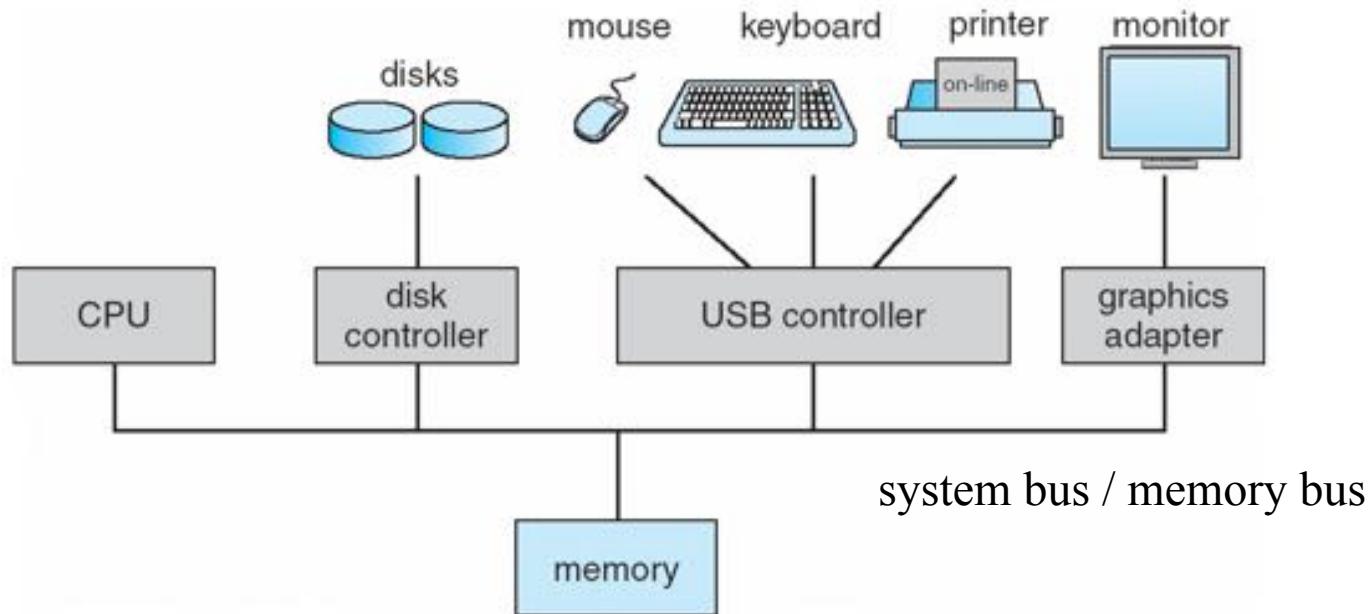
Computer System Structure

- Computer system divided into four components:
 - Hardware – provides basic computing resources
 - ◆ CPU, memory, I/O devices
 - Operating system
 - ◆ controls and coordinates use of hardware among various applications and users
 - Application programs
 - ◆ use system resources to solve computing problems
 - ◆ word processors, compilers, web browsers, database systems, video games, ...
 - Users
 - ◆ People, machines, other computers



Basic Computer System Organization

- Computer-system operation
 - One or more CPUs, device controllers connect through common bus providing access to shared memory
 - Concurrent execution of CPUs and devices competing for memory cycles



Computer System Architecture

- Most systems use a single general-purpose processor
 - Most systems have special-purpose processors as well
- Multiprocessors systems growing in importance
 - Advantages include:
 - ◆ increased throughput
 - ◆ economy of scale
 - ◆ increased reliability – graceful degradation or fault tolerance
 - Two types:
 - ◆ asymmetric multiprocessing – each processor has a specific task
 - ◆ symmetric multiprocessing – each processor performs all tasks
 - Can be used to run parallel applications

Computer System Operation

- I/O devices and the CPU can execute concurrently
- Each device controller is in charge of a device type
- Each device controller has a local buffer
- CPU moves data from/to main memory to/from local buffers associated with I/O devices
- I/O is from the device to local buffer of controller
- Device controller informs CPU that it has finished its operation by causing an interrupt

What Operating Systems Do?

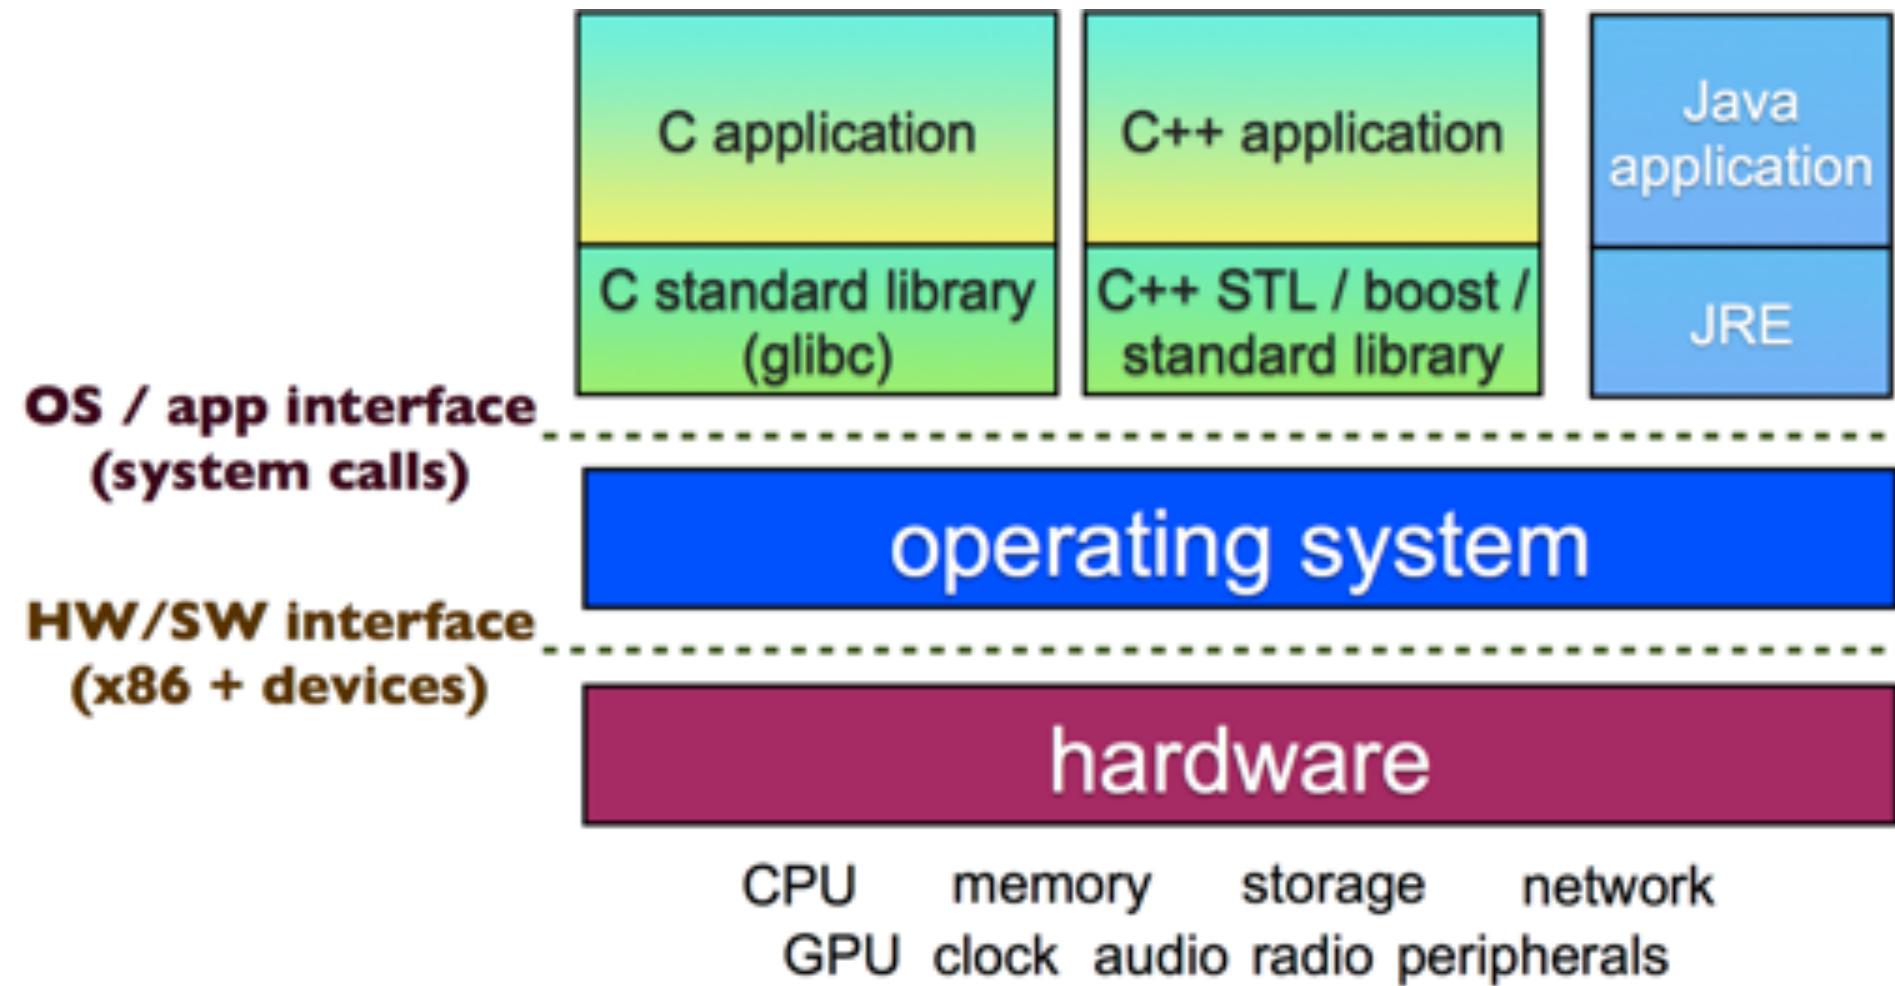
- Depends on the point of view
- Users perspective
 - Convenience, ease of use, good performance
 - Dedicated resources
- Server perspective
 - Keep all users happy
 - Shared resources
- Both perspectives are valid
- Different OS requirements for different environments
- Common OS fundamentals at core

Operating System Definition

- OS is a *resource allocator*
 - Manages all resources
 - Decides between conflicting requests
 - ◆ efficient and fair resource use
- OS is a *control program*
 - Controls execution of programs to prevent errors and improper use of the computer
- OS provides both resource allocation and control
- Still can have different perspective on scope
 - An OS is everything a vendor ships you when you order an operating system
 - An OS is the one program running at all times (kernel), and everything else is a system program or application

Where does the OS fit in?

Users



What does an OS do again?

- Acts as a resource manager
- Processes
 - Hides programs from one another
- Traffic cop
 - Resource management
 - Who gets to run, when?
- Memory management
 - Protection from other programs' mistakes
- Security
 - Protection from other programs' malice
- System call interface
 - Abstract, simplified interface to services
 - Like a function library, but communicates with the OS
- Portability
 - Programs don't have to take into account details of their environment
- Device management
- Communication
 - Between processes
 - To devices & networks

Operating System History

- 1950s: Simplify operators' job
- 1960s: Structure, concepts, concurrency, theory, ...
- 1970s: Small and flexible (UNIX)
- 1980s: Individual user systems (PCs)
- 1990s: Internet, distributed systems
- 2000s: Security, Linux/Windows/Mac OS
- 2010s: Embedded, distributed, mobile, parallel,
virtual, ...
- Timeline of operating systems (Wikipedia)
https://en.wikipedia.org/wiki/Timeline_of_operating_systems

Operating Systems (1950s)

- Primitive systems
 - Little memory
 - Programs stored on tape
- Single user
 - Batch processing
 - Computer execute one function at a time
- No overlap of I/O and computation
- OS for IBM 704
 - Atlas Supervisor (University of Manchester)
 - University of Michigan Executive System (UMES)
- You had to wear a suit or dress to use the OS!

IBM 704



Operating Systems (1960s)

- Multiprogramming
 - Timesharing
 - Multiple programs run concurrently
- Many operating systems concepts invented
 - Virtual memory, hierarchical file systems, synchronization, security and many more
- End up with slow, complex systems on limited hardware (Multics)
- Multics was an MIT project which partnered with Bell Labs and GE to try to support multiple users
- You had to have glasses to design the OS!



Edsger W. Dijkstra



Peter J. Denning

Operating Systems (1970s)

- Becoming more available
 - UNIX (Bell Labs)
 - ◆ first OS written in a high-level language
- Becoming more flexible
 - Extensible system
 - Community forms beyond developers (precursor to open source movement)
- Performance focus
 - Optimization of algorithms from 1960s
- You had to have beards to program the OS!



Brian W. Kernighan (L)
Dennis M. Ritchie (R)



Operating Systems (1980s)

□ Critical mass reached

- A variety of well-known systems, concepts
- UNIX fragments (BSD, SysV, ...)
- Open source begins to emerge

□ PC Emerges

- Simple, single user
- No network
- Simple OSes: DOS

□ Graphical User Interfaces

- X Windows
- Apple Macintosh



Bill Gates



Steve Jobs



Operating Systems (1990s)

- Connect to Internet
 - “Real OSes” for PCs
 - ◆ NT/2000+, Linux, Mac OS X
 - Web emergence
- Server systems galore
 - Mainframes even re-emerge
- Complex systems and requirements
 - Multiprocessors, memory hierarchy, interconnects, ...
 - Parallel
 - Real-time
 - Distributed



Linux B. Torvalds and Penguin

More Operating System Humor



Operating Systems (2000s)

- More dimensions (problems) and range of scales
 - Distributed systems
 - Multicore
 - Security
 - Ubiquitous
 - Virtual Machines
 - Embedded
 - Mobile
- More dominance of OSes (Linux, MacOS, Windows)



The Holy Trinity

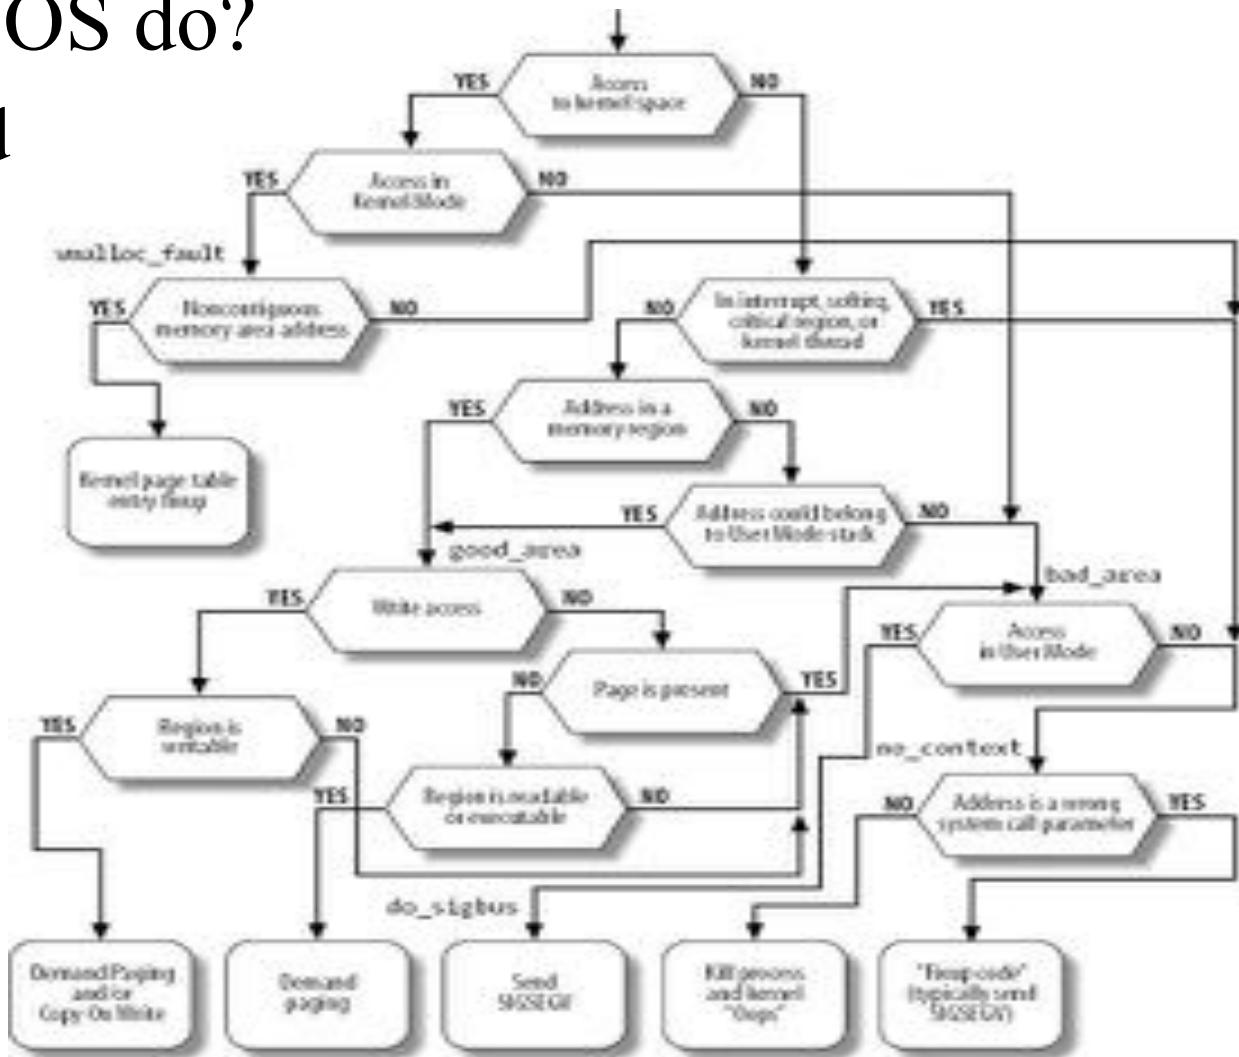
Operating Systems (2010s)

- New, interesting challenges ahead
- Cloud computing is real
- Mobile devices are proliferating
- Internet of things
- Virtual machines
- Parallel computing
- Extreme-scale supercomputing
- OSes of various sorts for various purposes
 - iOS, Android, ...
- Good time to learn OS principles and techniques



Under the OS Covers – OS Functions

- What does the OS do?
 - Mostly behind the scenes...
- Consider how page faults are handled
- What are page faults?



Page Fault Handling

- A page fault is caused when access to a virtual memory location is not found in physical memory
- What happens?
 - Trap is generated by the hardware
 - Handler in OS determines how to obtain memory
 - If page is still on disk, then handler
 - ◆ allocates physical page
 - ◆ makes I/O request to disk via file system and driver
 - Driver copies page from disk into new physical page
 - OS restarts the process at the trapped instruction

Page Fault Handling – Considerations

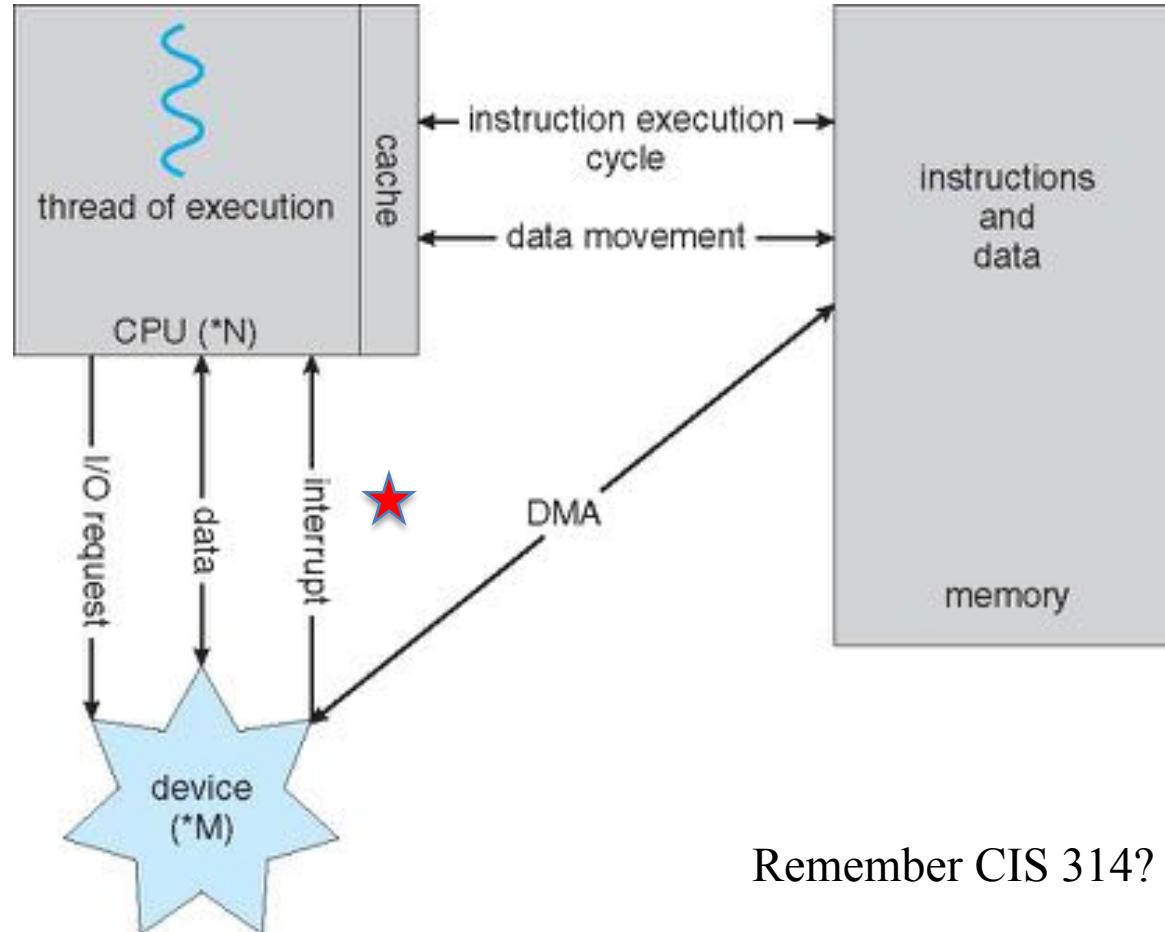
- Suppose there are multiple processes
 - More complex of a situation than with just 1 process
- OS has to make trade-offs
 - What if there are no physical pages available?
 - There may be multiple outstanding disk requests, so what order should they be processed?
 - How does the OS interact with hardware effectively?
 - What process should be run while waiting for the page fault to be handled?
 - Many others...

Learning about Operating Systems

- OS has many protocols like page fault handling
 - You will need to know them
- OS designers add layers of abstraction to simplify programming (e.g., virtual memory)
 - You will need to understand these concepts
- The design of protocols using these concepts involves trade-offs (e.g., disk performance)
 - You will need to understand why OS protocols are written the way that they are

Device I/O – How a modern computer works?

□ A von Neumann architecture



Remember CIS 314?


Common Functions of Interrupts

- Interrupt transfers control to the interrupt service routine, generally through an interrupt vector
- The operating system preserves the state of the CPU by storing registers and the program counter
- Determines which type of interrupt has occurred:
 - Polling
 - Vectored interrupt system
- Interrupt architecture must save the address of the interrupted instruction
- A trap or exception is a software-generated interrupt caused either by an error or a user request
- An operating system is interrupt driven

I/O Structure

- After I/O starts, control returns to user program ...
 - Only upon I/O completion
 - ◆ wait instruction idles the CPU until the next interrupt
 - ◆ wait loop (contention for memory access)
 - ◆ at most one I/O request is outstanding at a time
 - Without waiting for I/O completion
 - ◆ system call – request to the OS to allow user to wait for I/O completion
 - ◆ device-status table contains entry for each I/O device indicating its type, address, and state
 - ◆ OS indexes into I/O device table to determine device status and to modify table entry to include interrupt

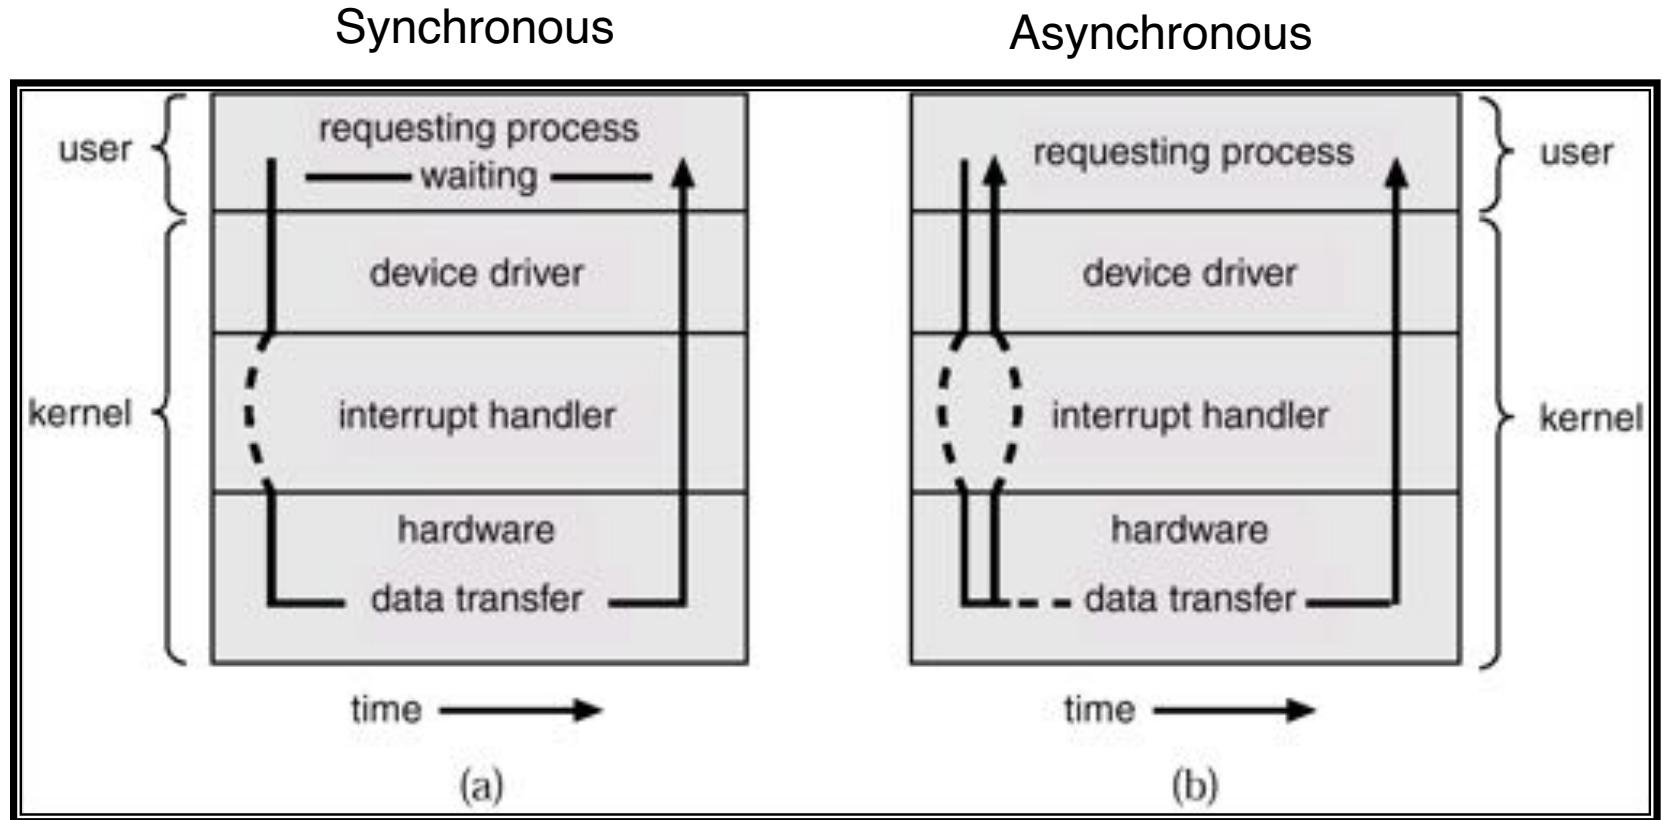
Interrupt Timeline



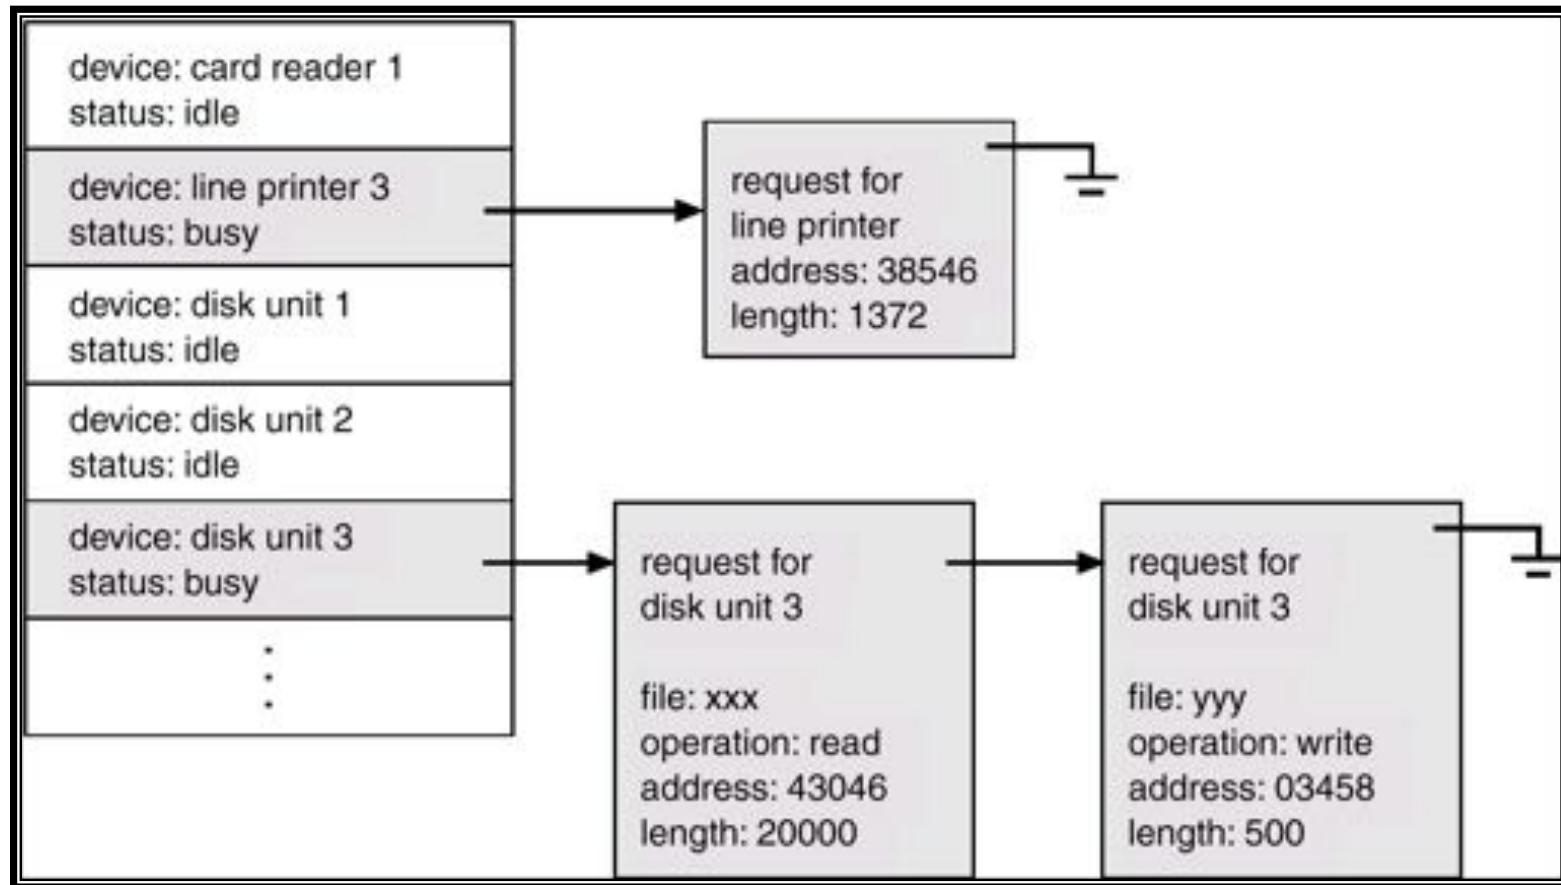
I/O Structure – Two Methods (1)

- Synchronous
 - After I/O starts, control returns to user program only upon I/O completion
 - A wait instruction idles the CPU until next interrupt
 - Or the CPU runs a wait loop (problems?)
 - No simultaneous I/O processing (only 1 request)
- Asynchronous
 - After I/O starts, control returns to user program without waiting for I/O completion
 - User program then waits for I/O complete
 - I/O interrupts upon completion

I/O Structure – Two Methods (2)




Device Status Table



Timer Interrupts for Kernel Control

- Timer to prevent infinite loop / process hogging resources
 - Timer is set to interrupt the computer after some time period
 - Keep a counter that is decremented by the physical clock
 - Operating system set the counter (privileged instruction)
 - When counter zero generate an interrupt
 - Set up before scheduling process to regain control or terminate program that exceeds allotted time



Process Management

- A process is a program in execution
 - A program is a passive entity
 - A process is an active entity (unit of work in the system)
- Process needs resources to accomplish its task
 - CPU, memory, I/O, files
 - Initialization data
- Process termination requires reclaim of any reusable resources
- Single-threaded process has one program counter specifying location of next instruction to execute
 - Process executes instructions sequentially
 - One at a time, until completion
- Multi-threaded process has one program counter per thread
- Typically system has many processes, some user, some operating system running concurrently on one or more CPUs
 - Concurrency by multiplexing the CPUs among the processes / threads

Process Management Activities

- OS is responsible for process management
- Activities include:
 - Creating and deleting both user and system processes
 - Suspending and resuming processes
 - Providing mechanisms for process synchronization
 - Providing mechanisms for process communication
 - Providing mechanisms for deadlock handling

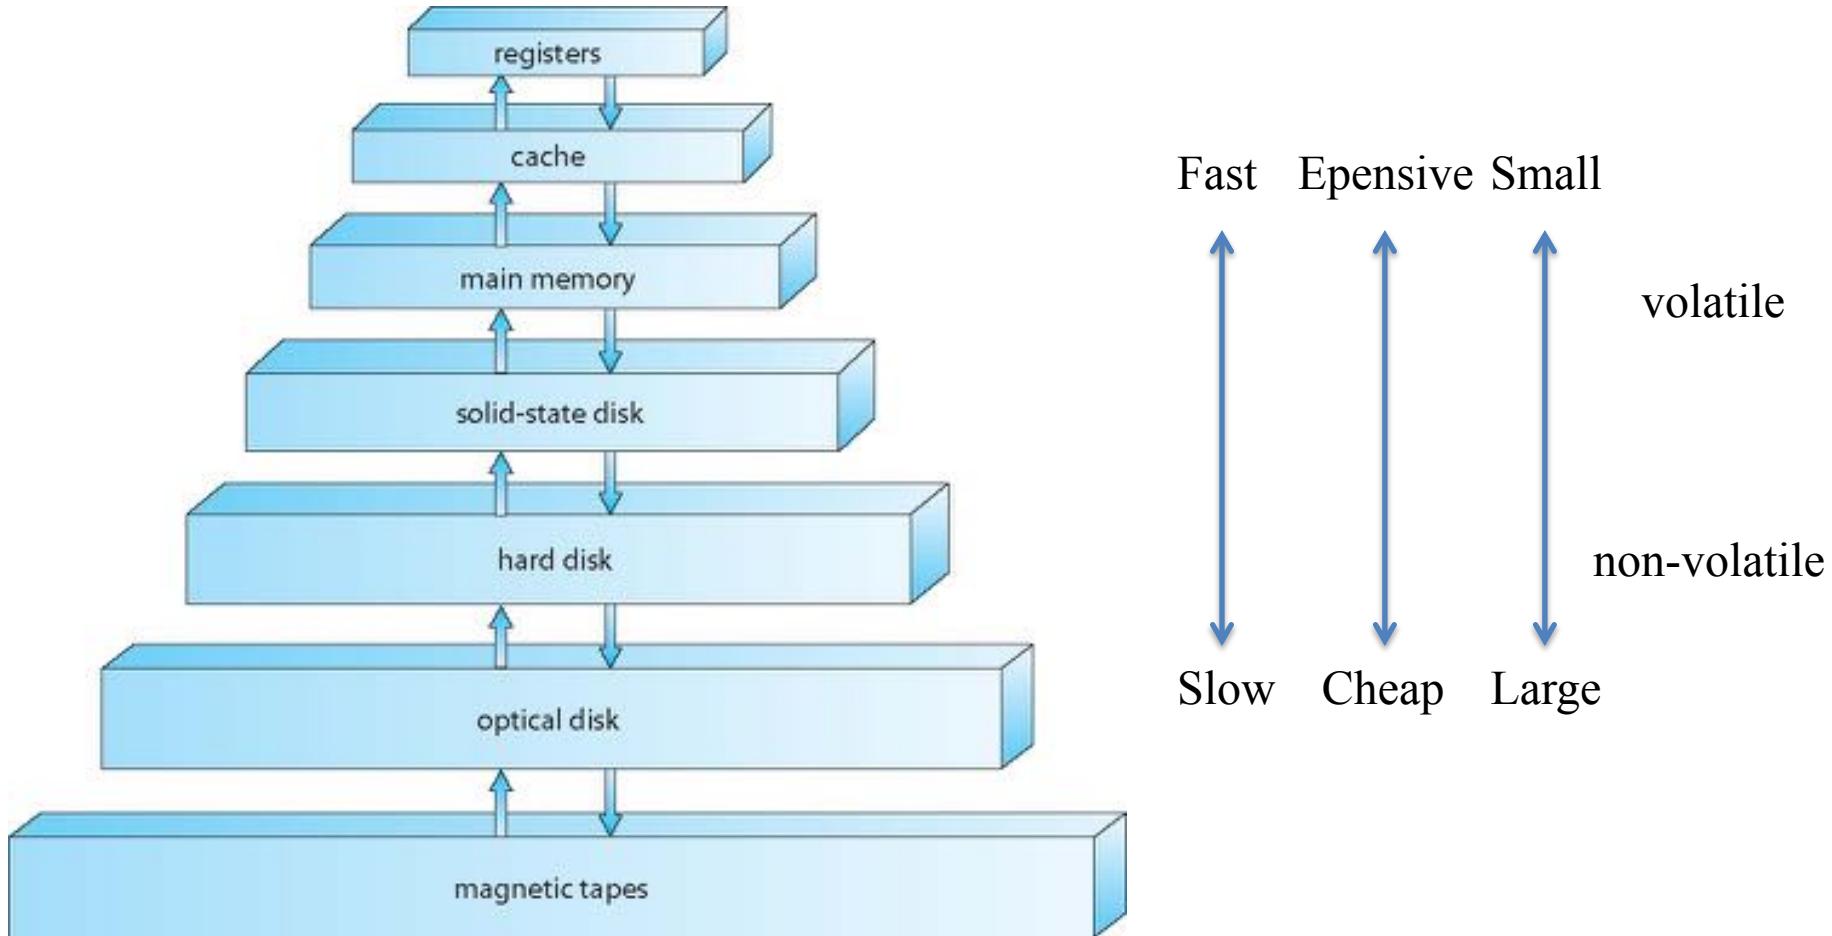
Scheduling

- Determine which task to perform given:
 - Multiple user processes
 - Multiple hardware components
- Provide effective performance
 - Responsive to users, CPU utilization
- Provide fairness
 - Do not starve low priority processes

Memory Management

- To execute a program all (or part) of the instructions must be in memory
- All (or part) of the data that is needed by the program must be in memory.
- Memory management determines what is in memory and when
 - Optimizing CPU utilization and computer response to users
- Memory management activities
 - Keeping track of which parts of memory are currently being used and by whom
 - Deciding which processes (or parts thereof) and data to move into and out of memory
 - Allocating and deallocating memory space as needed

Storage Hierarchy



Storage Management

- OS provides uniform, logical view of information storage
 - Abstracts physical properties to logical storage unit - file
 - Each medium is controlled by device
 - ◆ varying properties include access speed, capacity, data-transfer rate, access method (sequential or random)
- File-System management
 - Files usually organized into directories
 - Access control on most systems determines who can access what
 - OS activities include
 - ◆ Creating and deleting files and directories
 - ◆ Primitives to manipulate files and directories
 - ◆ Mapping files onto secondary storage
 - ◆ Backup files onto stable (non-volatile) storage media

Protection and Security

- Protection – any mechanism for controlling access of processes or users to resources defined by the OS
- Security – defense of the system against internal and external attacks
 - Huge range, including denial-of-service, worms, viruses, identity theft, theft of service
- Systems generally first distinguish among users, to determine who can do what
 - User identities (user IDs, security IDs) include name and associated number, one per user
 - User ID then associated with all files, processes of that user to determine access control
 - Group identifier (group ID) allows set of users to be defined and controls managed, then also associated with each process, file
 - Privilege escalation allows user to change to effective ID with more rights

OS Structures

- Multiprogramming (batch system)
 - Single user cannot keep CPU and I/O devices busy at all times
 - Multiprogramming organizes jobs (code and data) so CPU always has one to execute
 - A subset of total jobs in system is kept in memory
 - One job selected and run via job scheduling
 - When it has to wait (for I/O for example), OS switches to another job
- Timesharing (multitasking)
 - CPU switches jobs so frequently
 - Response time should be < 1 second
 - Each user has at least one program executing in memory
 - If several jobs ready to run at the same time
 - Virtual memory allows processes not completely in memory

Computing Environments – Traditional

- Stand-alone general purpose machines
- But blurred as most systems interconnect with others (i.e., the Internet)
- Portals provide web access to internal systems
- Network computers (thin clients) are like Web terminals
- Mobile computers interconnect via wireless networks
- Networking becoming ubiquitous – even home systems use firewalls to protect home computers from Internet attacks

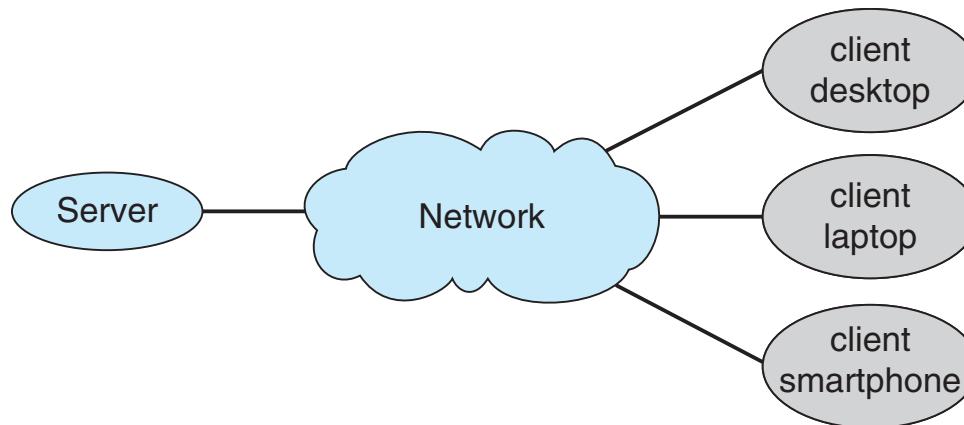
Computing Environments – Distributed

- Distributed computing
 - Collection of separate, possibly heterogeneous, systems networked together
 - ◆ network is a communications path, TCP/IP most common
 - Local Area Network (LAN)
 - Wide Area Network (WAN)
 - Metropolitan Area Network (MAN)
 - Personal Area Network (PAN)
 - Network Operating System provides features between systems across network
 - ◆ communication scheme allows systems to exchange messages
 - ◆ illusion of a single system

Computing Environments – Client-Server

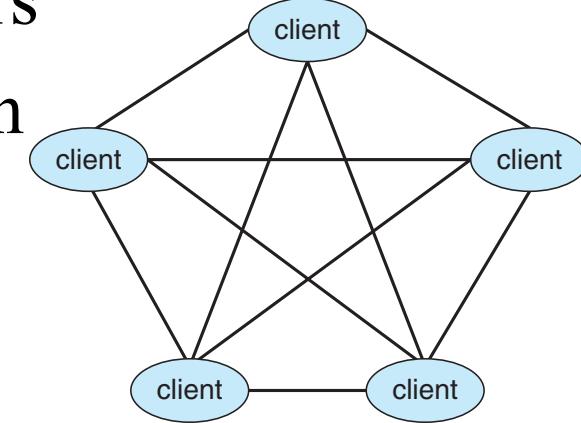
□ Client-Server Computing

- Dumb terminals supplanted by smart PCs
- Many systems now servers, responding to requests generated by clients
 - ◆ compute-server system provides an interface to client to request services (i.e., database)
 - ◆ file-server system provides interface for clients to store and retrieve files



Computing Environments – Client-Server

- Another model of distributed system
- P2P does not distinguish clients and servers
 - Instead all nodes are considered peers
 - May each act as client, server or both
 - Node must join P2P network
 - ◆ registers its service with central lookup service on network, or
 - ◆ broadcast request for service and respond to requests for service via discovery protocol
 - Examples include Napster and Gnutella, Voice over IP (VoIP) such as Skype



Computing Environments – Mobile

- Handheld smartphones, tablets,
- What is the functional difference between them and a “traditional” laptop?
- Extra feature – more OS features (GPS, gyroscope)
- Allows new types of apps like augmented reality
- Use IEEE 802.11 wireless, or cellular data networks for connectivity
- Leaders are Apple iOS and Google Android
- “Internet of things”

Computing Environments – Virtualization (1)

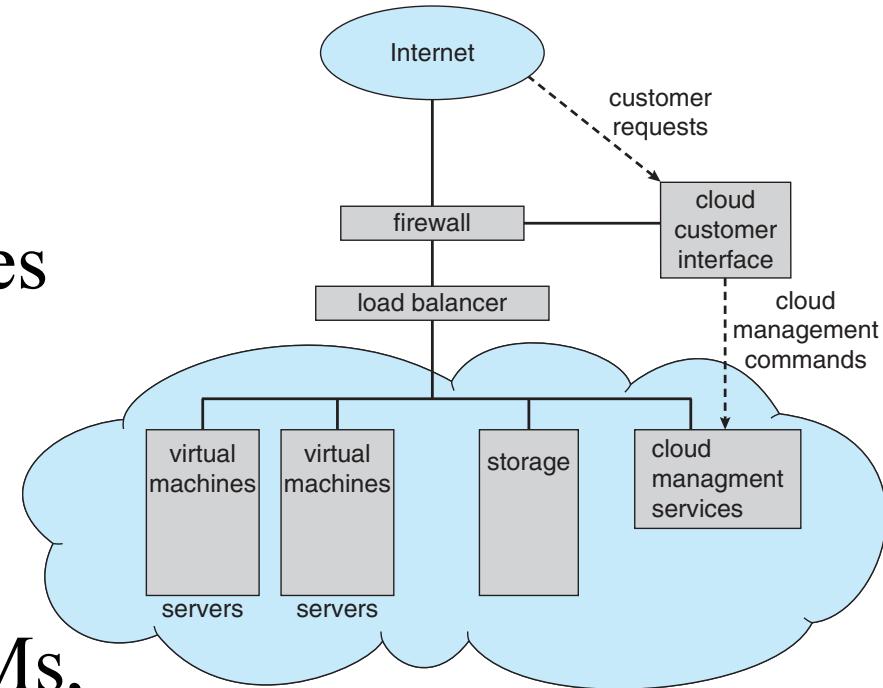
- Allows OSes to run applications within other OSes
- Emulation used when source CPU type different from target type (i.e. PowerPC to Intel x86)
 - Generally slowest method
 - When computer language not compiled to native code – Interpretation
- Virtualization – OS natively compiled for CPU, running guest OSes also natively compiled
 - Consider VMware running WinXP guests, each running applications, all on native WinXP host OS
 - VMM (virtual machine Manager) provides virtualization services

Computing Environments – Virtualization (2)

- Use cases involve laptops and desktops running multiple OSes for exploration or compatibility
 - Apple laptop running Mac OS X host, Windows as a guest
 - Developing apps for multiple OSes without having multiple systems
 - QA testing applications without having multiple systems
 - Executing and managing compute environments within data centers
- VMM can run natively, in which case they are also the host
 - There is no general purpose host then (VMware ESX and Citrix XenServer)

Computing Environments – Cloud Computing

- Delivers computing, storage, even applications as a service across a network
 - IaaS, PaaS, SaaS
- Logical extension of virtualization because it uses virtualization as the base for its functionality.
- Cloud computing environments composed of traditional OSes, plus VMs, plus cloud management tools
 - Internet connectivity requires security like firewalls
 - Load balancers spread traffic across multiple applications



Open Source Operating Systems

- Operating systems made available in source-code format rather than just binary closed-source
- Counter to the copy protection and Digital Rights Management (DRM) movement
- Started by Free Software Foundation (FSF), which has “copyleft” GNU Public License (GPL)
- Examples include GNU/Linux and BSD UNIX (including core of Mac OS X), and many more
- Can use VMM like VMware Player (Free on Windows), VirtualBox (open source and free)
 - Use to run guest operating systems for exploration



Lab 1 Assignment

- Find Lab 1 on the CIS 415 web site
- Gets you familiar with the lab environment
- Programming assignments will be using a virtual machine image that will be run inside VirtualBox
 - Practice using VirtualBox and the VM
- Canvas will be used for managing and submitting your assignments and projects
- There is a survey as part of Lab 1 assignment
- The first lab meeting will focus on making sure everyone completes the Lab 1 assignment