



CIS 415

Operating Systems

Processes

Prof. Allen D. Malony

Department of Computer and Information Science
Spring 2020



UNIVERSITY OF OREGON

Logistics

- Project 1 posted
 - Due Sunday, April 19, 11:59pm
- Labs
 - Intended to support projects
 - Lab 2 to be posted shortly
- Read OSC Chapter 3

Outline

- Process concept
- Process operation
- System calls to create processes
- Process management
- Process scheduling

Overview of Processes

- We have programs, so why do we need processes?
- Questions that we explore
 - How are processes created?
 - ◆ from binary program to executing process
 - How is a process represented and managed?
 - ◆ process creation, process control block
 - How does the OS manage multiple processes?
 - ◆ process state, ownership, scheduling
 - How can processes communicate?
 - ◆ interprocess communication, concurrency, deadlock

Superview and User Modes

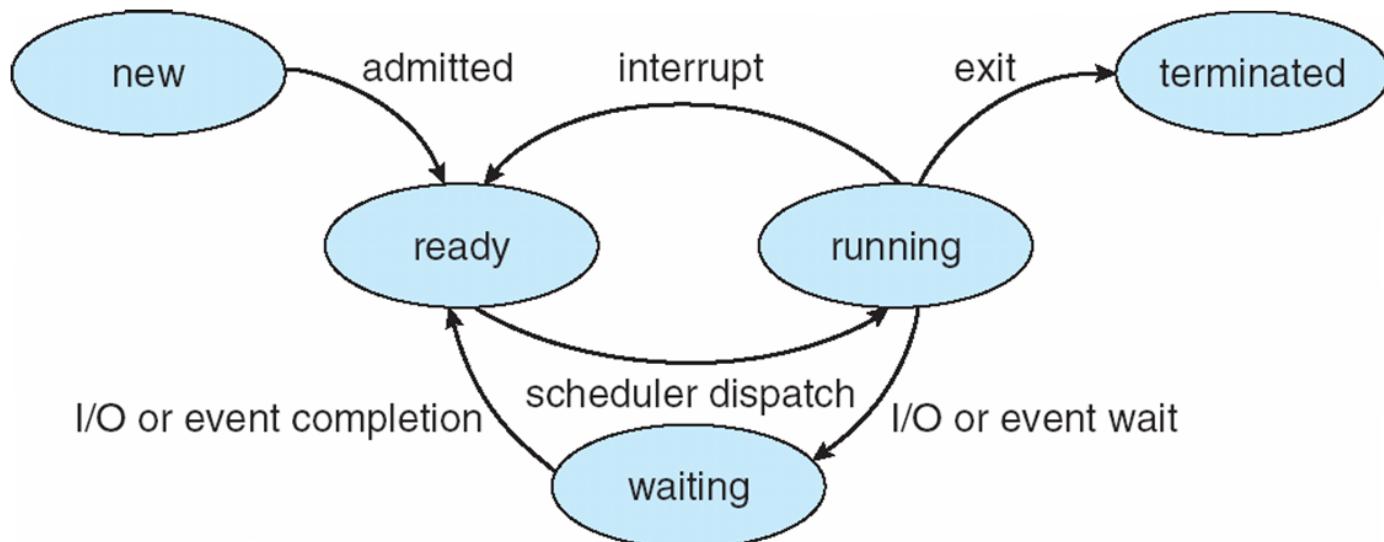
- OS runs in “supervisor” mode
 - Has access to protected (privileged) instructions only available in that mode (kernel executes in ring 0)
 - Allows it to manage the entire system
- OS “loads” programs into processes
 - Run in user mode
 - Many processes can run in user mode at the same time
- How does OS get programs loaded into processes in user mode and keep them straight?

Process Concept

- An operating system executes a variety of programs:
 - Batch system – jobs
 - Time-shared systems – user programs or tasks
- Terms *job* and *process* used almost interchangeably
- A process is a *program in execution*
 - Process execution can result in more processes being created
- Multiple parts of a process
 - Program code (image, text)
 - Execution state (program counter, processor registers, ...)
 - Stack containing temporary data (e.g., call stack frames)
 - ◆ function parameters, return addresses, local variables, ...
 - Data section containing global variables
 - Heap containing memory dynamically allocated during run time

Process Execution State

- As a process executes, it changes state
 - *New*: The process is being created (starting state)
 - *Running*: Instructions are being executed
 - *Waiting*: The process is waiting for some event to occur
 - *Ready*: The process is waiting to run
 - *Terminated*: The process has finished execution (end state)

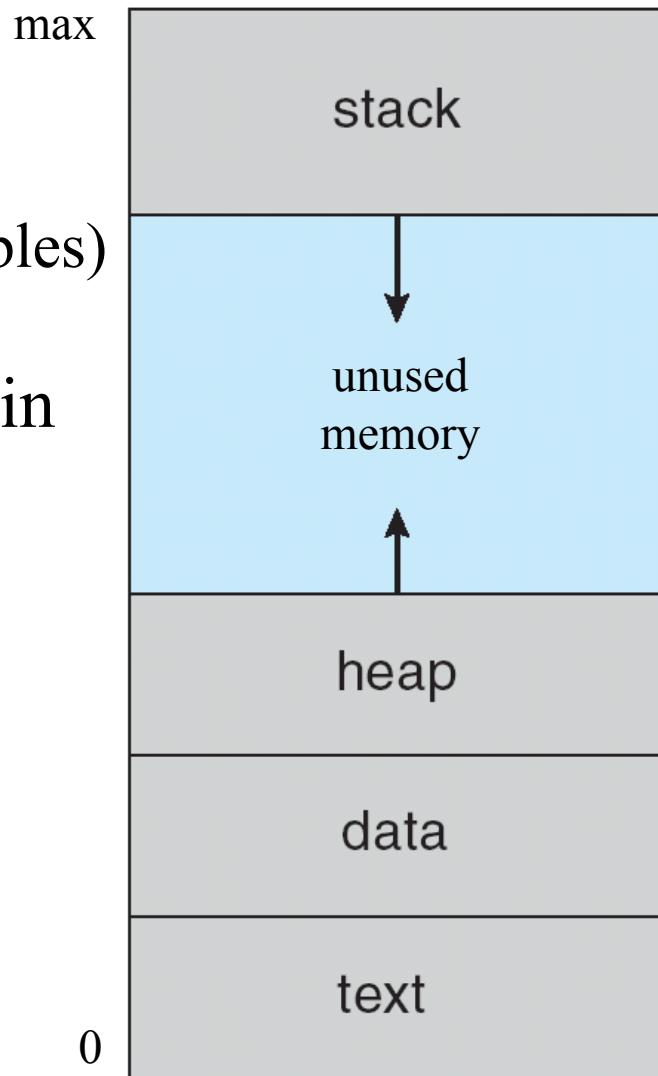


Process State

- Process state consists of:
 - Address space (what can be addressed by a process)
 - Execution state (what is need to execute on the CPU)
 - Resources being used (by the process to execute)
- Address space contains code and data of a process
- Processes are individual *execution contexts*
 - Threads are also include here (we will talk about later)
- Resources are physical support necessary to execute
 - Memory: physical memory, address translation support
 - Storage: disk, files, ...
 - Processor: CPU (at least 1)

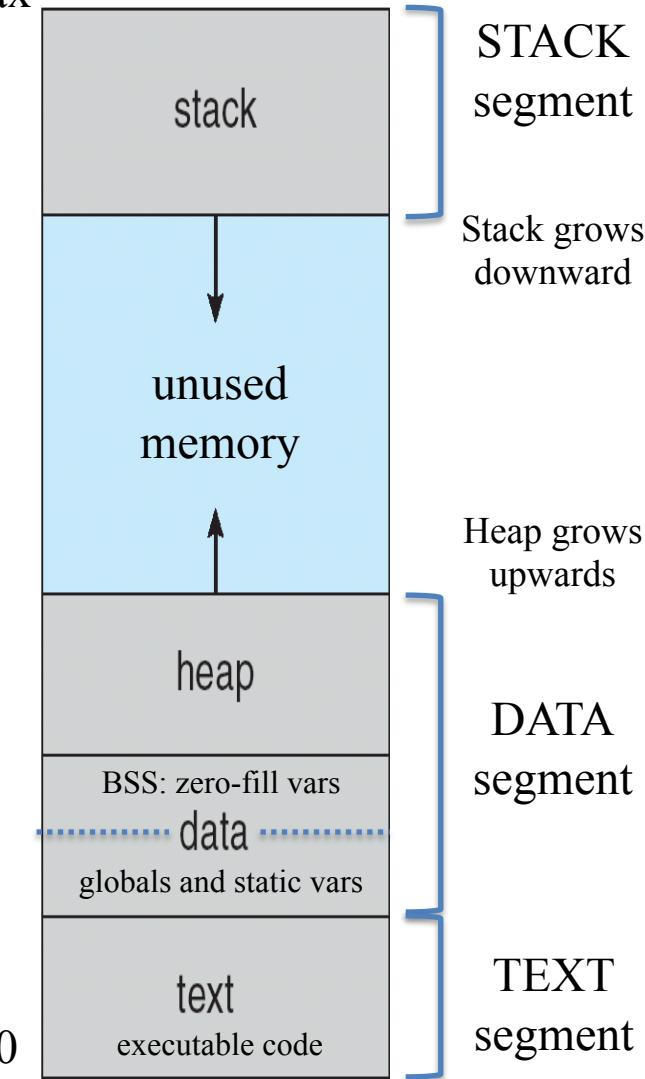
A Process in Memory

- A process has to reference memory for different purposes
 - Instructions
 - Stack (subroutine “frames”, local variables)
 - Data: static and dynamic (heap)
- Shows where are these things located in “logical” memory for a process?
- Logical memory
 - What can be referenced by an address
 - # address bits in instruction addresses determine logical memory size
 - A “logical address” is from 0 to the size of logical memory (max)
- Compiler and OS determine where things get placed in logical memory



Process Address Space

- All locations addressable by process max
 - Also called *logical address space*
 - Every process has one
- Restrict addresses to different areas
 - Restrictions enforced by OS
 - **Text segment** is where read only program instructions are stored
 - **Data segment** hold the data for the running process (read/write)
 - ◆ heap allows for dynamic data expansion
 - **Stack segment** is where the stack lives
- Process (logical) address space starts at 0 and runs to a high address

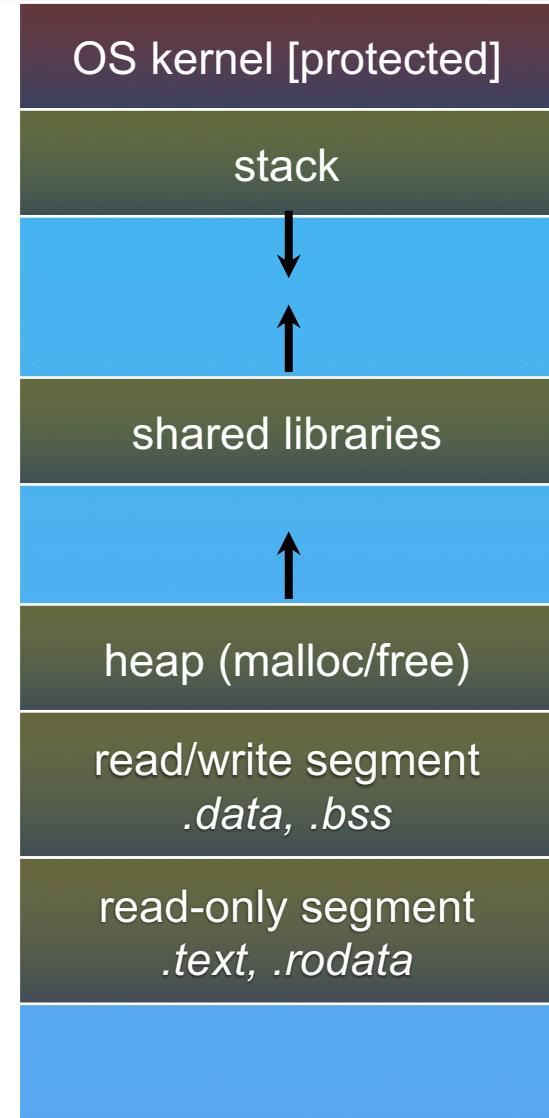


Process Address Space

- Program (Text)
- Global Data (Data)
- Dynamic Data (Heap)
 - Grows up
- Thread-local Data (Stack)
 - Grows down
- Each thread has its own stack
- # address bits determine the addressing range

0xFFFFFFFF

0x00000000



Process Address Space in Action

```
int value = 5;                                Global  
  
int main()  
{  
    int *p;                                      Stack  
  
    p = (int *)malloc(sizeof(int));                Heap  
  
    if (p == 0) {  
        printf("ERROR: Out of memory\n");  
        return 1;  
    }  
  
    *p = value;  
    printf("%d\n", *p);  
    free(p);  
    return 0;  
}
```

Process Address Space in Action

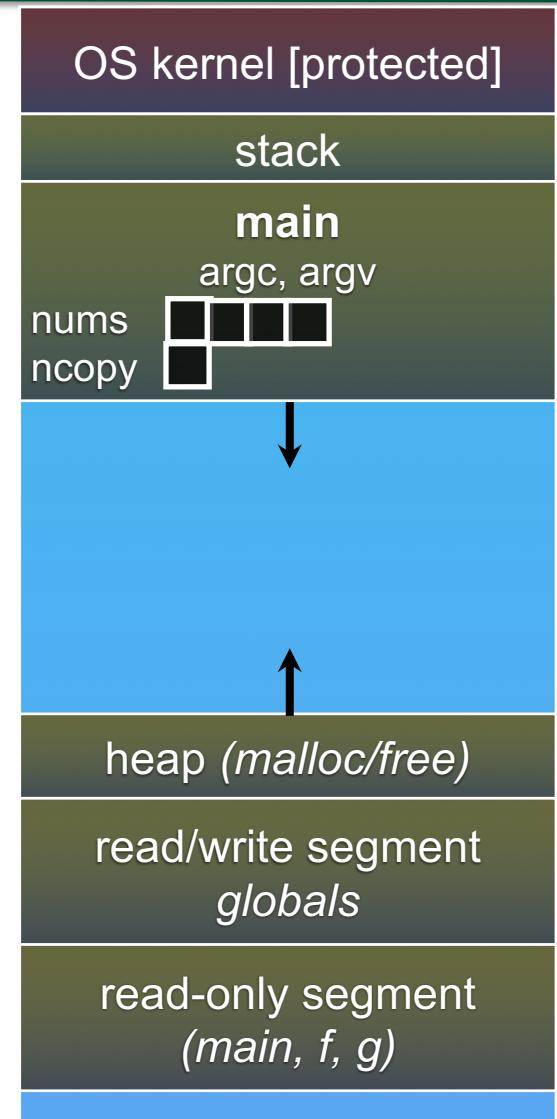
```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(nncpy);
    return 0;
}
```



Process Address Space in Action

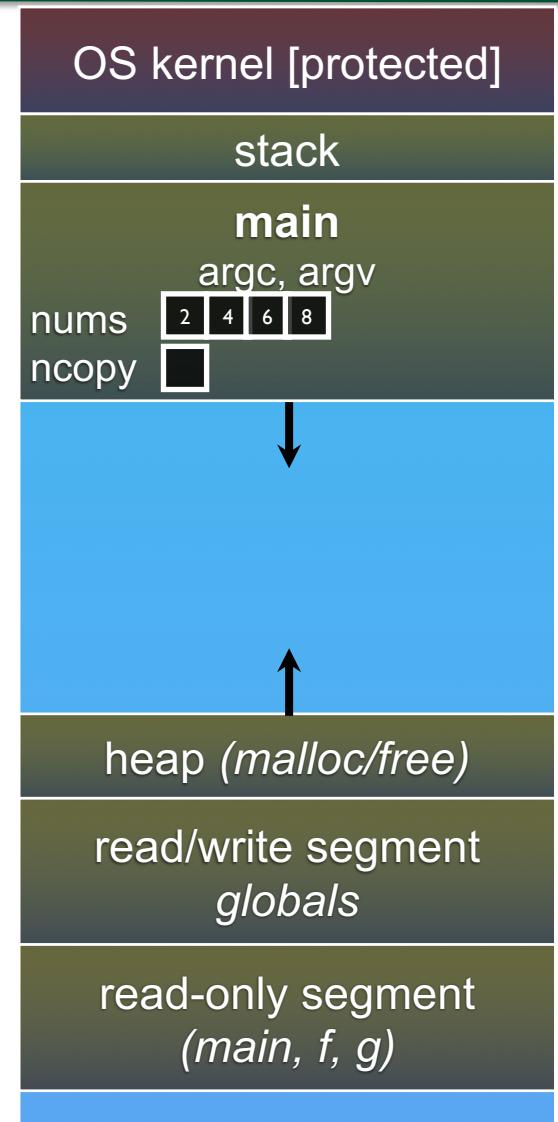
```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    freencpy;
    return 0;
}
```



Process Address Space in Action

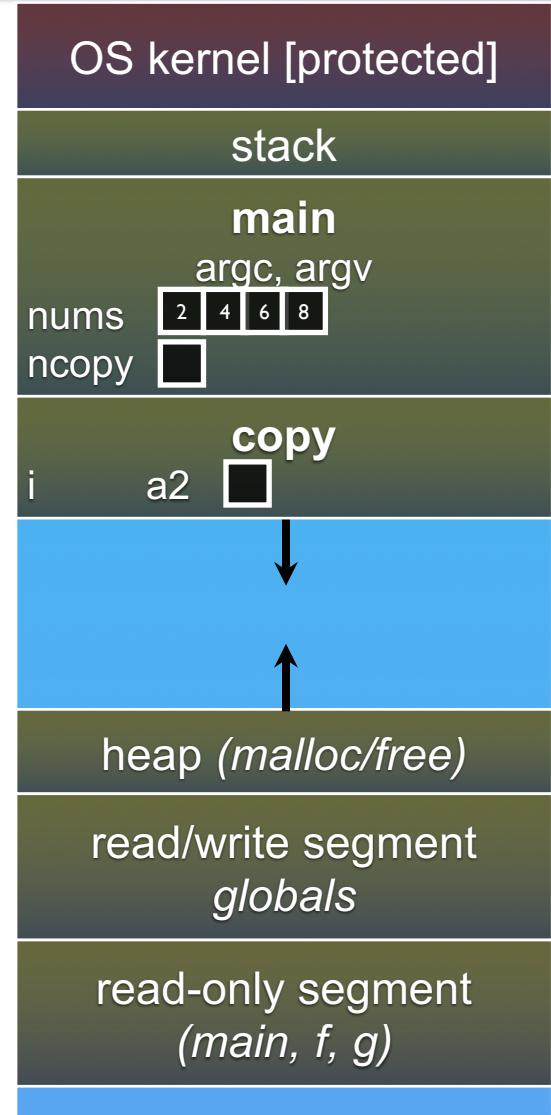
```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(nncpy);
    return 0;
}
```



Process Address Space in Action

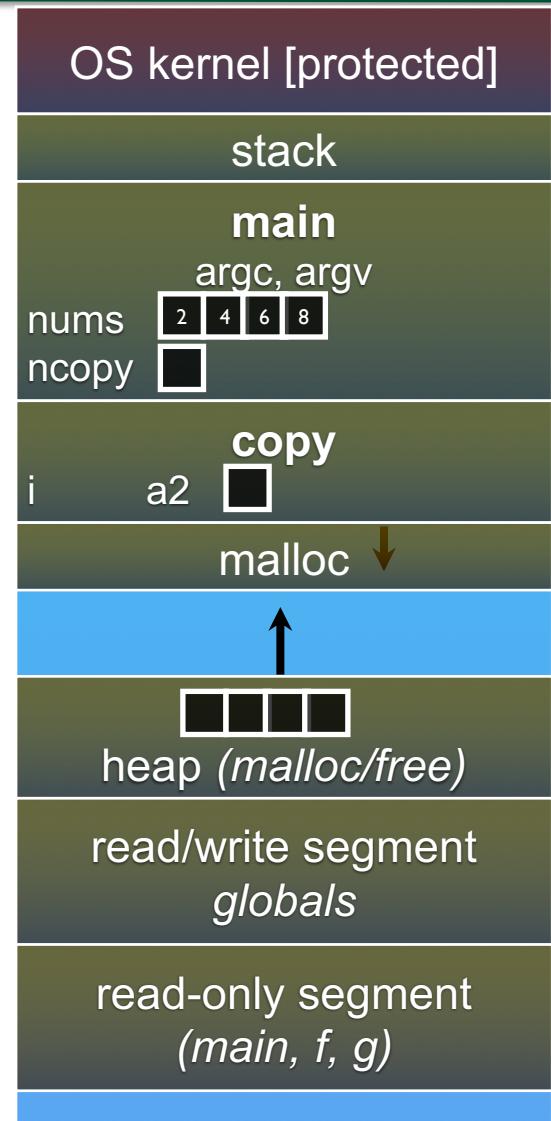
```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    → a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(ncpy);
    return 0;
}
```



Process Address Space in Action

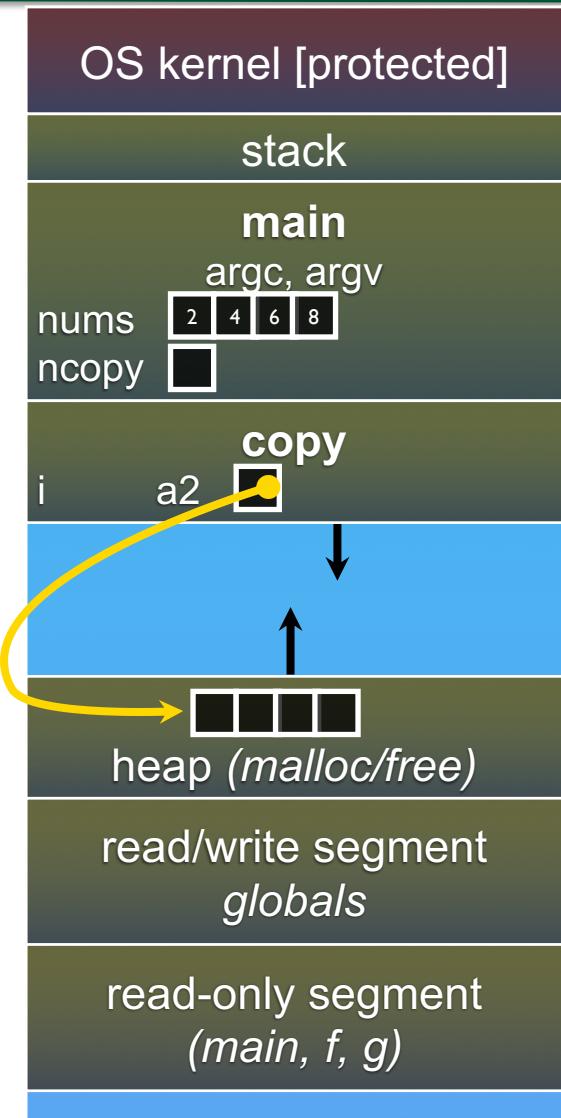
```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    → a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2, 4, 6, 8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(nncpy);
    return 0;
}
```



Process Address Space in Action

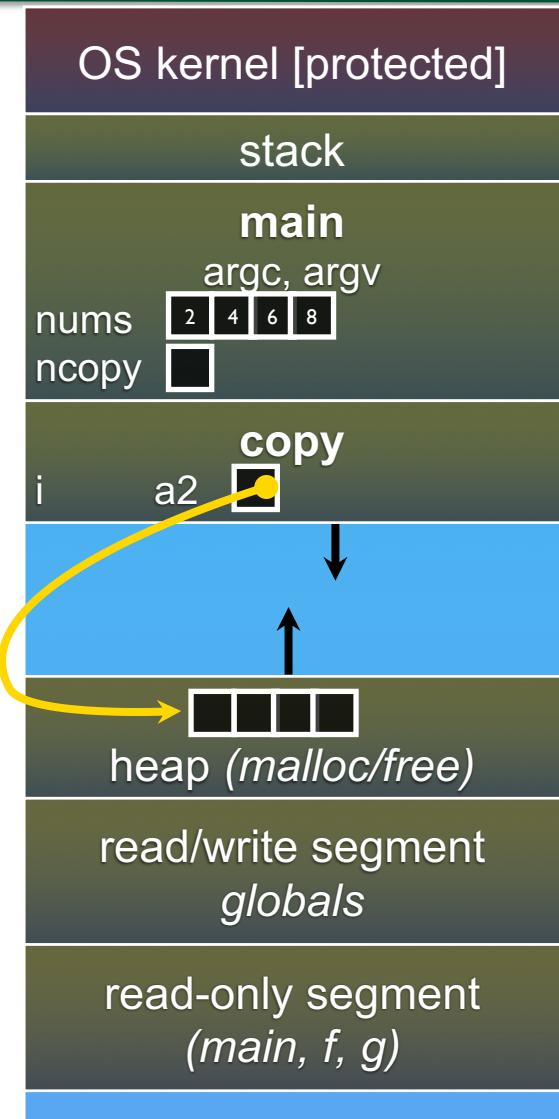
```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(ncpy);
    return 0;
}
```



Process Address Space in Action

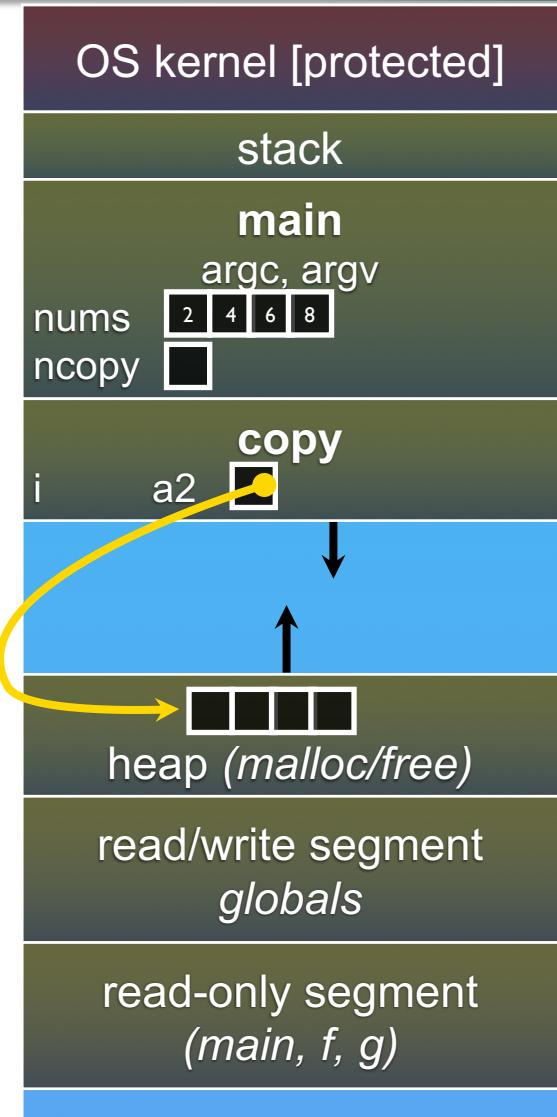
```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(ncpy);
    return 0;
}
```



Process Address Space in Action

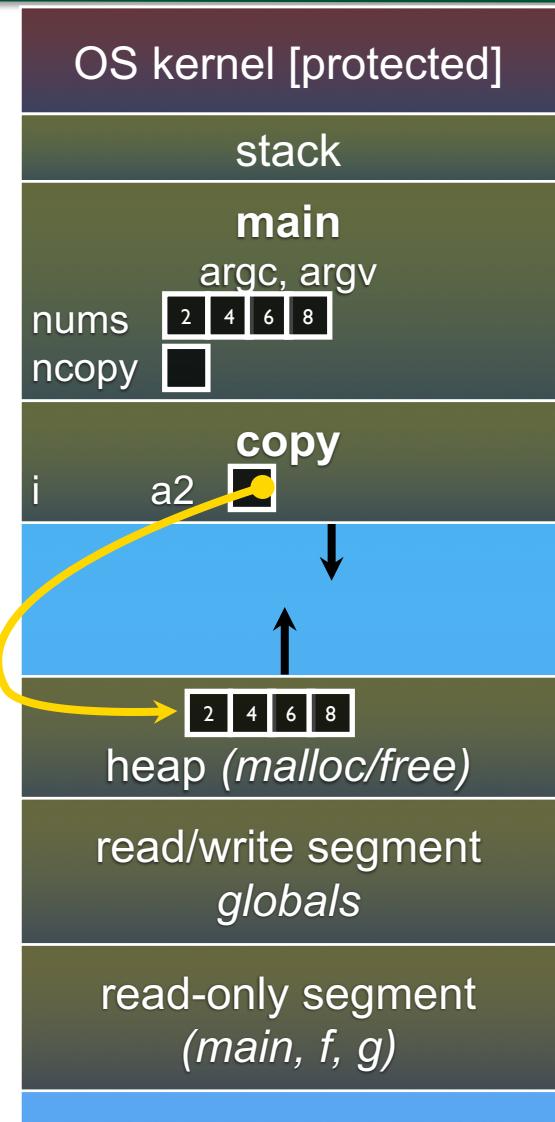
```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(ncpy);
    return 0;
}
```



Process Address Space in Action

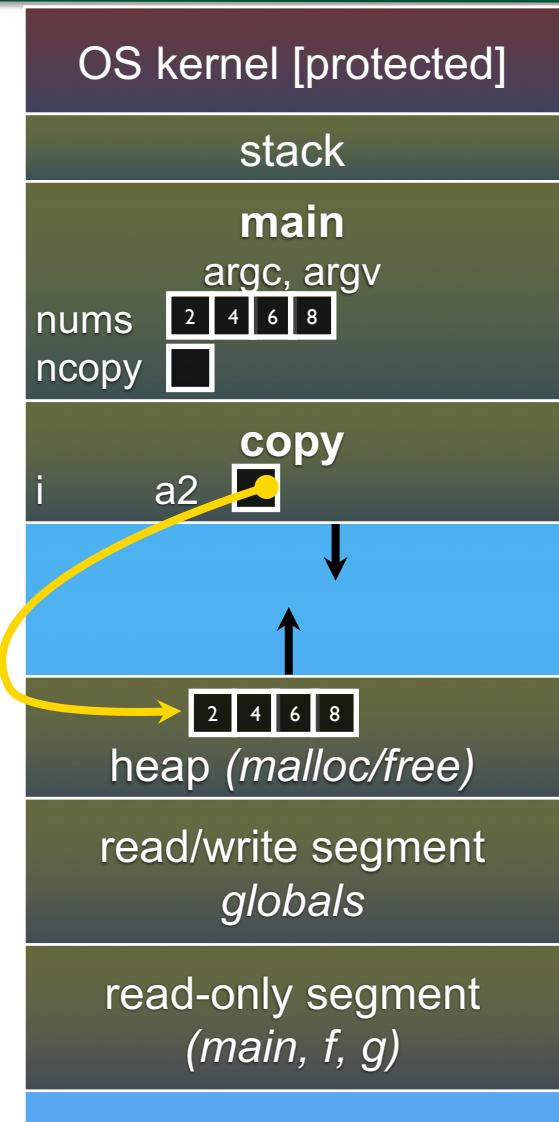
```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(ncpy);
    return 0;
}
```



Process Address Space in Action

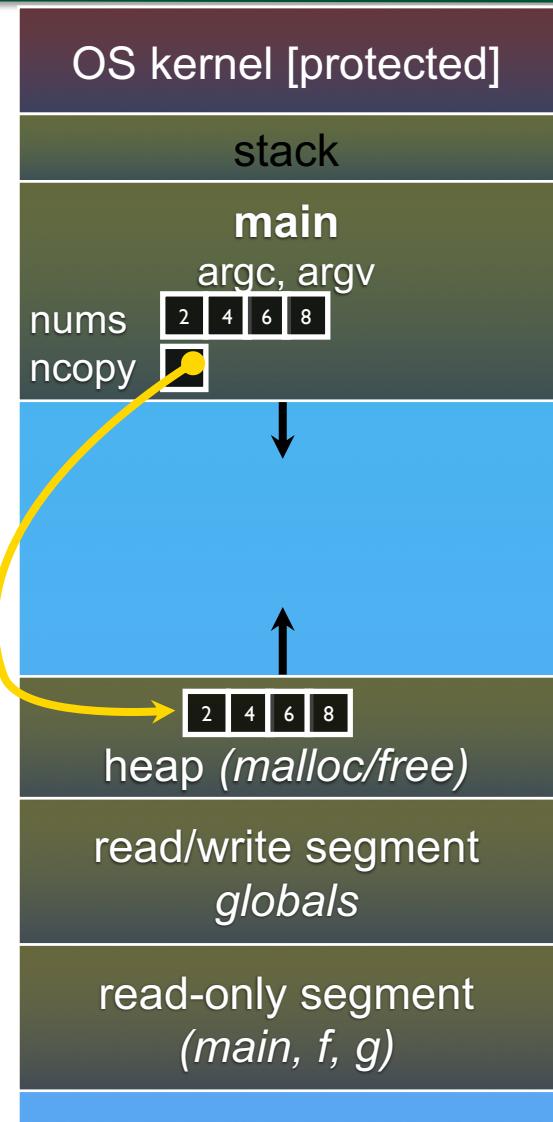
```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(ncpy);
    return 0;
}
```



Process Address Space in Action

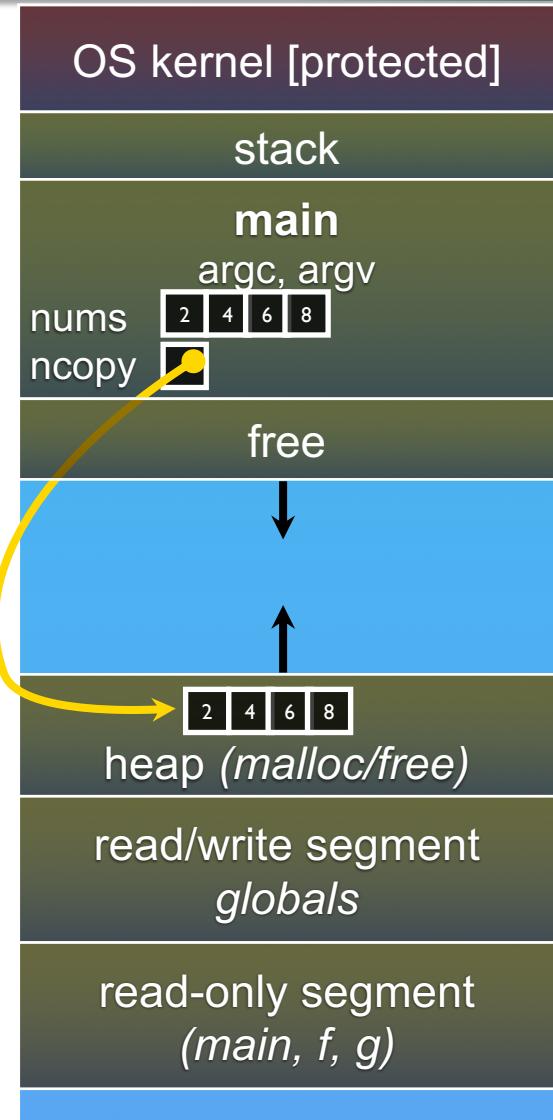
```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(ncpy);
    return 0;
}
```



Process Address Space in Action

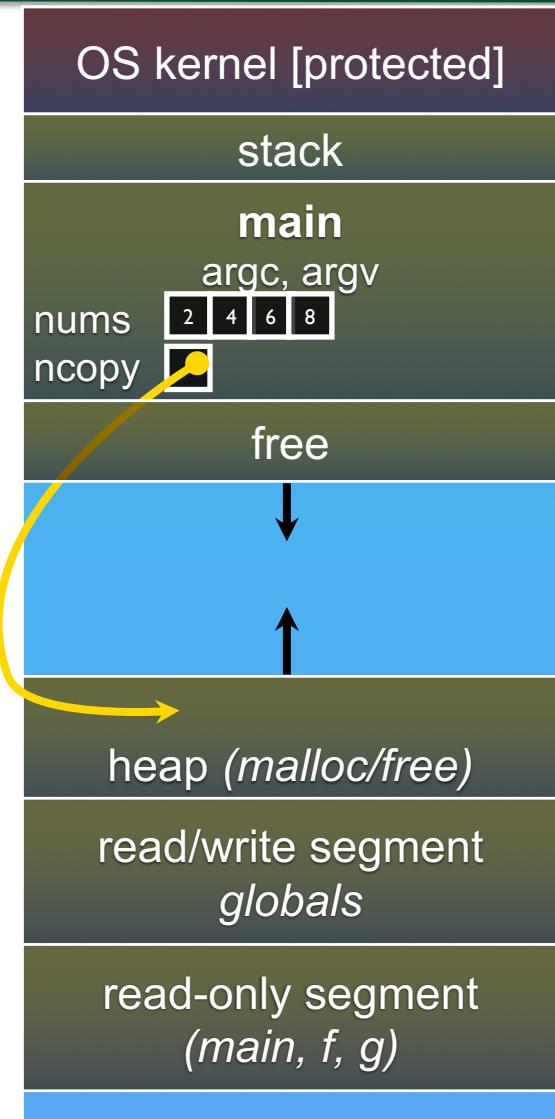
```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(ncpy);
    return 0;
}
```



Process Address Space in Action

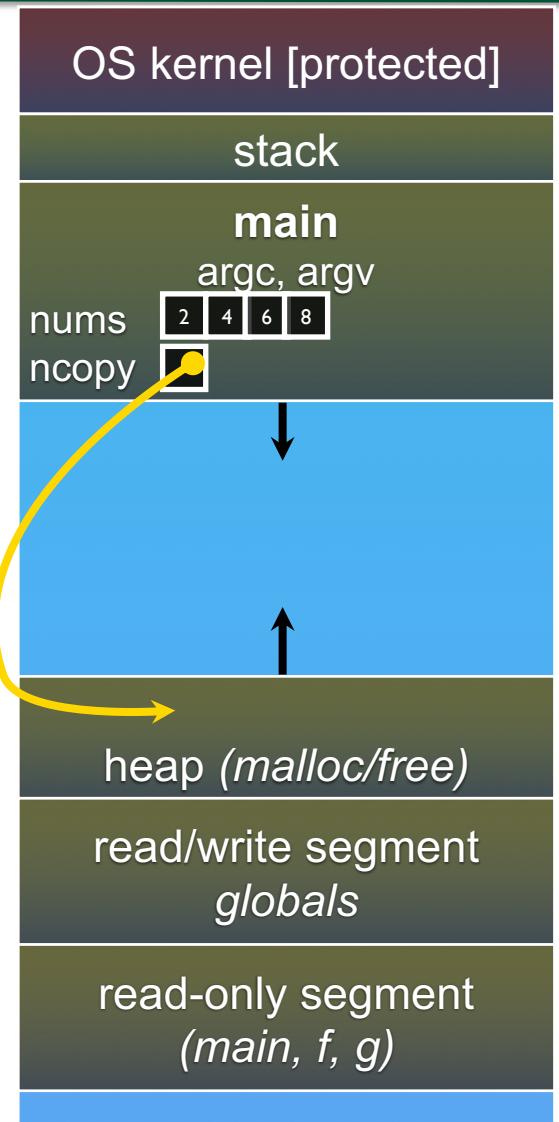
```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(ncpy);
    return 0;
}
```



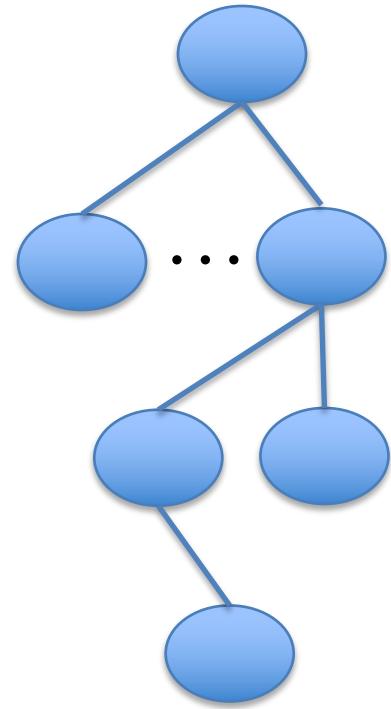
Process Creation

- What happens?
 - New process object in the OS kernel is created
 - ◆ build process data structures
 - Allocate address space (abstract resource)
 - ◆ later, allocate actual memory (physical resource)
 - Add to execution (ready) queue
 - ◆ make it runnable
- Hmm, who created the first process and how?
- Is the OS a process?



Process Creation Options (Parent and Child)

- Process hierarchy options
 - Parent process create children processes
 - Child processes can create other child processes
 - Tree of processes
- Resource sharing options
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution options
 - Parent and children execute concurrently
 - Parent waits until children terminate
- Address space
 - Child duplicate of parent
 - Child has a program loaded into it



Executing a Process

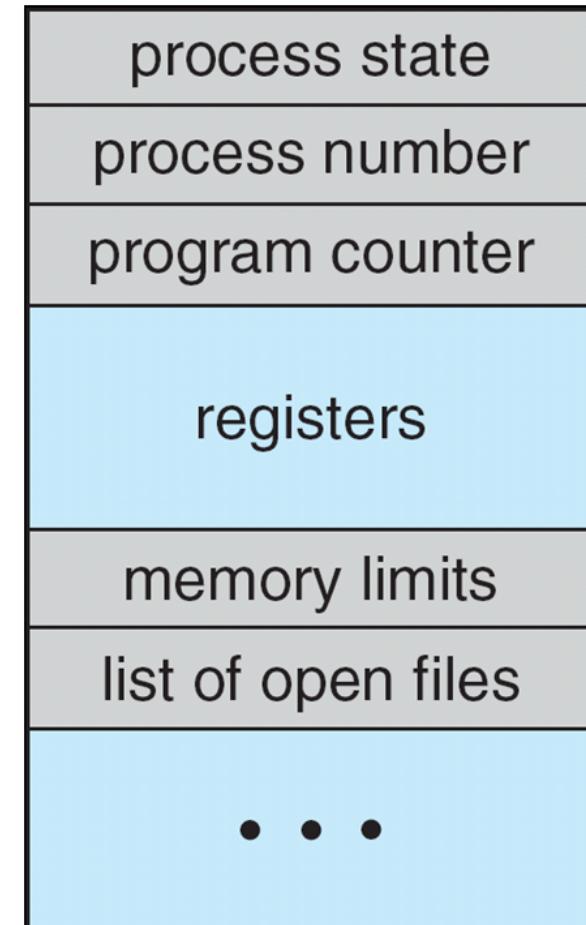
- What is required?
- Registers store state of execution in CPU
 - Stack pointer
 - Data registers
- Program count to indicated what to execute?
 - CPU register holding address of next instruction
- A “thread of execution”
 - Able to execute instructions
 - Has its own stack
 - Each process has at least 1 thread of execution

Executing a Process

- A process thread executes instructions found in the process's address space ...
 - Usually the text segment
- ... until a trap or interrupt
 - Time slice expires (timer interrupt)
 - Another event (e.g., interrupt from other device)
 - Exception (program error, oops)
 - System call (switch to kernel mode)

Process Control Block (PCB)

- Information associated with each process
 - Also called the *task control block*
- Process state: running, waiting, ...
- Program counter
 - Location of instruction to next execute
- CPU registers for process thread
- CPU scheduling information
 - Priorities, scheduling queue pointers, ...
- Memory-management information
 - Memory allocated to the process
- Accounting information
 - CPU used, clock time elapsed, ...
- I/O status information
 - I/O devices allocated to process
 - List of open files

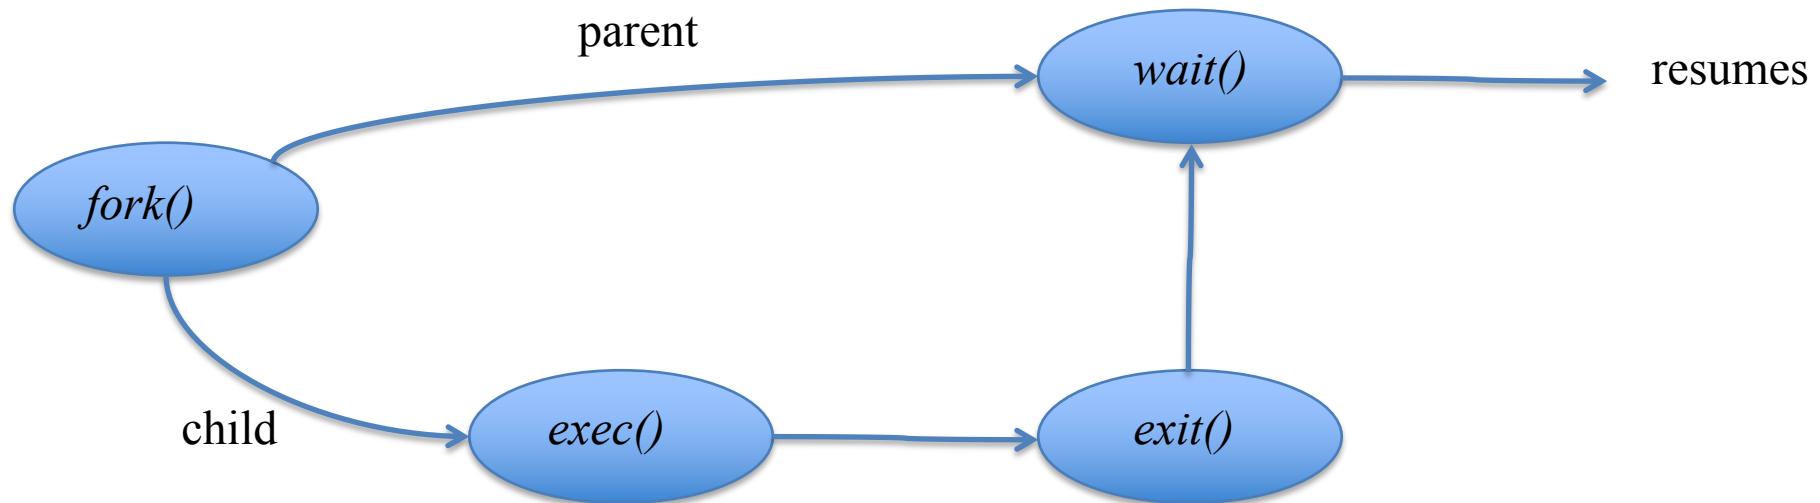


process-specific data
structure in the OS

Program Creation System Calls

- *fork()*
 - Copy address space of parent and all threads
- *forkl()*
 - Copy address space of parent and only calling thread
- *vfork()*
 - Do not copy the parent's address space
 - Share address space between parent and child
- *exec()*
 - Load new program and replace address space
 - Some resources may be transferred (open file descriptors)
 - Specified by arguments

Process Creation with New Program



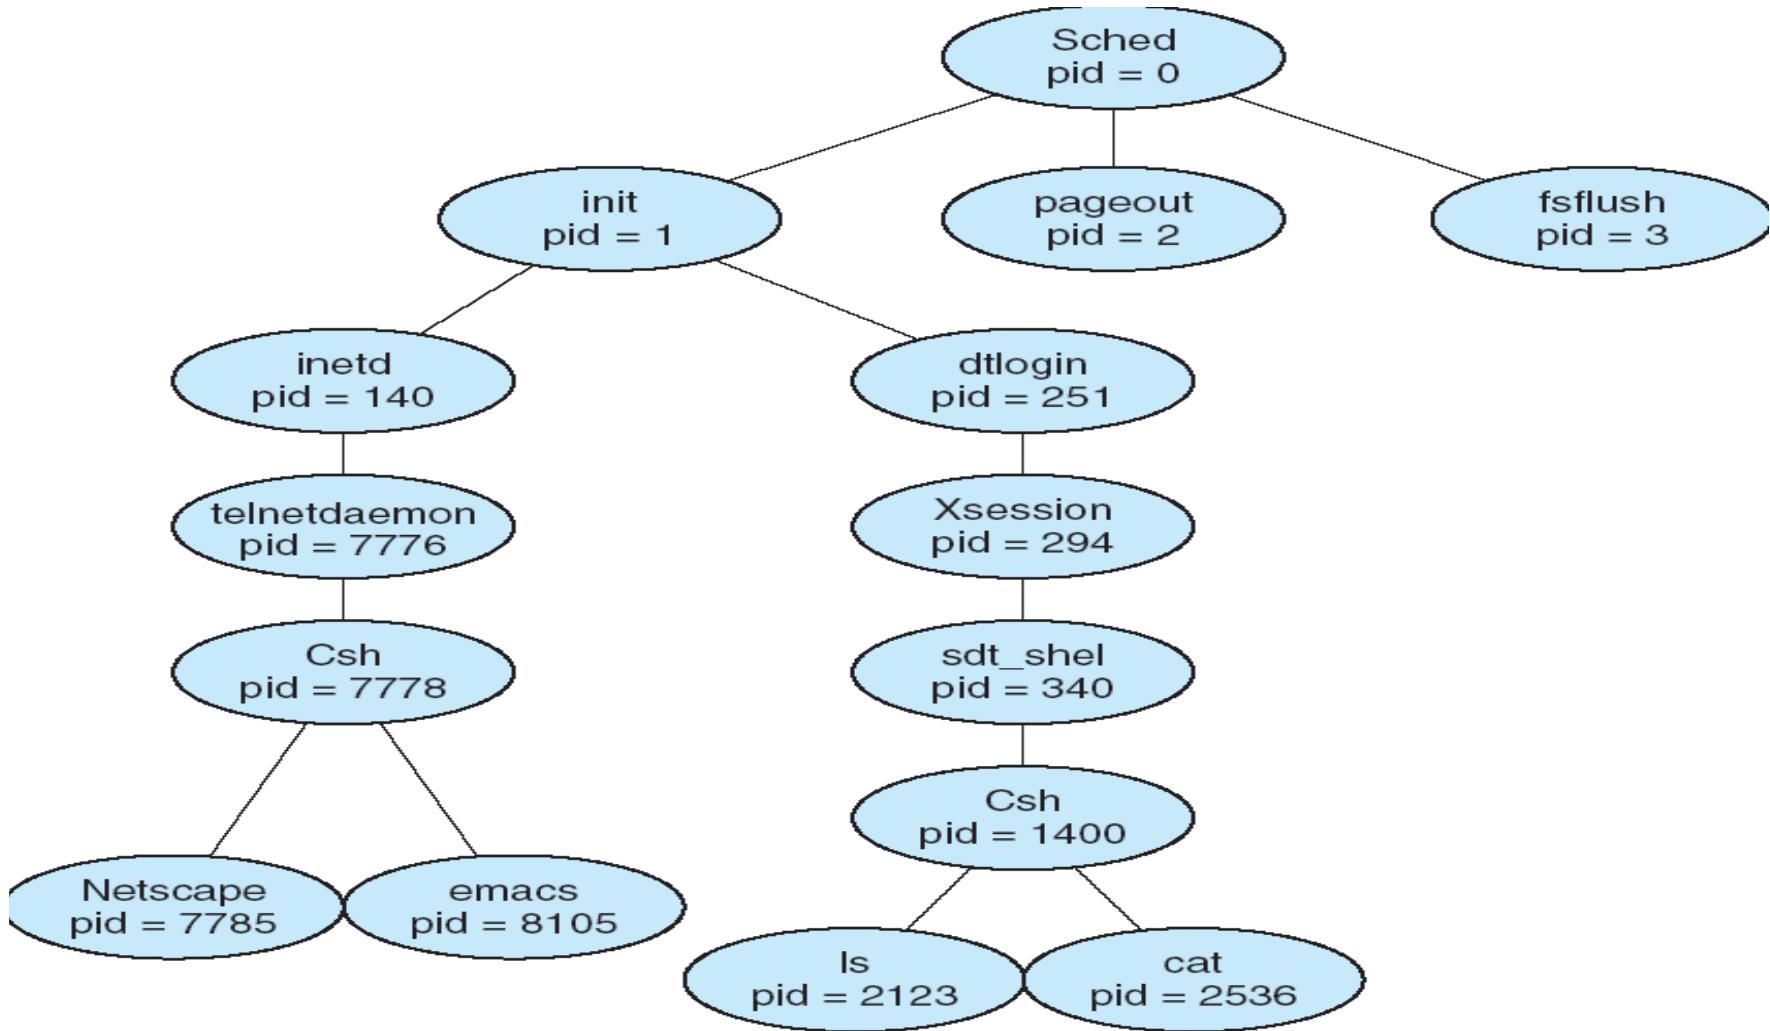
- Parent process calls *fork()* to spawn child process
 - Both parent and child return from *fork()*
 - Continue to execute the same program
- Child process calls *exec()* to load a new program

C Program Forking Separate Process

```
int main( )
{
pid_t pid;
/* fork another process */
pid = fork( );
if (pid < 0) { /* error occurred */
    fprintf(stderr, "Fork Failed");
    exit(-1);
}
else if (pid == 0) { /* child process */
    execvp("/bin/ls", "ls", NULL); /* exec a file */
}
else { /* parent process */
    /* parent will wait for the child to complete */
    wait(NULL);
    printf ("Child Complete");
    exit(0);
}
}
```

execl, execlp, execle,
execv, execvp, execvpe
all execute a file and are frontends
to execve

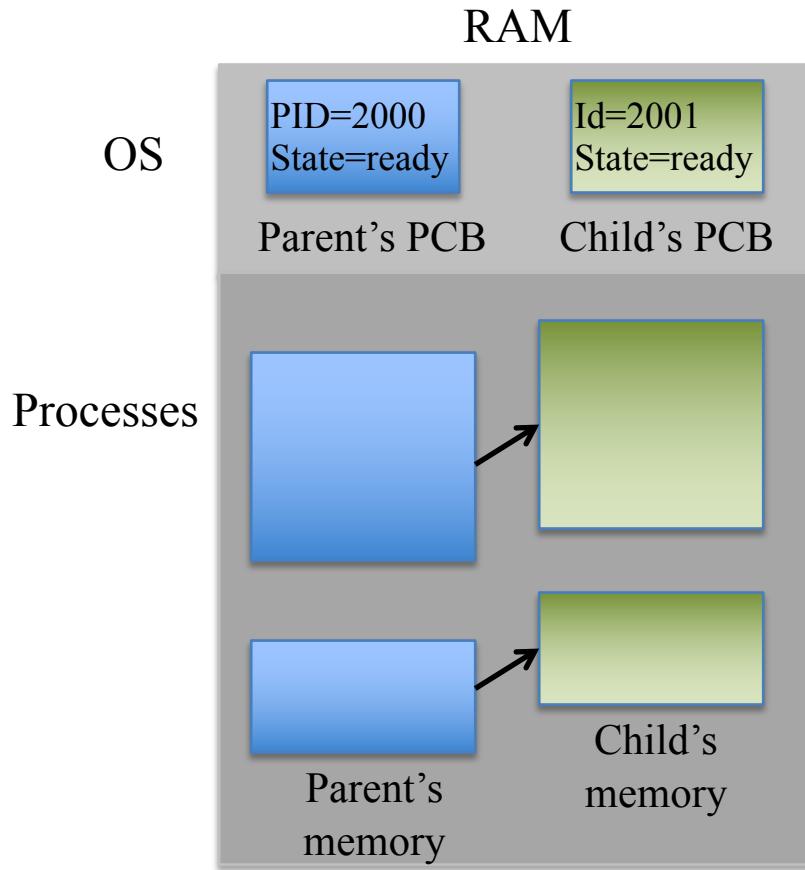
A tree of processes on a typical system



Process Termination

- Process executes last statement and asks the operating system to delete it (*exit()*)
 - Output data from child to parent (via *wait()*)
 - Process' resources are deallocated by operating system
- Parent may terminate execution of children processes (*abort()*)
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - If parent is exiting
 - ◆ some operating systems do not allow child to continue if parent terminates
 - ◆ all children terminated - cascading termination

Process Layout



1. PCB with new PID created
2. Memory allocated for child initialized by copying over from the parent
3. If parent had called `wait()`, it is moved to a waiting queue
4. If child had called `exec()`, its memory is overwritten with new code and data
5. Child added to ready queue and is all set to go now!

Effects in memory after parent calls `fork()`

Relocatable Memory

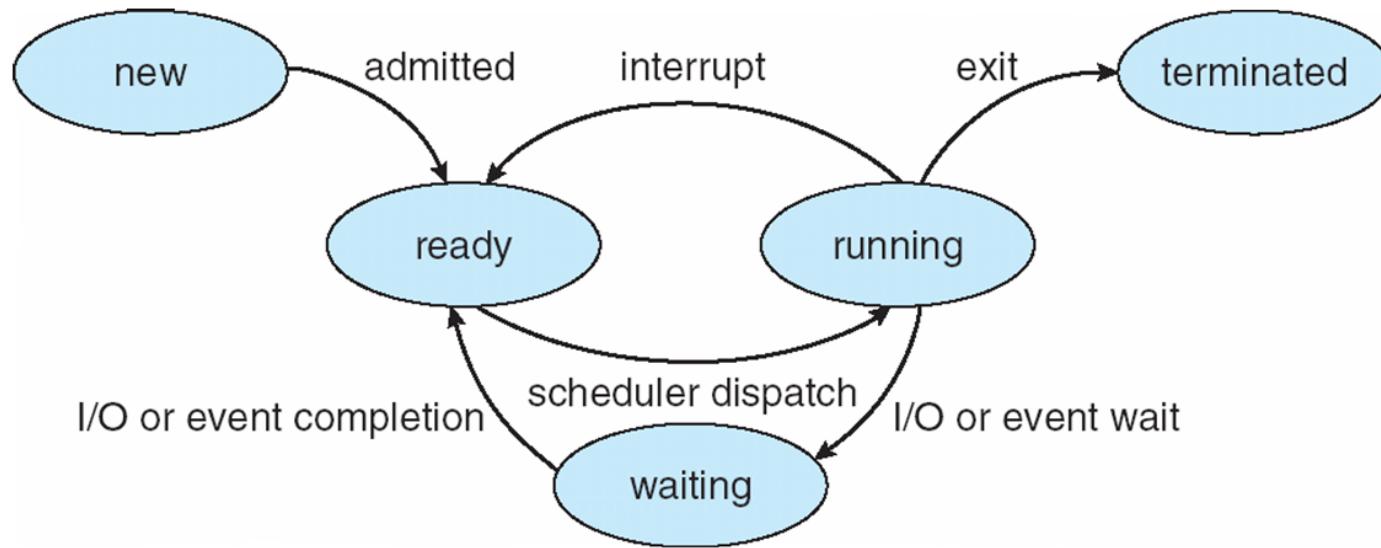
- Program instructions use addresses that logically start at 0
- Cannot place all programs in memory at physical address 0
- Relocation is the mechanism needed that enables the OS to place a program in an arbitrary location in memory
 - Gives the programmer the impression that they own the processor and the memory
- Program is loaded into memory at specific locations
 - Need some form of address translation (relocation) to do this
 - ◆ base-limit, segmentation, paging, virtual memory
 - We will talk about this later
- Also, may need to share program code across processes

Process State

- What do we need to track about a process?
 - How many processes?
 - What's the state of each of them?
- Process table
 - Kernel data structure tracking processes on system
- Process control block
 - Structure for tracking process context

Scheduling Processes

- Processes transition among execution states



Process States

□ Running

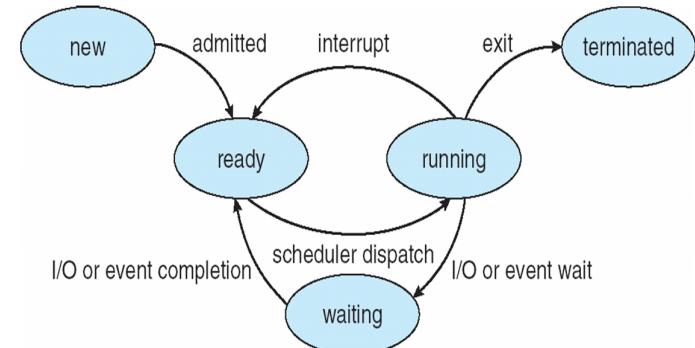
- Process is executing in the processor and in memory with all resources to run

□ Ready

- Process in memory with all resources to run, but is waiting for dispatch onto a processor

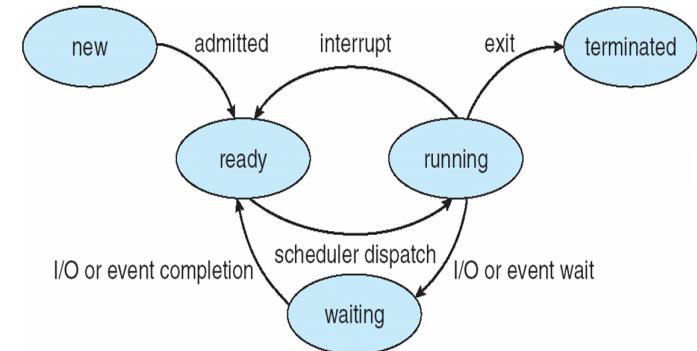
□ Waiting

- Process is not running, instead waiting for some event to occur



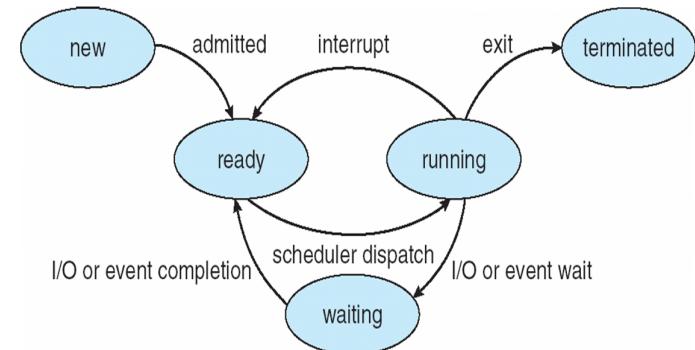
State Transitions

- New Process ==> Ready
 - Allocate resources necessary to run
 - Place process on process queue (usually at end)
- Ready ==> Running
 - Process is at the head of process queue
 - Process is scheduled onto an available processor
- Running ==> Ready
 - Process is interrupted
 - ◆ usually by a timer interrupt
 - Process could still run, in that it is not waiting something
 - Placed back on the process queue



State Transitions: Page Fault Handling

- Running ==> Waiting
 - Either something exceptional happened that caused an interrupt to occur (e.g., page fault exception) ...
 - ... or the process needs to wait on some action (e.g., it made a system call or requested I/O)
 - Process must wait for whatever event happened to be serviced
- Waiting ==> Ready
 - Event has been satisfied so that the process can return to run
 - Put it on the process queue
- Ready ==> Running
 - As before...



State Transitions: Other Issues

□ Priorities

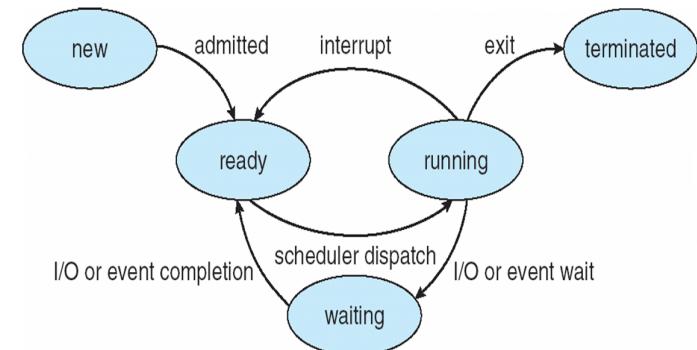
- Can provide policy indicating which process should run next
 - ◆ more when we discuss scheduling...

□ Yield

- System call to give up processor voluntarily
- For a specific amount of time (sleep)

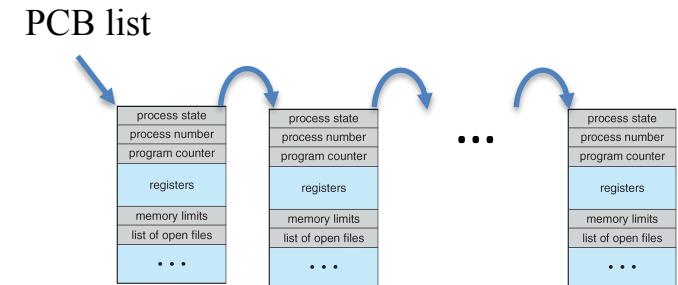
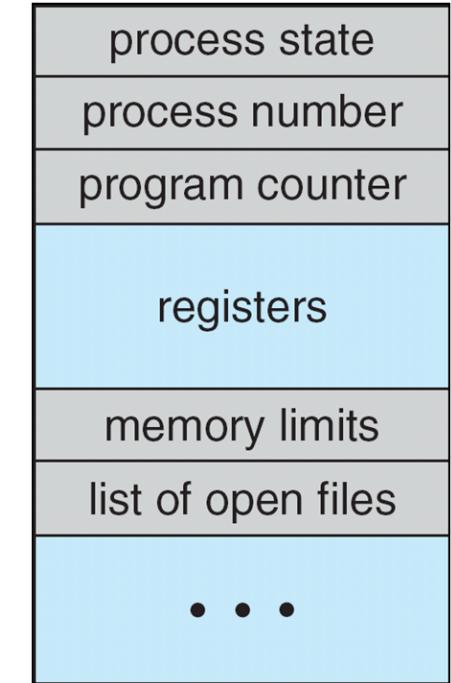
□ Exit

- Terminating signal (Ctrl-C)



Process Control Block

- Information (metadata) associated with every process kept by the OS (kernel)
 - Process state
 - Program counter
 - CPU registers
 - Scheduling information
 - Memory management information
 - Accounting information
 - I/O status information
- OS maintains a list of PCBs for all processes
- Process state lists link in the PCBs



Per Process Control Information

- Process state
 - Ready, running, waiting (momentarily)
- Links to other processes
 - Children
- Memory Management
 - Segments and page tables
- Resources
 - Open files
- And much more...

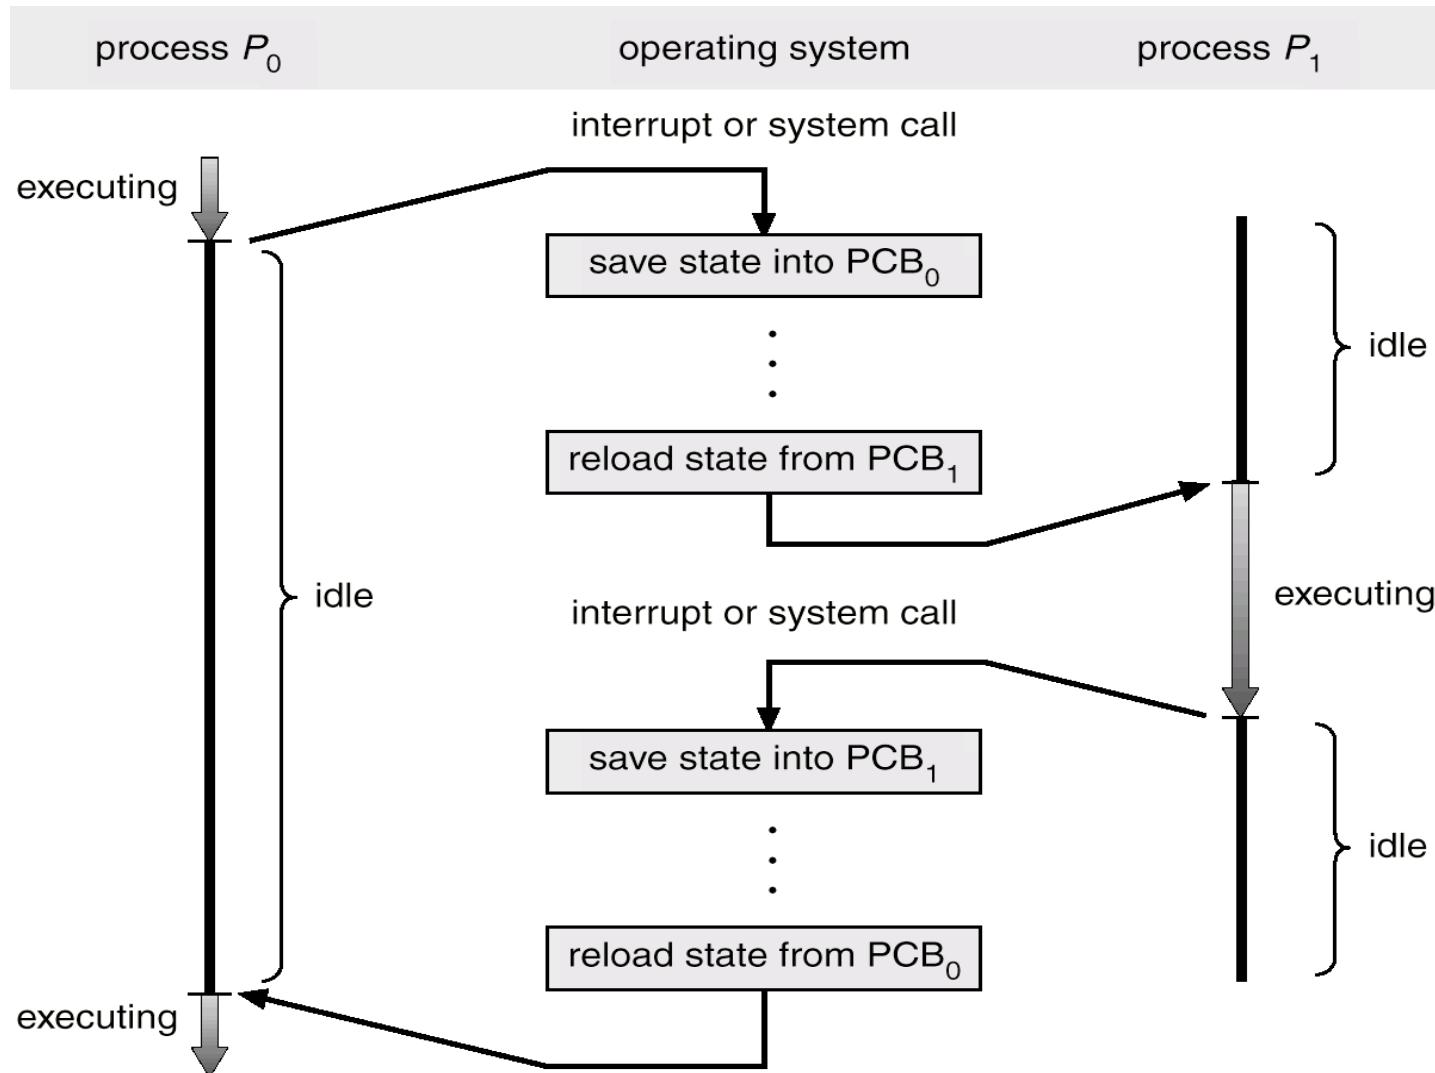
/proc File System

- Linux and Solaris
 - `ls /proc`
 - Process information pseudo-file system
 - Does not contain “real” files, but runtime system information
 - ◆ System memory
 - ◆ Devices mounted
 - ◆ Hardware configuration
 - A directory for each process
- Various process information
 - `/proc/<pid>/io`
I/O statistics
 - `/proc/<pid>/environ`
Environment variables (in binary)
 - `/proc/<pid>/stat`
Process status and info
- Check out “man proc”

Context Switch

- OS switches from one execution context to another
 - One process to another process
 - Interrupt handling
 - Process to kernel (mode transition, not context switch)
- Current process to new process
 - Save the state of the current process
 - ◆ process control block: describes the state of the process in the CPU
 - Load the saved context for the new process
 - ◆ load the new process's process control block into OS and registers
 - Start the new process
- Does this differ if we are running an interrupt handler?

Context Switch



Context Switch Performance

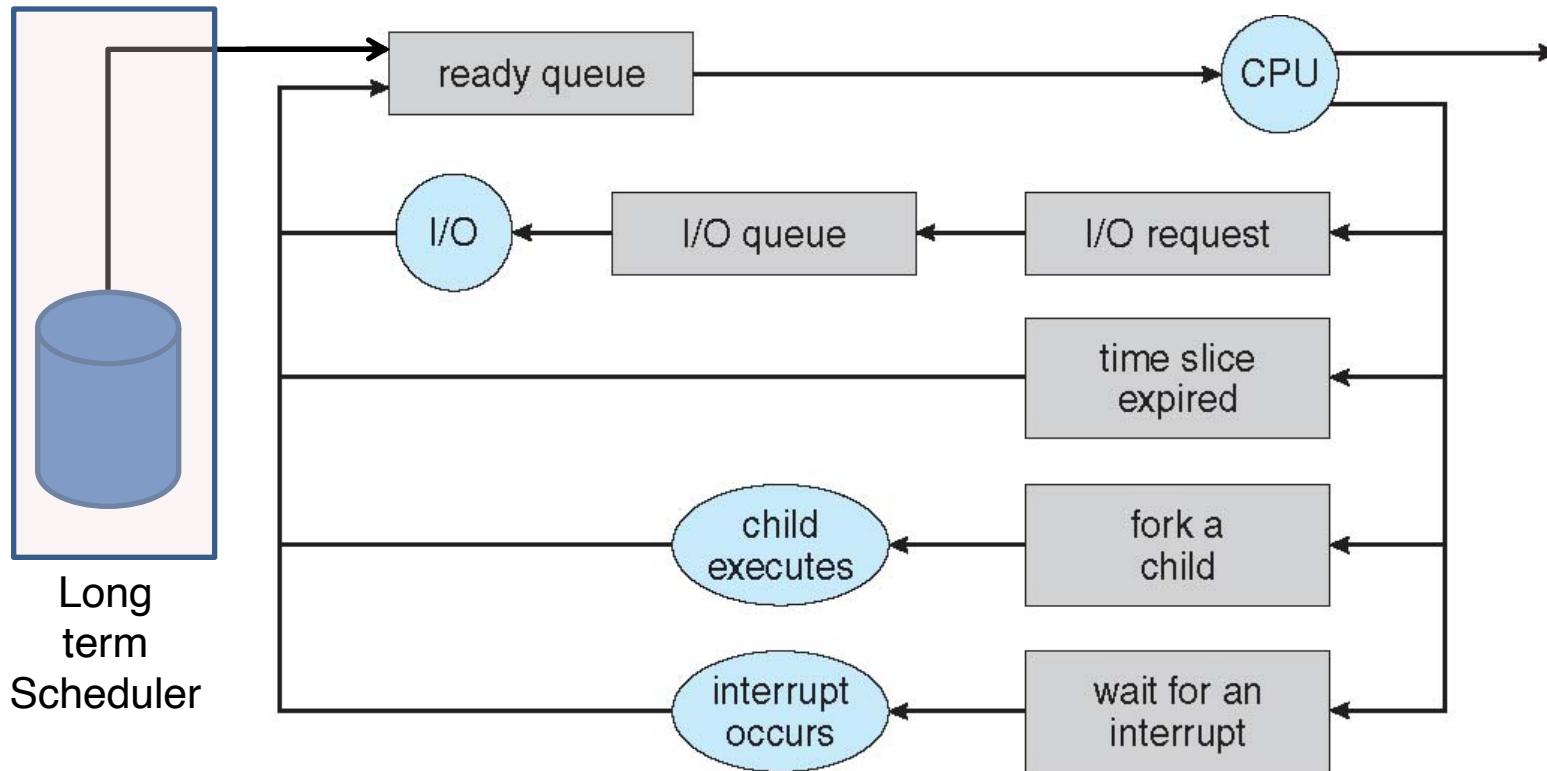
- No useful work is being done during a context switch (i.e., no user process is running)
 - Want to speed up context switch processing
 - If a system call can be done in user mode, then the OS does not have to context switch
- Hardware support helps in context switching
 - Multiple hardware register sets
 - Be able to quickly set up the processor
- However, hardware optimization may conflict
 - Managing address translation tables is difficult
 - Different virtual to physical mappings on different processes

Process Scheduling

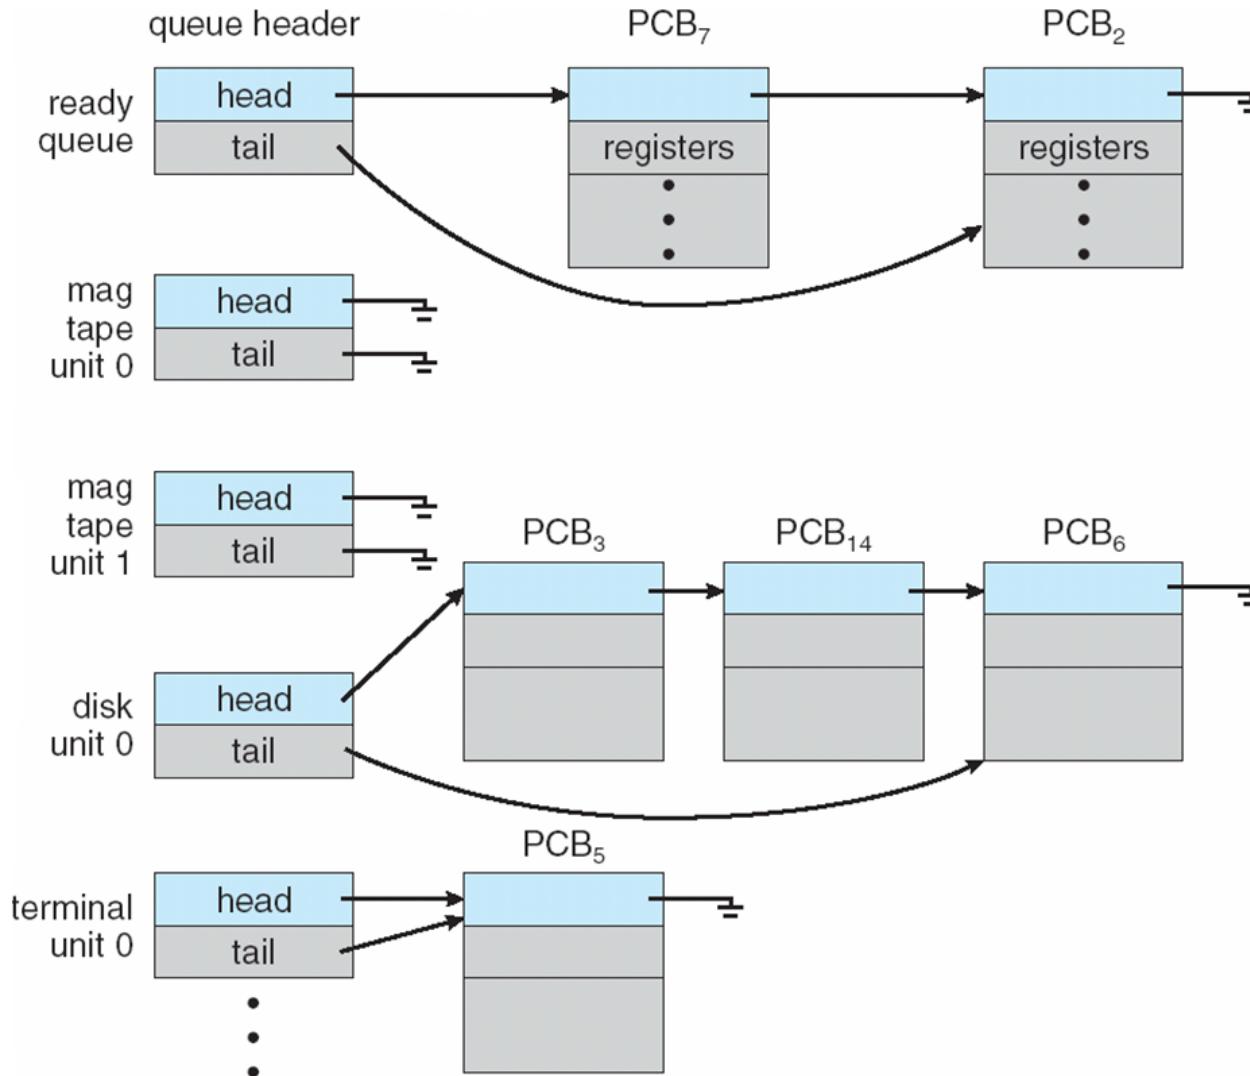
- Maximize CPU use, quickly switch processes onto CPU for time sharing
- Process scheduler selects among available processes for next execution on CPU
- Maintains scheduling queues of processes
 - *Job queue* – set of all processes in the system
 - *Ready queue* – set of all processes residing in main memory, ready and waiting to execute
 - *Device queues* – set of processes waiting for an I/O device
 - Processes migrate among the various queues

Representation of Process Scheduling

- Process scheduling queueing diagram represents:
 - queues, resources, flows
- Processes move through the queues



Ready Queue And Various I/O Device Queues



Schedulers

- *Short-term scheduler* (process scheduler)
 - Selects which process should be executed next
 - Allocates CPU to running process
- *Medium-term scheduler* (multiprogram scheduler)
 - Manages process (jobs) in execution
 - Moves partially executed jobs to/from disk storage
 - Adjusts the degree of multiprogramming
- *Long-term scheduler* (job scheduler)
 - Selects which jobs should be allowed to run
 - Loads process and makes it ready to run

Project 1 – Building a Pseudo-Shell

- A “shell” in Linux (Unix) is an interactive, command-line program that support to a user to run commands, manage control flow, and interact with the OS at a higher level
 - Popular shells: bash, (t)csh, ksh, zsh, ...
- Shells must parse and process user commands
- This project will build a simple shell from scratch
 - Single threaded
 - Synchronous
 - Supports a simple command interface
 - Implements a limited set of functions
- Objective is to develop an interactive program that uses parsing/tokenization and system calls to implement command functions
- Project 2 will add more capabilities for process control

Next Class

- Process scheduling
- Interprocess communication (IPC)



Extra Slides

Process Actions in Client-Server (1)

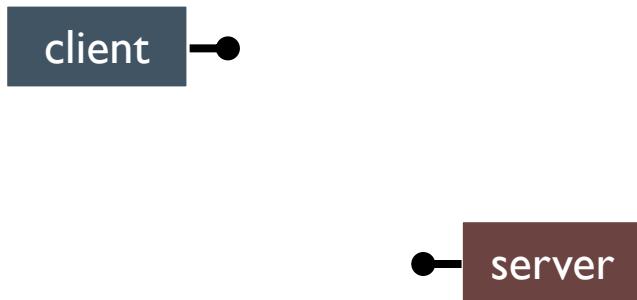
- Example of forking to create a new process
- Consider a web server



server

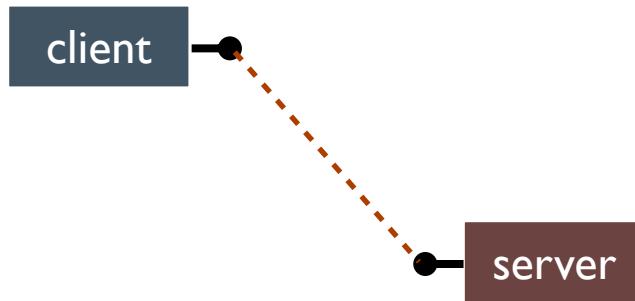
Process Actions in Client-Server (2)

- A remote “client” wants to connect and look at a webpage hosted by the web server



Process Actions in Client-Server (3)

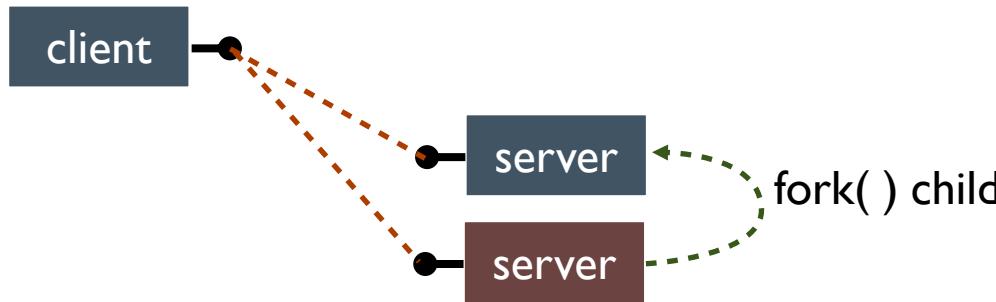
- An interprocess communication is made and a connection is established



- What if the web server only served this client until it was done looking at web pages?
- How can the web server support multiple “concurrent” clients?

Process Actions in Client-Server (4)

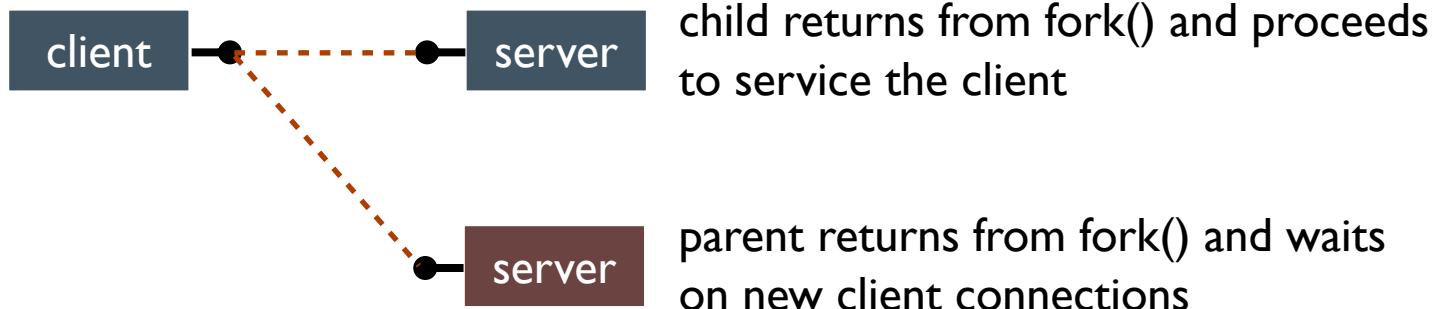
- Create more “concurrency” by creating more processes!



- A `fork()` copies the parent PCB and establish a child process initialized with that PCB
- All of the parent's state is inherited by the child, including the connection to the client

Process Actions in Client-Server (5)

- Responsibility for “servicing” the client is handed off to the “child” server process



- The “parent” server process goes back to waiting for new clients

Process Actions in Client-Server (6)

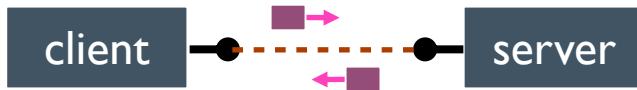
- The “parent” server should disengage from the client by releasing its (duplicate) connection



- The “child” server is now completely responsible for the client connection
- There is still a logical relationship between the child / parent server processes

Process Actions in Client-Server (7)

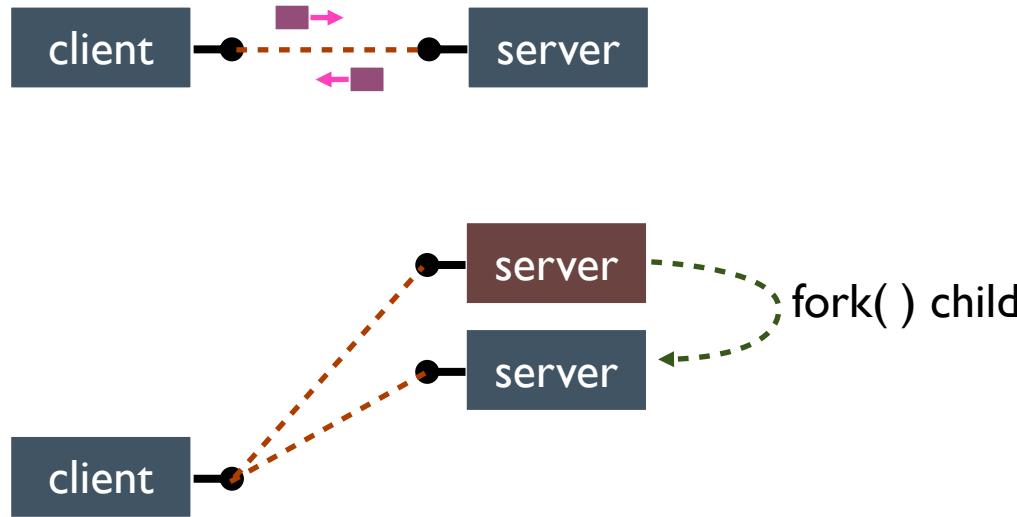
- Communication takes places between the client and “child” server process



- The “parent” server process is not involved

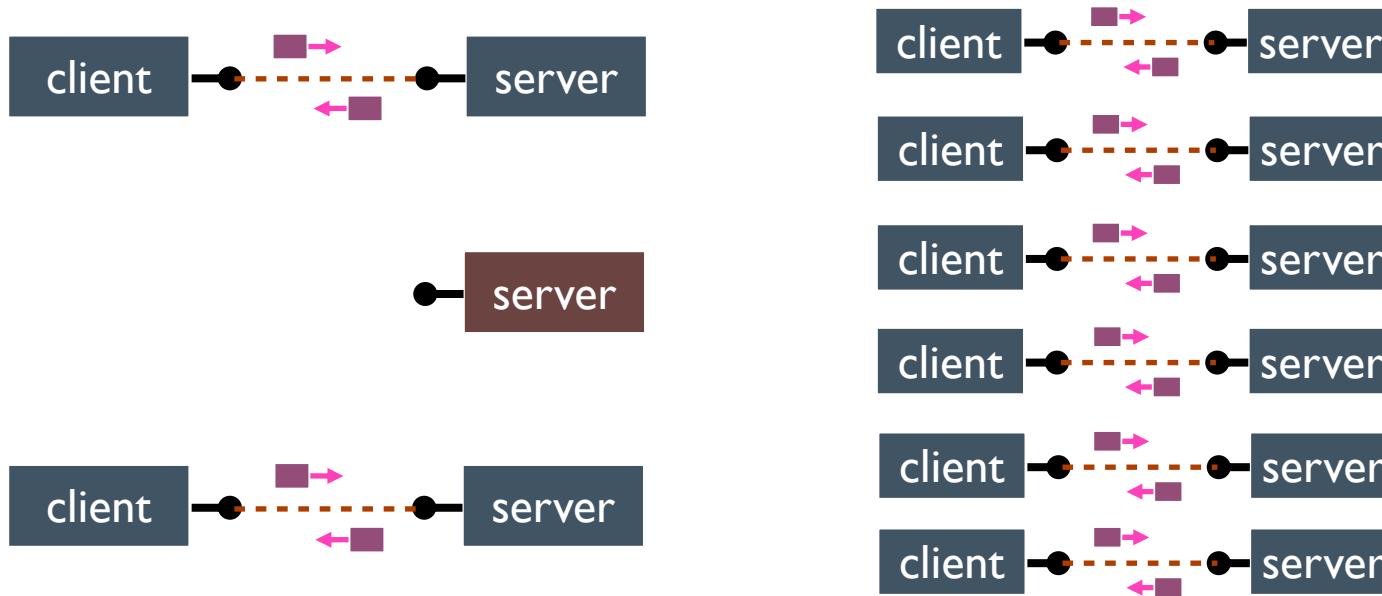
Process Actions in Client-Server (8)

- Now, this can procedure can continue as new client connections come in



Process Actions in Client-Server (9)

- In this way, the amount of “concurrency” increases with more server processes



- It gives both logical isolation between the servers, as well as simplifies the design and functionality

Process Actions in Client-Server (10)

- Many types of servers operate this way
- A objective is to avoid blocking in the server
 - If one server is stalled, others can run

