



CIS 415

Operating Systems

Interprocess Communication

Prof. Allen D. Malony

Department of Computer and Information Science

Spring 2020



UNIVERSITY OF OREGON

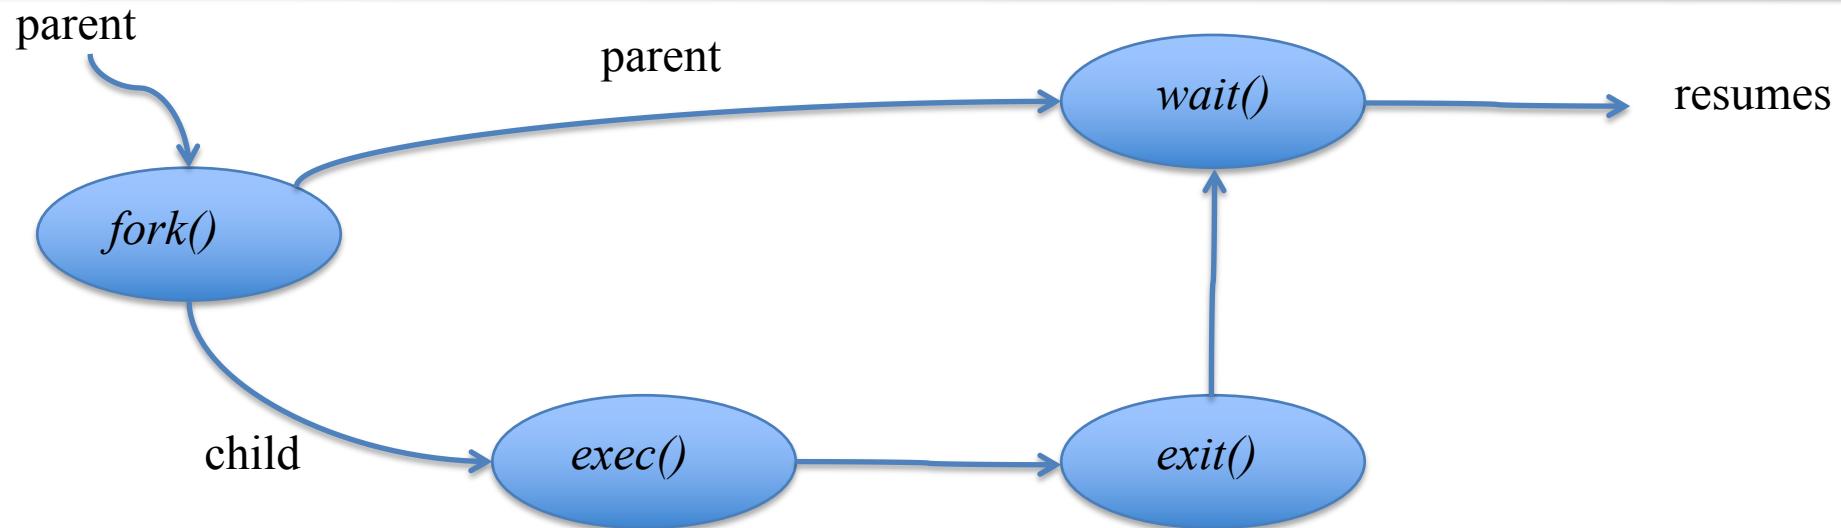
Logistics

- Get started on Project 1
 - Do not put it off
- Lab 2 tomorrow
- Read OSC Chapter 3
- Look at all lectures slides, even if we do not have time to cover in class

Outline

- Quick review of *fork()* and *exec()*
- Brief introduction to process scheduling
- Interprocess communication
- Remote procedure calls

Process Creation with New a Program



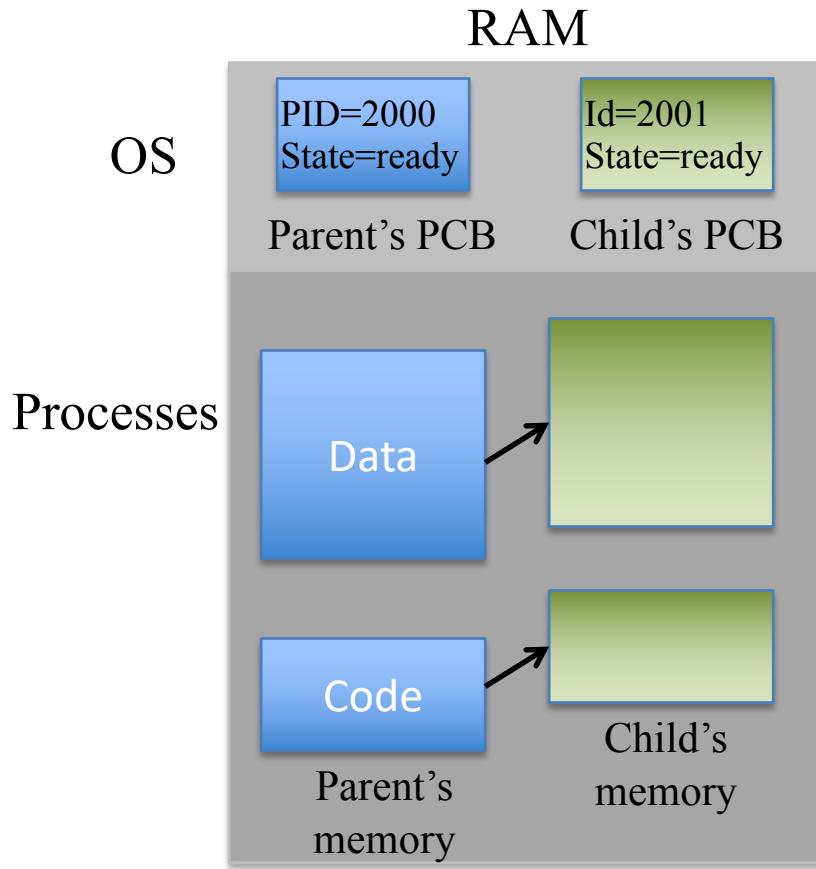
- Parent process calls *fork()* to spawn child process
 - Both parent and child return from *fork()*
 - Continue to execute the same program
- Child process calls *exec()* to load a new program

C Program Forking Separate Process

```
int main( )
{
pid_t pid;
/* fork another process */
pid = fork( );
if (pid < 0) { /* error occurred */
    fprintf(stderr, "Fork Failed");
    exit(-1);
}
else if (pid == 0) { /* child process */
    execvp("/bin/ls", "ls", NULL); /* exec a file */
}
else { /* parent process */
    /* parent will wait for the child to complete */
    wait(NULL);
    printf ("Child Complete");
    exit(0);
}
}
```

execl, execlp, execle,
execv, execvp, execvpe
all execute a file and are frontends
to execve

Process Layout



1. PCB with new PID created
2. Memory allocated for child initialized by copying over from the parent
3. If parent had called *wait()*, it is moved to a waiting queue
4. If child had called *exec()*, its memory is overwritten with new code and data
5. Child added to ready queue and is all set to go now!

Effects in memory after parent calls *fork()*

Analogy

- Doing a *fork()* is analogous to cloning
 - Copy is exactly the same
 - Completely able to live independently
 - Each has its own ability to execute
 - Each has its own resources
 - ◆ gets a copy of the parent's process address space (unless you use another variant of *fork()*)
 - ◆ gets a copy of the parent's PCB (includes its memory management information)
 - Still executes the same code
- What if you want to change the clone's behavior
- Doing an *exec()* is analogous to replacing the brain
 - Executing new program code



Process Communication

- Process cooperation is a fundamental aspect of an OS and of the computing environment it is maintaining on behalf of executing applications
- Processes need to interact and share information
- Process model is a useful way to isolate running programs (separate resources, state, and so on)
 - It can simply programs (no need to worry about other processes)
 - But processes do not always work in isolation, nor do we want them to
 - Sometimes it is easier to design programs if there are multiple processes working together
- How to support process interoperation, communication, and sharing of information
- Discuss a variety of ways
 - Not talk about sharing of files or signals



Process Communication (Interoperation)

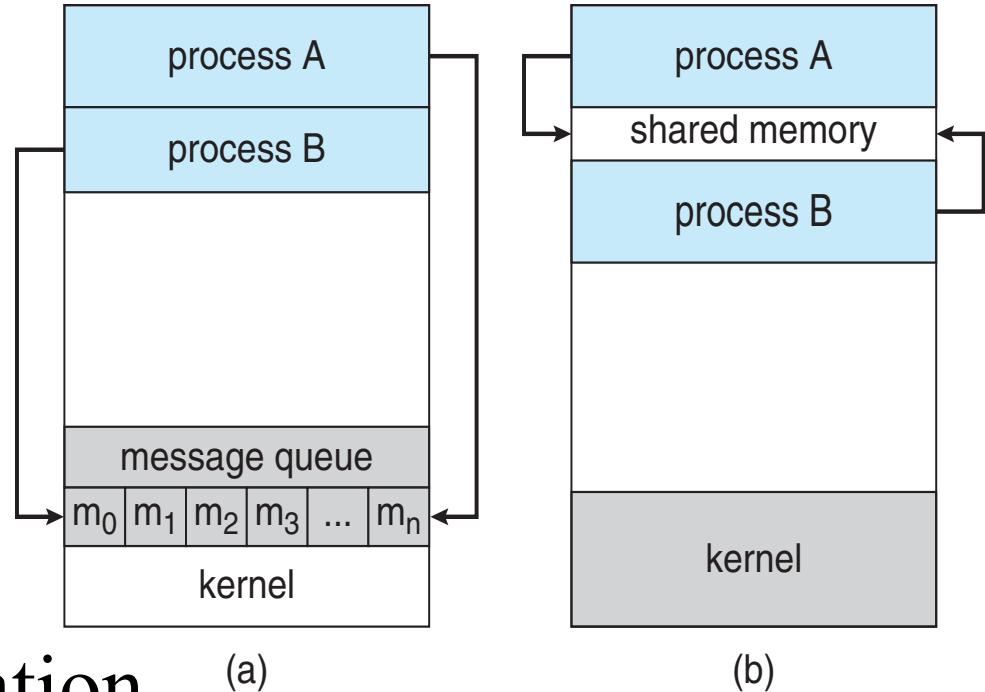
- When is communication necessary?
- Lots of examples in operating systems
 - Kernel/OS access to user process data
 - Processes sharing data via shared memory
 - Processes sharing data via system calls
 - Processes sharing data via file system
 - Threads with access to same data structures
- In general, there are numerous examples in computer science where interoperation is important
 - DB transactions, distributed computing, parallelism

Interprocess Communication (IPC)

- Mechanism for processes to communicate and synchronize
- Logically, we want some sort of *messaging system*
- Logically, an IPC facility would provide 2 operations for processes to communicate:
 - Send(message) (message size fixed or variable)
 - Receive(message)
- However, first there needs to be a communication channel
 - If processes P and Q wish to communicate they open a channel
 - Then exchanging messages can proceed
- How is the communication channel (link) realized?
 - Physically, there are different alternatives
 - Logically, need abstract interfaces, protocols, and properties

IPC Mechanisms

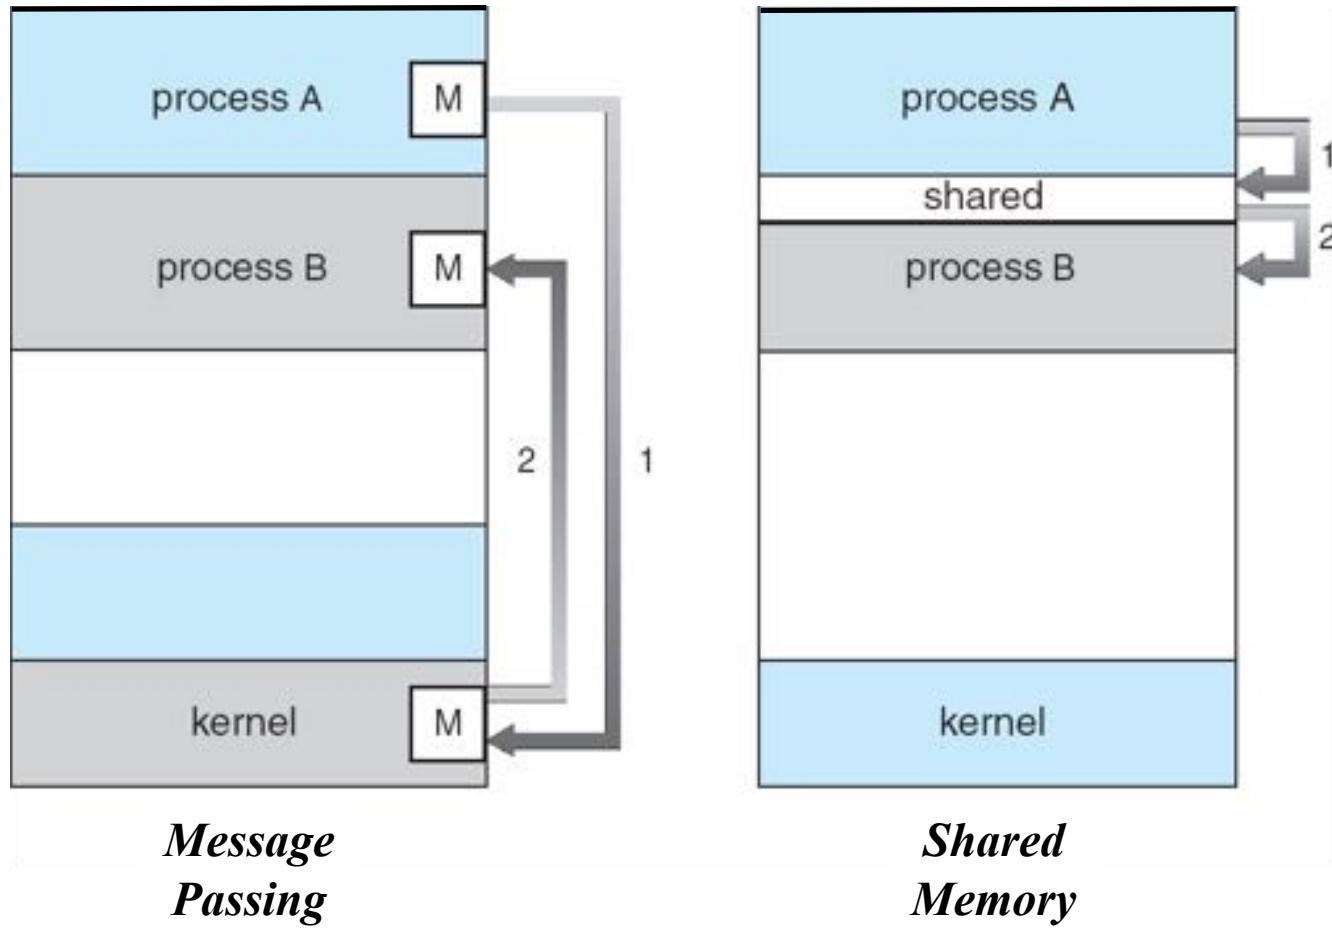
- Interprocess communication (IPC) supports the exchange of data between processes
- Two fundamental methods (OS supported):
 - *Shared memory*
 - ◆ pipes, shared buffer
 - *Message Passing*
 - ◆ mailboxes, sockets
- Which one would you use and why?
- Depends on the application



Implementation Questions

- How are links established?
- Can a link be associated with >2 processes?
- How many links can there be between every pair of communicating processes?
- What is the capacity of the link?
- How is the data flow regulated on the link?
- Is the message size fixed or variable?
- Is the link unidirectional or bi-directional?
- Does an actual network link have to be used or can the memory system be used somehow?

Communication Models



Think Abstractly about the Problem

- Consider two processes sharing a memory region
 - *Producer* writes
 - *Consumer* reads
- Producer action
 - While the buffer not full ...
 - ... stuff can be written (added) to the buffer
- Consumer actions
 - When stuff is in the buffer ...
 - ... it can be read (removed)
- Must manage where new stuff is in the buffer
- Can think of the buffer as being
 - Bounded (Problems?)
 - Unbounded (Realistic?)



Shared Memory -- Producer

```
item nextProduced;                                Circular buffer

while (1) {
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing ... buffer is full */
    buffer[in] = nextProduced;
    in = (in + 1) % BUFFER_SIZE;
}
```

Assume that the consumer is modifying the `out` variable.
Do you see any problems here?

Shared Memory -- Consumer

```
item nextConsumed;
```

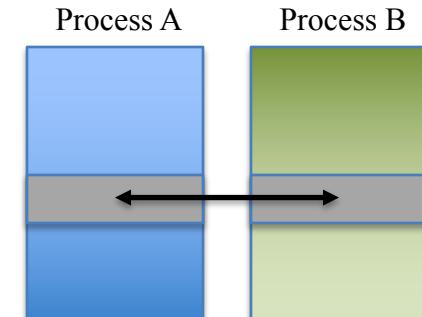
Circular buffer

```
while (1) {  
    while (in == out)  
        ; /* do nothing ... buffer is empty */  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
}
```

Assume that the producer is modifying the `in` variable.
Do you see any problems here?

IPC with Shared Memory

- Communicate by reading/writing from a specific memory location (*in logical memory*)
 - Setup a shared memory region (*segment*) in your process
 - Permit others to attach to the shared memory region
- *shmget(2)* -- create shared memory segment
 - Permissions (read and write)
 - Size
 - Returns an identifier for the segment
- *shmat(2)* -- attach to existing shared memory segment
 - Specify identifier
 - Location in local address space
 - Permissions (read and write)
- Also, operations for detach and control



IPC with Pipes

- Producer-Consumer mechanism for data exchange
 - $prog1 \mid prog2$ (shell notation for pipe)
 - Output of $prog1$ becomes the input to $prog2$
 - More precisely, a connection is made so that the *standard output* of $prog1$ is connected to *standard input* of $prog2$
- OS sets up a fixed-size buffer (*OS manages it*)
 - System calls: $pipe(2)$, $dup(2)$, $popen(2)$
- Producer
 - Write to buffer, if space available
- Consumer
 - Read from buffer if data available

Management of Pipes

□ Buffer management

- A finite region of memory (array or linked-list)
- Wait to produce if no room
- Wait to consume if empty
- Produce and consume complete items

□ Access to buffer

- Write adds to buffer (updates end of buffer)
- Reader removes stuff from buffer (updates start of buffer)
- Both are updating buffer state

□ Issues

- What happens when end is reached (e.g., in finite array)?
- What happens if reading and writing are concurrent?
- Who is managing the pipe?

IPC with Message Passing

- Mechanisms for processes to communicate and to synchronize actions
- Messaging system
 - Processes communicate without shared variables
 - Use messages instead
- Establish communication link
 - Producer sends on link
 - Consumer receives on link
- IPC Operations
 - $Send(P, message)$: send a message to process P
 - $Receive(Q, message)$: receive a message from process Q
- Issues
 - What if a process wants to receive from any process?
 - What if communicating processes are not ready at same time?
 - What size message can a process receive?
 - Can other processes receive the same message from one process?

Synchronous Messaging

- Direct communication from one process to another
- Synchronous send
 - $\text{Send}(P, \text{message})$
 - Producer must wait for the consumer to be ready to receive the message
- Synchronous receive
 - $\text{Receive}(ID, \text{message})$
 - ID could be any process
 - Wait for someone to deliver a message
 - Allocate enough space to receive message
- Synchronous means that both have to be ready!
 - Otherwise process (sender or receiver) is blocked

Properties of Direct Communication Links

- Links are established automatically
- A link is associated with exactly one pair of communicating processes
- Between each pair there exists exactly one link
- The link may be unidirectional, but is usually bi-directional
- Messaging
 - Processes must name each other explicitly

Asynchronous Messaging

- Indirect communication from one process to another
- Asynchronous send
 - *Send(M, message)*
 - Producer sends message to a buffer M (like a mailbox)
 - No waiting (modulo busy mailbox)
- Asynchronous receive
 - *Receive(M, message)*
 - Receive a message from a specific buffer (get your mail)
 - No waiting (modulo busy mailbox)
 - Allocate enough space to receive message
- Asynchronous means that you can send/receive when you're ready
 - What are some issues with the buffer?

Properties of Indirect Communication Link

- Link established only if processes share a mailbox
- A link may be associated with many processes
- Each pair of processes may share several communication links
- Link may be unidirectional or bi-directional
- Messages
 - Messages are directed and received from mailboxes (also referred to as ports)
 - Each mailbox has a unique ID
 - Processes can communicate only if they share a mailbox

Synchronization

- Message passing may be either *blocking* or *non-blocking*
- Blocking is considered synchronous
 - Blocking send
 - ◆ sender is blocked until the message is received
 - Blocking receive
 - ◆ receiver is blocked until a message is available
- Non-blocking is considered asynchronous
 - Non-blocking send
 - ◆ sender sends the message and continue
 - Non-blocking receive
 - ◆ receiver receives: a valid message or a null message
- Different combinations possible
 - If both send and receive are blocking, we have a *rendezvous*

Client-Server Communication

- Sockets
- Remote procedure calls
- Remote method invocation (a la Java)

IPC with Sockets

- Defined as an end point for communication
 - Connect one socket to another (*TCP/IP*)
 - Send/receive message to/from another socket (*UDP/IP*)
- Sockets are named by
 - IP address (roughly, machine)
 - Port number (service: ssh, http, ...)
 - Concatenate the two (**161.25.19.8:1625**)
- Communication consists of a pair of sockets
- Semantics
 - Bidirectional link between a pair of sockets
 - Messages: unstructured stream of bytes
- Connection between
 - Processes on same machine (UNIX domain sockets)
 - Processes on different machines (TCP or UDP sockets)
 - User process and kernel (netlink sockets)

Files and File Descriptors

- POSIX system calls for interacting with files
 - *open()*, *read()*, *write()*, *close()*
 - *open()* returns a *file descriptor*
 - ◆ an integer that represents an open file
 - ◆ inside the OS, it's an index into a table that keeps track of any state associated with your interactions, such as the file position
 - ◆ you pass the file descriptor into *read*, *write*, and *close*
 - File descriptors are kept as part of the process information in the process control block

Networks and Sockets

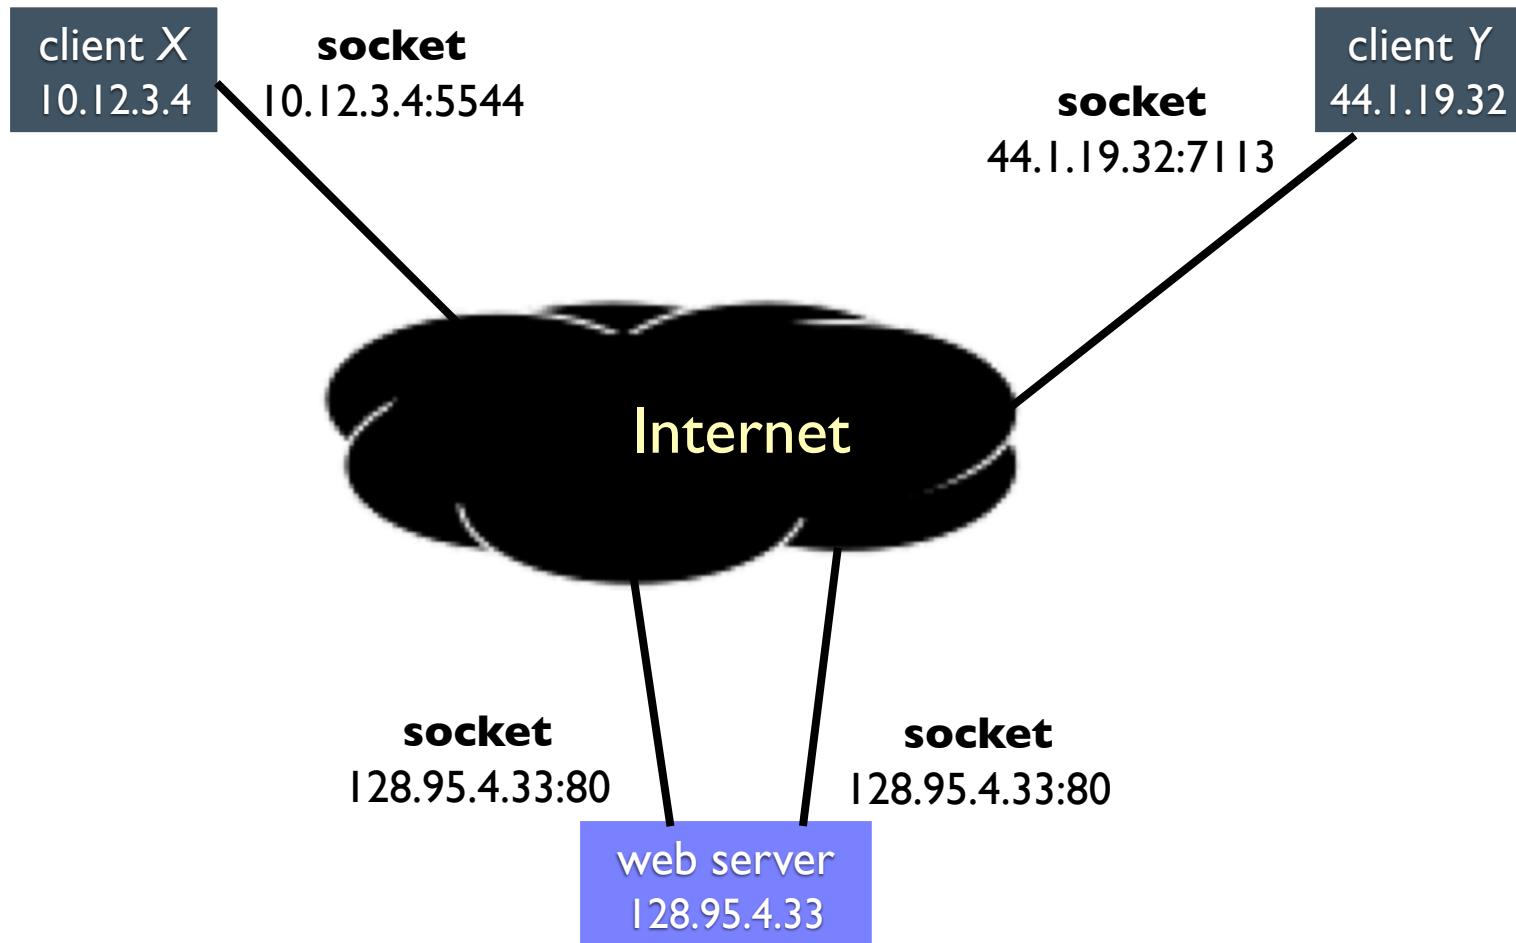
- UNIX likes to make all I/O look like file I/O
 - The good news is that you can use *read()* and *write()* to interact with remote computers over a network!
- File descriptors are used for network communications
 - The socket is the file descriptor
- Just like with files....
 - Your program can have multiple network channels (sockets) open at once
 - You need to pass *read()* and *write()* the socket file descriptor to let the OS know which network channel you want to use

Examples of Sockets

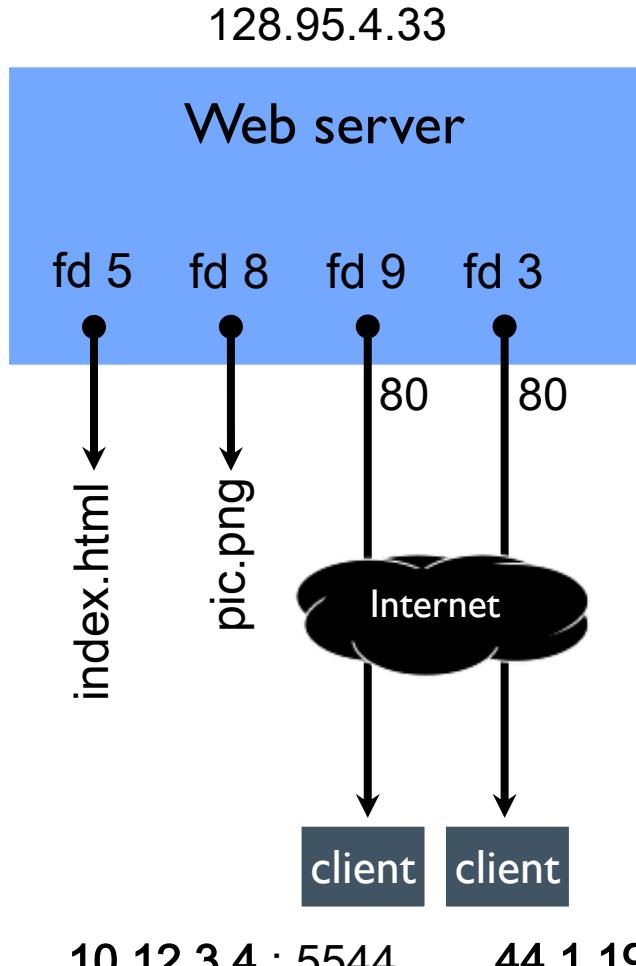
- HTTP / SSL
- email (POP/IMAP)
- ssh
- telnet



IPC and Sockets



File Descriptors



file descriptor	type	connected to?
0	pipe	stdin (console)
1	pipe	stdout (console)
2	pipe	stderr (console)
3	TCP socket	local: 128.95.4.33:80 remote: 44.1.19.32:7113
5	file	index.html
8	file	pic.png
9	TCP socket	local: 128.95.4.33:80 remote: 10.12.3.4:5544

OS's file descriptor table

Types of Sockets

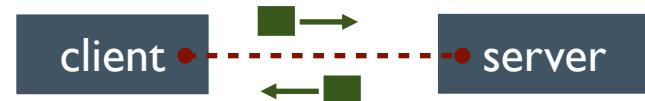
- Stream sockets
 - For connection-oriented, point-to-point, reliable bytestreams
 - ◆ uses TCP, SCTP, or other stream transports
- Datagram sockets
 - For connection-less, one-to-many, unreliable packets
 - ◆ uses UDP or other packet transports
- Raw sockets
 - For layer-3 communication
 - ◆ raw IP packet manipulation

Stream Sockets

- Typically used for client / server communications
 - But also for other architectures, like peer-to-peer



- Client
 - An application that establishes a connection to a server

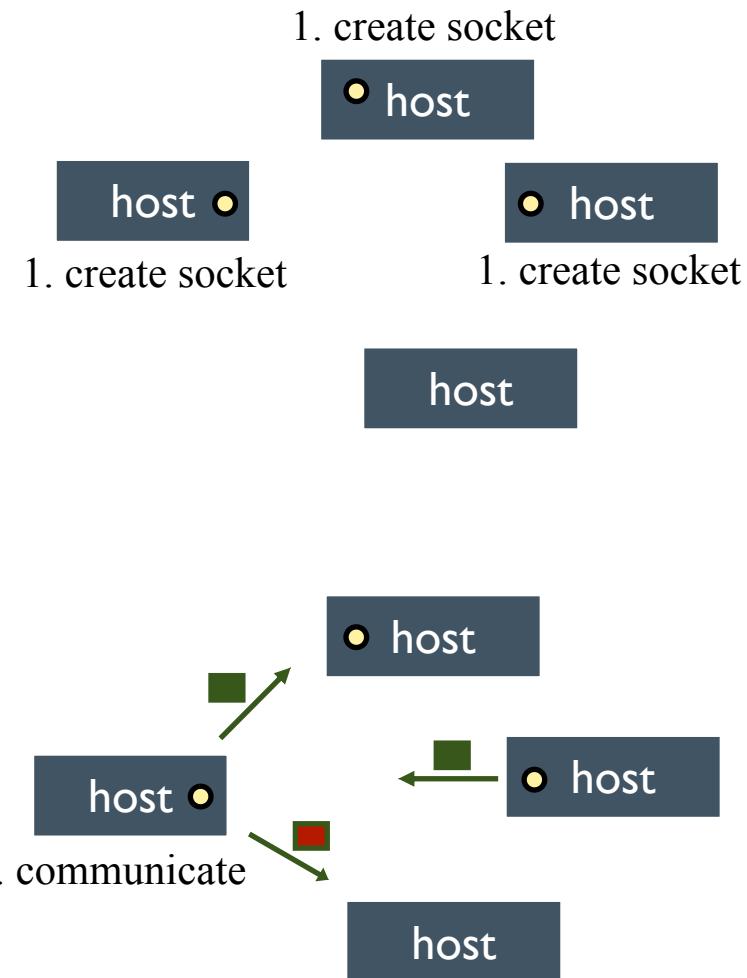


- Server
 - An application that receives connections from clients



Datagram Sockets

- Used less frequently than stream sockets
 - They provide no flow control, ordering, or reliability



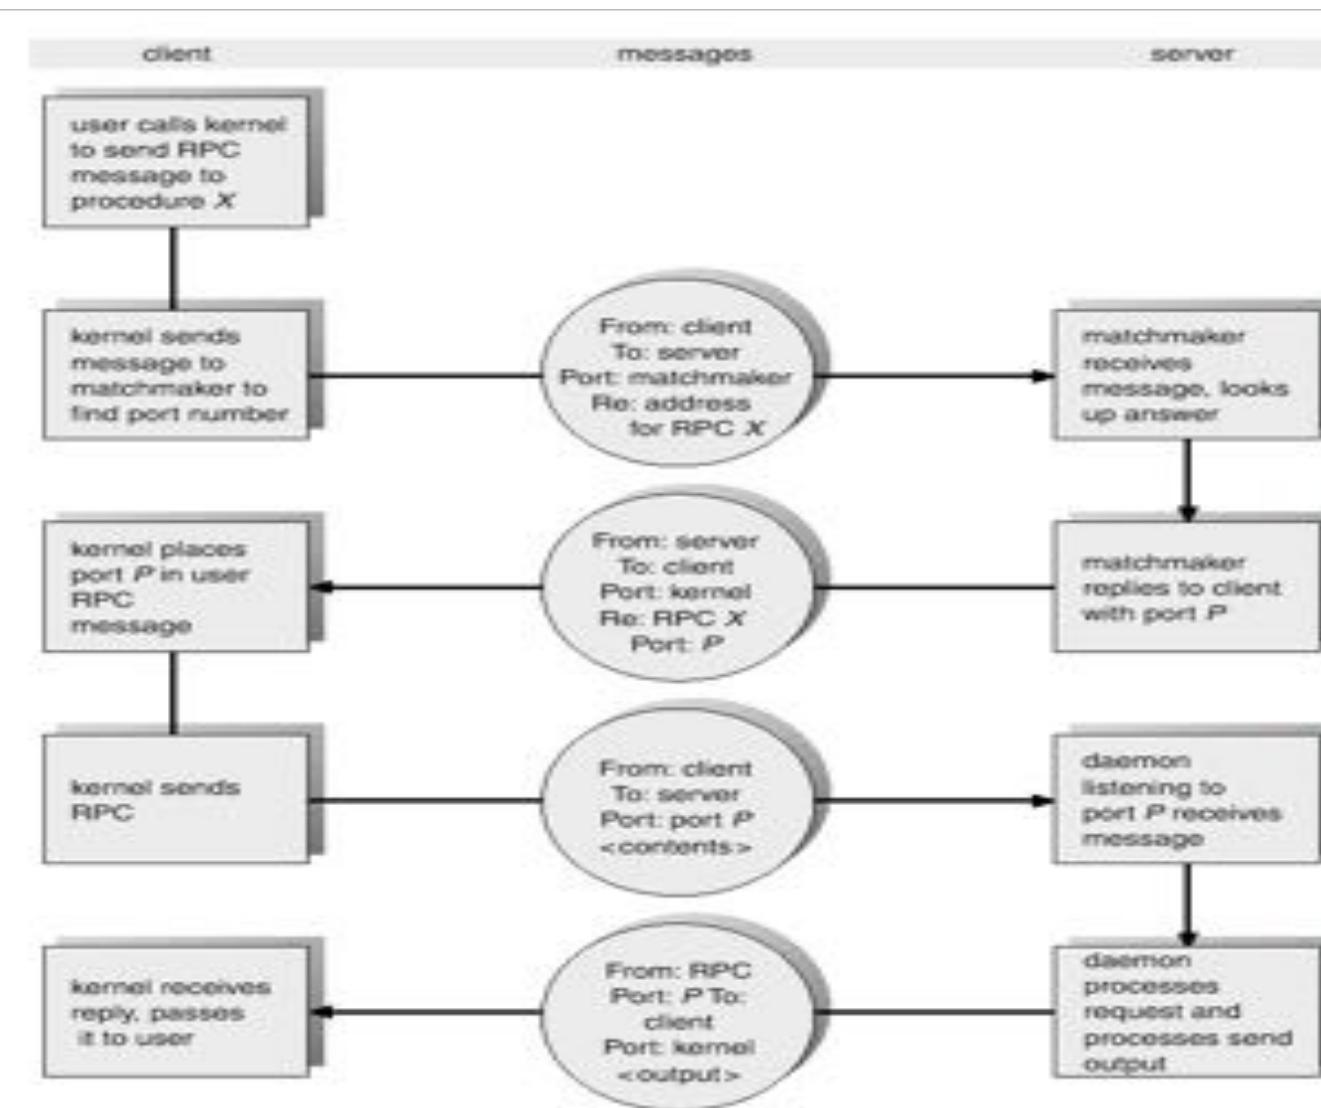
Issues using Sockets

- Communication semantics
 - Reliable or not
- Naming
 - How do we know a machine's IP address? DNS
 - How do we know a service's port number?
- Protection
 - Which ports can a process use?
 - Who should you receive a message from?
 - ◆ Services are often open -- listen for any connection
- Performance
 - How many copies are necessary?
 - Data must be converted between various data types

Remote Procedure Calls (RPC)

- Procedure calls between processes on network
 - Looks like a “normal” procedure call
 - However, the called procedure is run by another process
- RPC mechanism
 - Client stub (proxy for actual procedure call on server)
 - Client stub “marshalls” arguments
 - Client stub find destination (server) for RPC
 - Client stub send call and marshalled arguments to server
 - Server stub (performs the actual procedure call)
 - Server stub “unmarshalls” arguments
 - Server stub calls procedure with arguments
 - Server stub “marshall” result and returns to client

Remote Procedure Calls

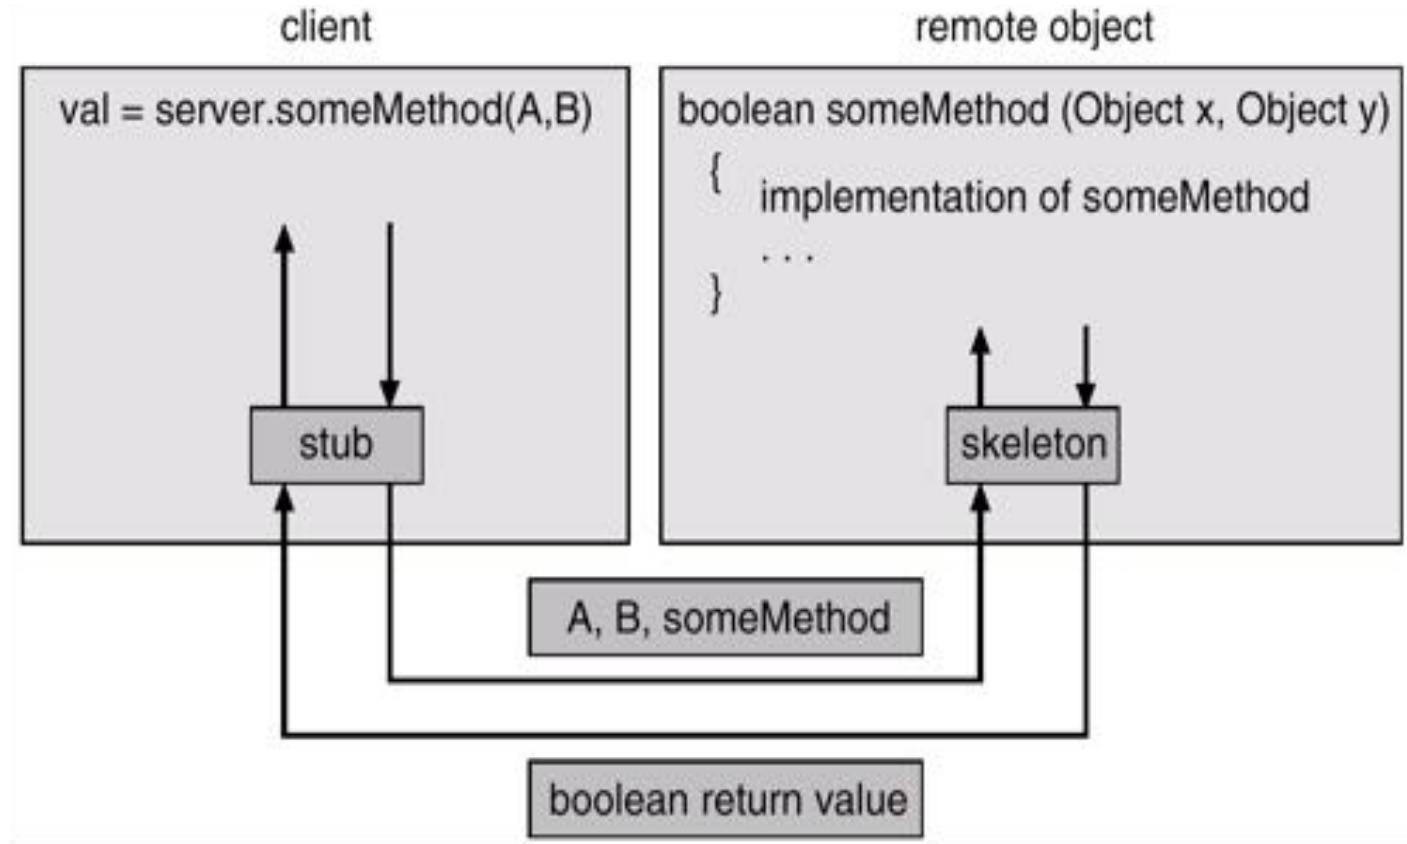


Remote Procedure Calls

- Supported by systems
 - CORBA
 - Java RMI
- Issues
 - Support to build client/server stubs
 - Marshalling arguments and code
 - Layer on existing mechanism (e.g., sockets)
 - Remote server crashes ... then what?
- Performance versus abstractions
 - What if the two processes are on the same machine?

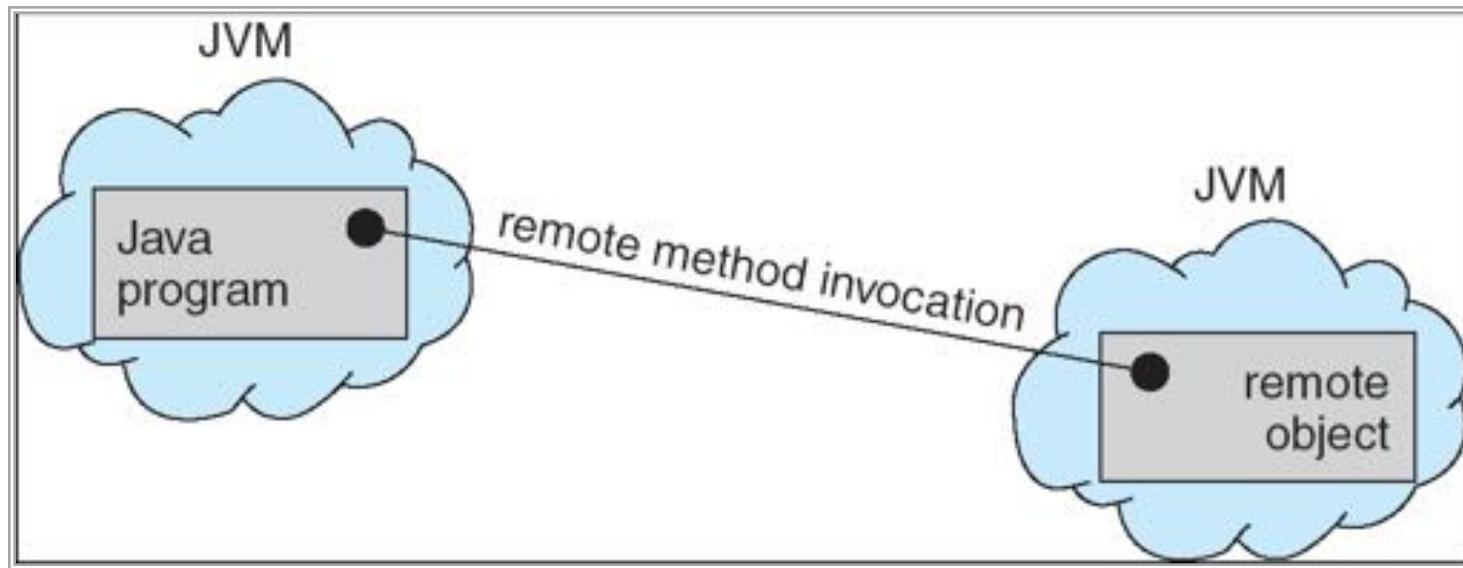
Remote Procedure Calls

□ Marshalling



Remote Method Invocation (RMI)

- RMI is a Java mechanism similar to RPCs
- RMI allows a Java program on one machine to invoke a method on a remote object



IPC Summary

- Lots of mechanisms
 - Pipes
 - Shared memory
 - Sockets
 - RPC
- Trade-offs
 - Ease of use, functionality, flexibility, performance
- Implementation must maximize these
 - Minimize copies (performance)
 - Synchronous vs Asynchronous (ease of use, flexibility)
 - Local vs Remote (functionality)

Summary

- Process
 - Execution state of a program
- Process Creation
 - fork and exec
 - From binary representation
- Process Description
 - Necessary to manage resources and context switch
- Process Scheduling
 - Process states and transitions among them
- Interprocess Communication
 - Ways for processes to interact (other than normal files)

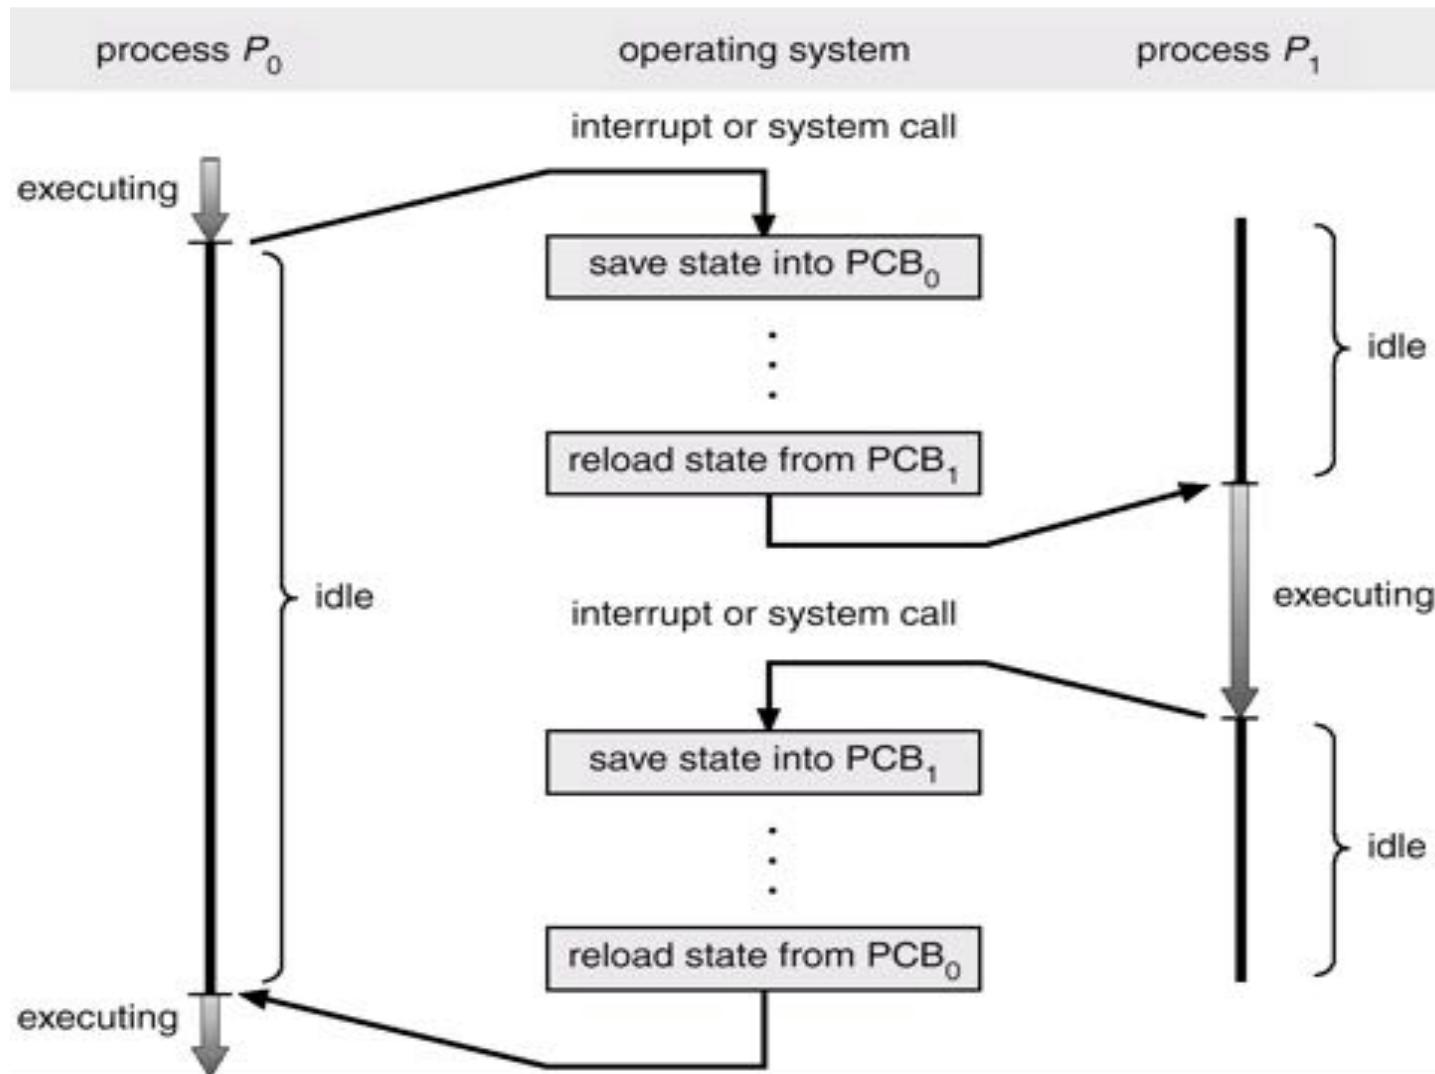
Next Class

□ Threads



Additional Slides

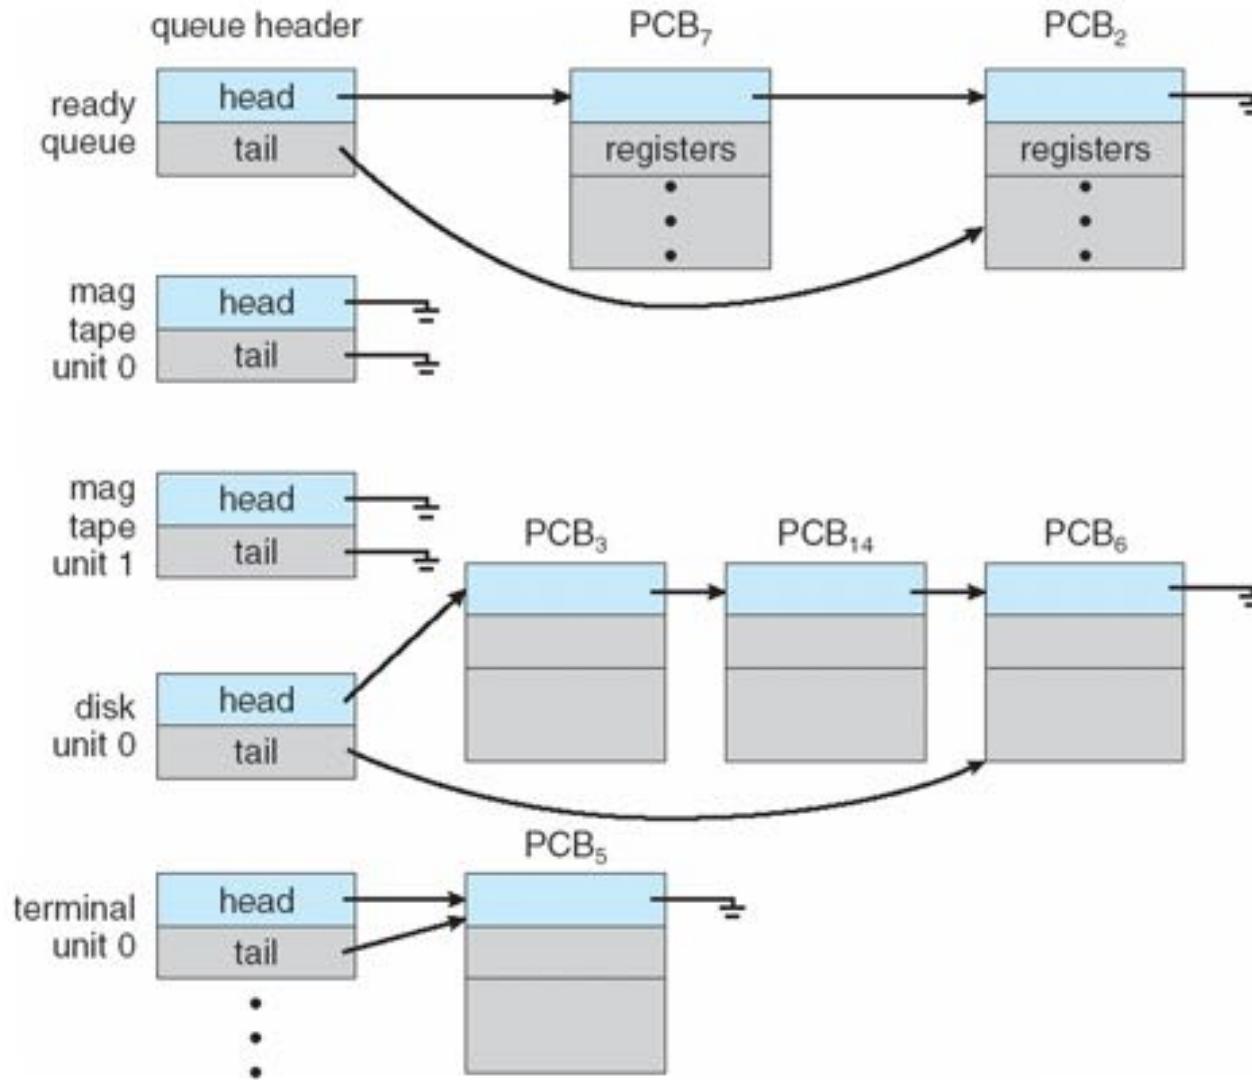
Context Switch



Process Scheduling

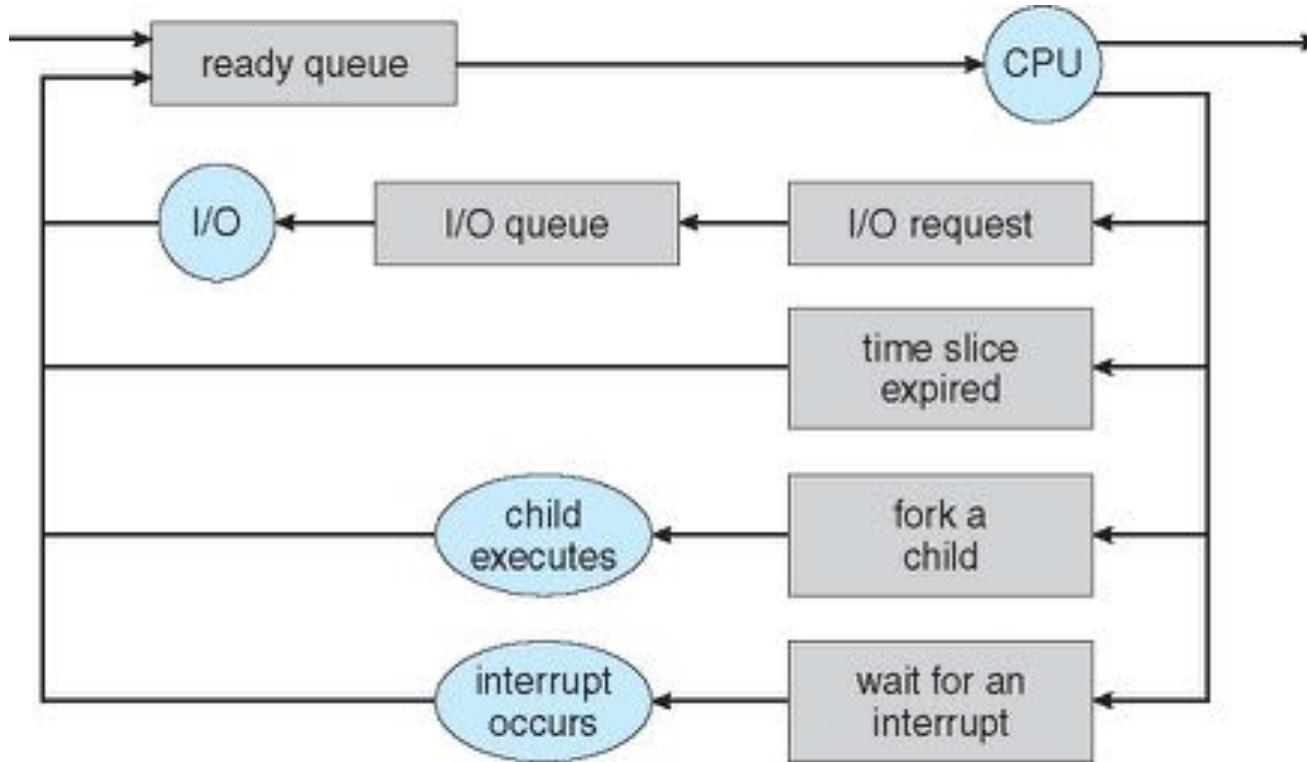
- Maximize CPU use, quickly switch processes onto CPU for time sharing
- Process scheduler selects among available processes for next execution on CPU
- Maintains scheduling queues of processes
 - *Job queue* – set of all processes in the system
 - *Ready queue* – set of all processes residing in main memory, ready and waiting to execute
 - *Device queues* – set of processes waiting for an I/O device
 - Processes migrate among the various queues

Ready Queue And Various I/O Device Queues



Representation of Process Scheduling

- Process scheduling queueing diagram represents:
 - queues, resources, flows
- Processes move through the queues



Schedulers

- *Short-term scheduler* (or CPU scheduler)
 - Selects which process should be executed next and allocates CPU
 - Sometimes the only scheduler in a system
 - Short-term scheduler is invoked frequently (milliseconds)
 - ◆ must be fast because involves context switching
- *Long-term scheduler* (or job scheduler)
 - Selects which programs should be launched
 - ◆ program is not yet running as a process
 - ◆ once a program is launched, its process is brought into the ready queue
 - Long-term scheduler is invoked infrequently (seconds, minutes)
 - ◆ does not have to be fast
 - The long-term scheduler controls the degree of multiprogramming
 - It is possible that a process can be removed temporarily

Process Mix

- Processes can be described as either:
 - I/O-bound process
 - ◆ spends more time doing I/O than computations
 - ◆ many short CPU bursts
 - CPU-bound process
 - ◆ spends more time doing computations
 - ◆ few very long CPU bursts
- There can be sub-categories of these
- Long-term scheduler strives for good process mix

Addition of Medium Term Scheduling

- Medium-term scheduler can be added if degree of multiple programming needs to decrease
 - Remove process from memory, store on disk, bring back in from disk to continue execution
 - Refer to this as job swapping

