



CIS 415

Operating Systems

Threads

Prof. Allen D. Malony

Department of Computer and Information Science
Spring 2020



UNIVERSITY OF OREGON

Logistics

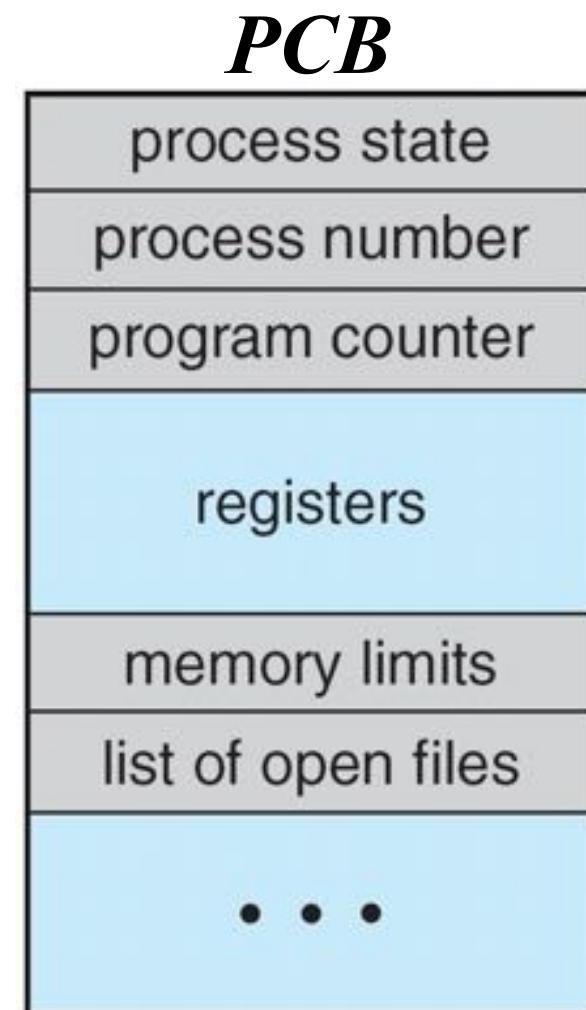
- Keep working hard on Project 1
 - Due Sunday, April 19, 11:59pm
- Lab this week covers systems call
- Read Chapter 4

Outline

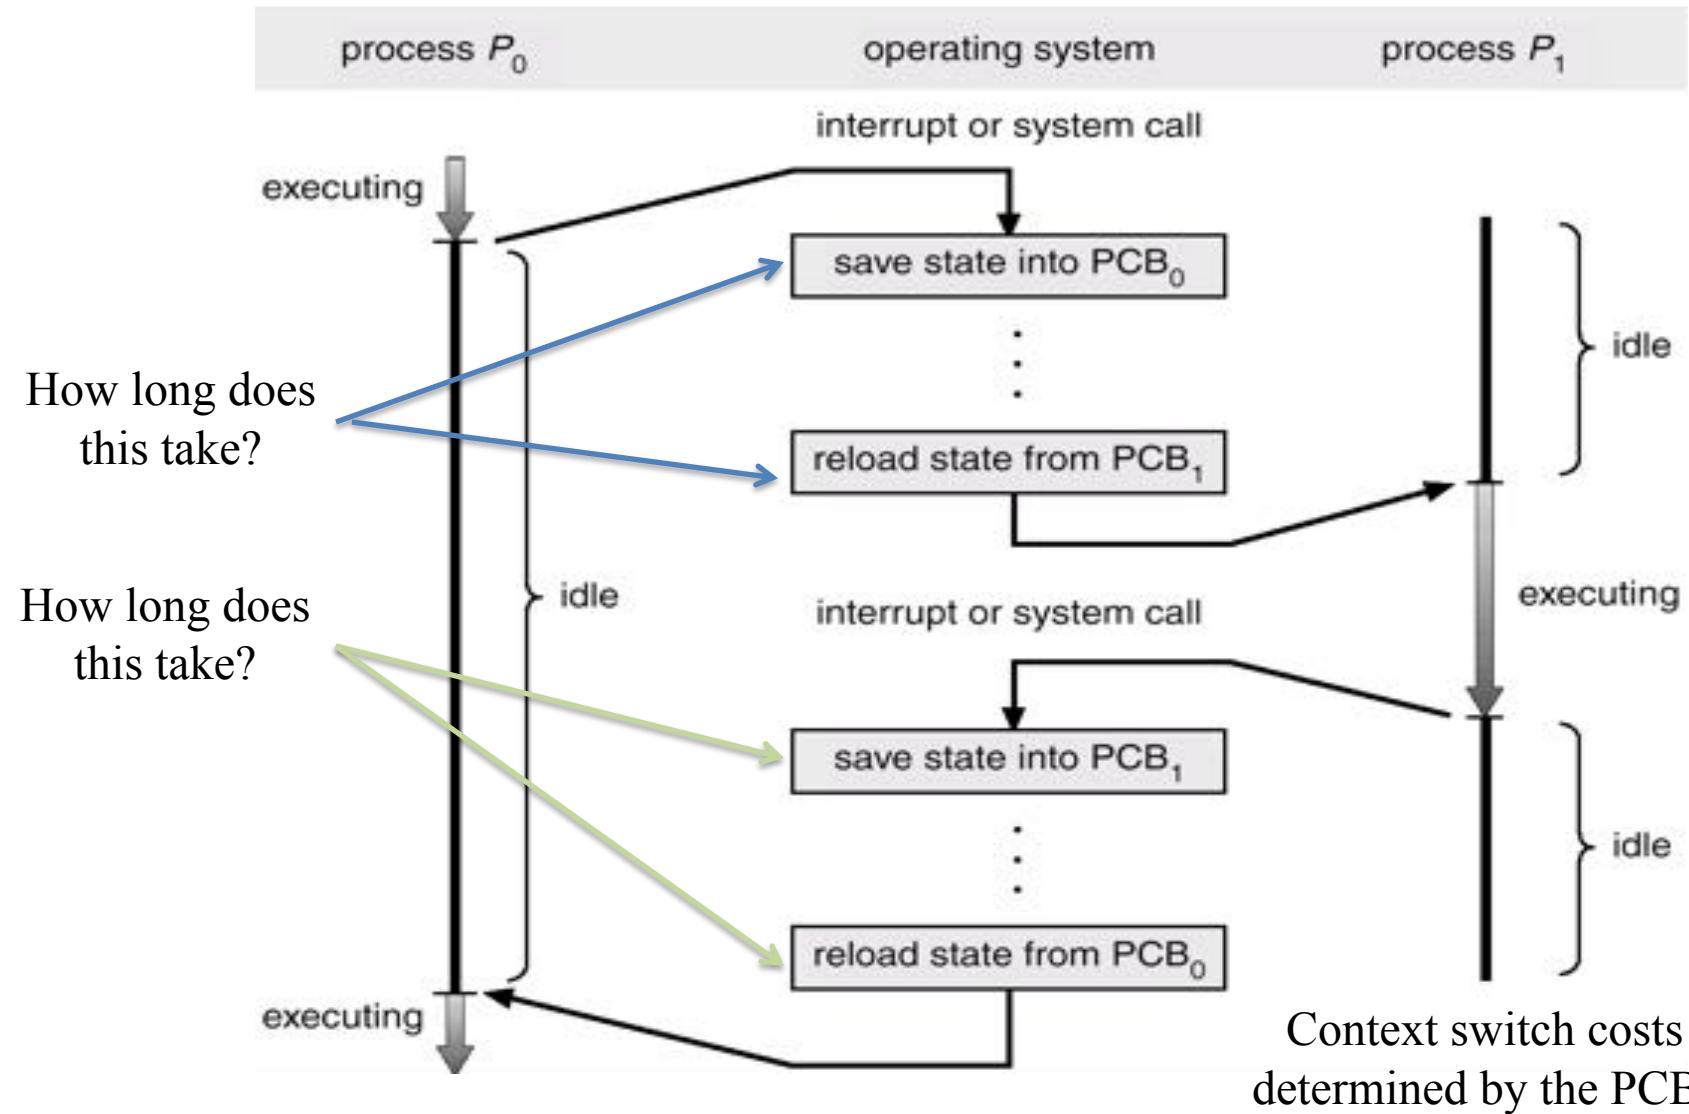
- Processes revisited
- Threads, threads, threads,

Process Control Block

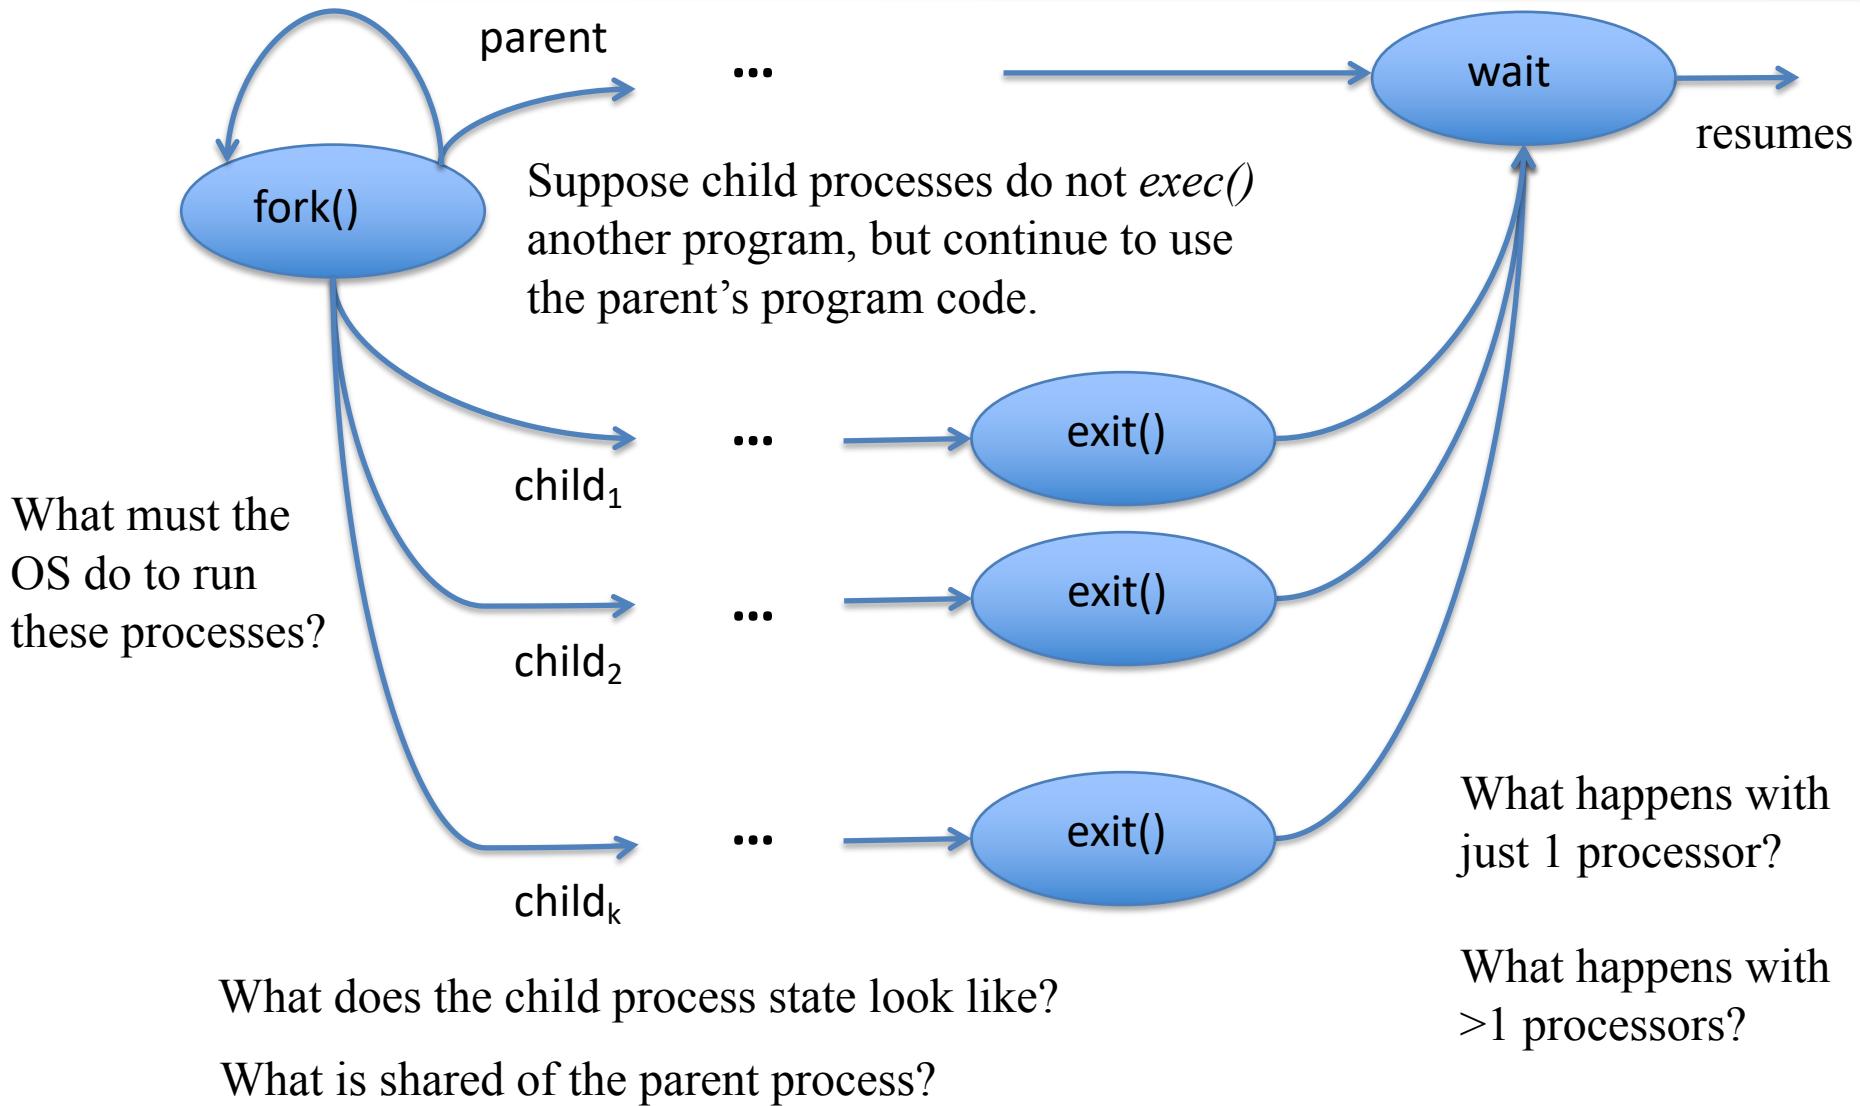
- Information associated with each process (aka *task control block*)
- Process state
 - Running, waiting, ...
- Program counter
 - Location of instruction to next execute
- CPU registers
 - Contents of all process-centric registers
- CPU scheduling information
 - Priorities, scheduling queue pointers, ...
- Memory-management information
 - Memory allocated to the process
 - Memory limits and other associated information
- Accounting information
 - CPU used, elapsed time, time limits, ...
- I/O status information
 - I/O devices allocated to process, list of open files



Process Context Switch



Program with Multiple Processes



Program Creation System Calls

- *fork()*
 - Copy address space of parent and all threads
- *vfork()*
 - Do not copy the parent's address space
 - Share address space between parent and child
 - While parent blocks, child must exit or call exec
- *exec()*
 - Load new program and replace address space
 - Some resources may be transferred (open file descriptors)
 - Specified by arguments
- *clone()*
 - Like fork() but child shares some process context
 - More explicit control over what is shared
 - Process address space can be shared!
 - Calls a function pointed to by argument

The child is still
considered a process,
but lighterweight!

Process Model

- Much of the OS's job is keeping processes from interfering with each other ... Why?
- Each process has its OWN resources to use
 - Program code to execute, address space, files, ...
- Processes are good for isolation (protection)
 - Prevent one process from affecting another process
- Processes are *heavyweight* ... Why do you think?
 - Pay a price for isolation
 - A full “process swap” is required for multiprocessing
 - There is lots of process state to save and restore
 - OS must context switch between them
 - ◆ intervene to save/restore all process state
- Is there an alternative?

Why Threads?

- A process is “a program in execution”
 - Memory address space containing code and data
 - Other resources (e.g., open file descriptors)
 - State information (PC, register, SP) => PCB details
- Consider a process in 2 respects (categories)
 - 1) Collection of resources required for execution
 - ◆ code, address space, open files, ...
 - 2) A “thread of execution”
 - ◆ current state of execution (CPU state)
 - ◆ where the execution is now (PC)
- Suppose we think about these separately!
 - Resources
 - Execution state



Terminology

□ *Multiprogramming*

- Running multiple programs on a computer concurrently
- Each program is a (set of) process(es)
- Processes of different programs are independent

□ *Multiprocessing*

- Running multiple processes on a computer concurrently
- OS manages mapping of processes to processor(s)

□ *Multithreading*

- Define multiple *execution contexts (threads of execution)* in a single address space (of a process)
- OS manages mapping of threads to an address space
- OS manages mapping of threads to processor(s)

What's a Thread?

- Thread of execution through a program on a CPU
 - Program counter *per thread*
 - Registers *per thread*
- Memory
 - Address space *process*
 - ◆ address space is shared!!!
 - Stack *per thread*
 - ◆ each thread has its own stack pointer
 - Heap *process*
 - + private dynamic memory *per thread*
- I/O
 - Share files, sockets, ... *process*

Why Multithreaded Applications?

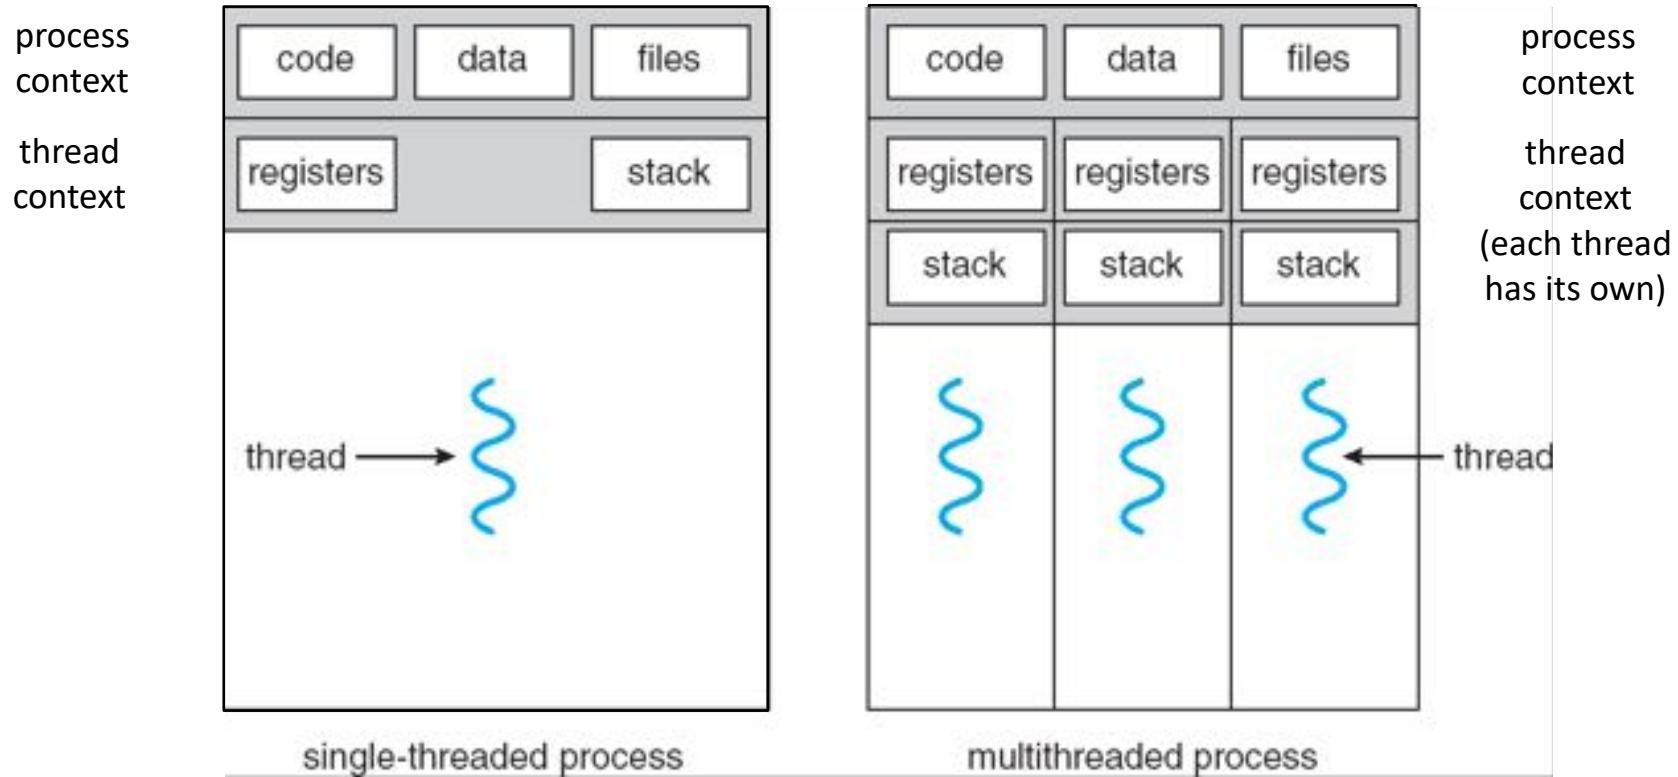
- Multiple threads sharing a common address space
 - More correctly, they share process resources, including memory!
- Why would you want to do that?
- Some applications could be written to support concurrency
 - How do you get concurrency?
 - One way is to create multiple processes
 - Use IPC to support process-level concurrency
- Some applications want to share data structures among concurrently executing parts of the computation
 - Is this possible with processes? => shared segments
 - Is it difficult with processes? => well, it is not that easy
 - Again, use IPC for process-level data sharing
- What is the problem? What is the solution?

Advantages of Threads

- Threads could be used if there is no need for the OS to enforce resource separation
 - This is a trust issue, in part (back off on isolation / protection)
 - However, introduces more issues with respect to concurrency
- Improve responsiveness
 - Possible to have a thread of execution that never blocks
 - More to come about this ...
- Resource sharing is facilitated
 - All threads in a process have equal access to resources
- Economy of resources
 - Thread-level resources are “cheaper” than process resources!
 - Threads are “lighter weight”
- Utilization of multiprocessors
 - Run multiple threads on multiple processes without the overhead of running multiple processes

Single-Threaded vs. Multithreaded

- Regular UNIX process can be thought of as a special case of a multithreaded process
 - A process that contains just one thread!
- Multithreaded process has multiple threads



Working with Threads

- In a C program
 - *main()* procedure defines the first thread
 - C programs always start at *main()*
- Now you want to create a second thread
 - Allocate resources to maintain a second execution context in same address space
 - ◆ think about what state will be needed for a thread
 - Want something similar to *fork()* but simpler
 - ◆ supply a procedure name when start the new thread
 - Remember this creates another thread of execution

Threads vs. Processes

- Easier to create than a new process
- Less time to terminate a thread than a process
- Less time to switch between two threads
 - Within the same process
- Less communication overheads
 - Communicating between the threads of one process is simple because the threads share everything
 - Address space is shared ...
 - ... thus memory is shared (Hmm ...)



Which is Cheaper?

- Create new process or create new thread (in existing process)?
- Context switch between processes or threads?
- Interprocess or interthread communication?
- Sharing memory between processes or threads?
- Terminating a process or terminating a thread (not the last one)?

Process creation method	Time (sec), elapsed (real)
<i>fork()</i>	22.27 (7.99)
<i>vfork() (faster fork)</i>	3.52 (2.49)
<i>clone()</i>	2.97 (2.14)

Time to create 100,000 processes (Linux 2.6 kernel, x86-32 system)
clone() creates a lightweight Linux process (thread)

Implications?

- Consider a web server on a Linux platform
- Measure 0.22 ms per *fork()*
 - Maximum of $(1000 / 0.22) = 4545.5$ connections/sec
 - 0.45 billion connections per day per machine
 - ◆ fine for most servers
 - ◆ too slow for a few super-high-traffic front-line web services
- Facebook serves $O(750 \text{ billion})$ page views per day
 - Guess ~1-20 HTTP connections per page
 - Would need 3,000 -- 60,000 machines just to handle *fork()*, without doing any work for each connection!
- What is the problem here?

Thread Attributes

- Global to process:
 - memory
 - PID, PPID, GID, SID
 - controlling term
 - process credentials
 - record locks
 - FS information
 - timers
 - resource limits
 - and more...
- Local to specific thread:
 - thread ID
 - stack
 - signal mask
 - thread-specific data
 - alternate signal stack
 - error return value
 - scheduling policy/priority
 - Linux-specific
 - (e.g., CPU affinity)

Threading Models

- *Programming* – library or system call interface
 - Kernel threading (most common)
 - ◆ thread management support in the kernel
 - ◆ invoked via system call
 - ◆ NOTE: CPU only runs kernel threads!!!
 - User-space threading
 - ◆ thread management support in user-space library
 - ◆ threads are “thread switched” by a user-level library
 - ◆ linked into your program
- *Scheduling* – application or kernel scheduling
 - May create user-level or kernel-level threads

Kernel Threads

- Thread management support in kernel
 - Sets of system calls for creating, invoking, and switching among threads
- Supported and managed directly by the OS
 - Thread objects in the kernel
- Nearly all OSes support a notion of threads
 - Linux -- thread and process abstractions are mixed
 - Solaris
 - Mac OS X
 - Windows XP
 - ...

User-space Threads

- Thread management support in user-space library
 - Sets of functions for creating, invoking, and switching among threads (all in user mode)
 - Need to switch threads stacks in user space ...
- Linked into your program
 - Thread libraries
- Examples
 - Qthreads (<http://www.cs.sandia.gov/qthreads/>)
 - GNU Pth (<http://www.gnu.org/software/pth/>)
 - Cilk (<https://en.wikipedia.org/wiki/Cilk>)

Implementing User-space Threading

- Threads can perform operations in user mode that are usually handled by the OS
 - Assumes cooperating threads so hardware enforcement of separation not required
- Idea:
 - Think of a “dispatcher” subroutine in the process that is called when a thread is ready to relinquish control to another thread
 - Manages stack pointer, program counter
 - Switches process’s internal state among threads

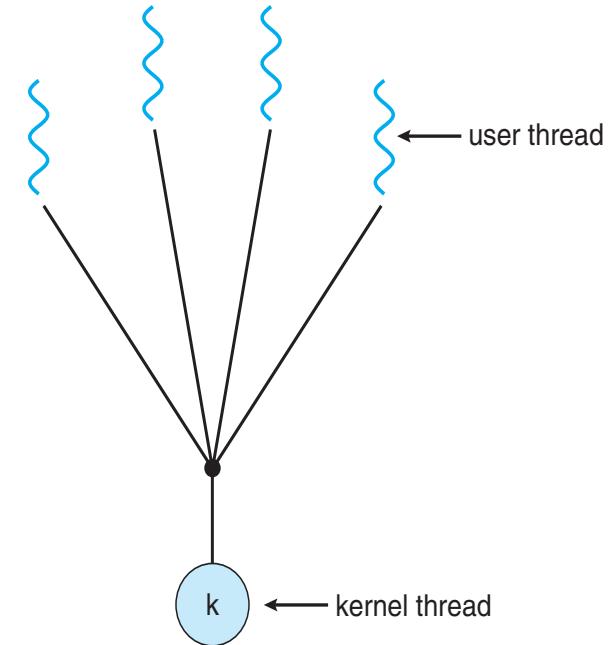
Many-to-One Thread Model

- Many user-level threads correspond to a single kernel thread

- Kernel is not aware of the mapping
 - Handled by a thread library

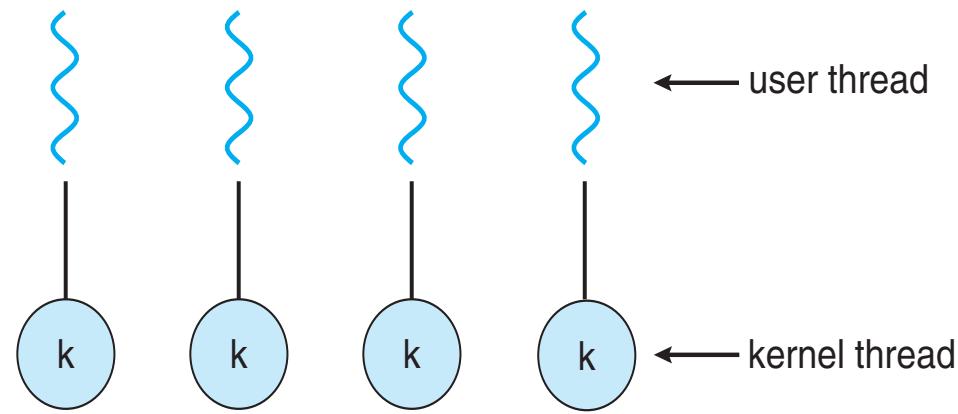
- How does it work?

- Create and execute a new thread
 - Upon yield, switch to another user thread in the same process
 - ◆ kernel is unaware
 - Upon wait, all threads are blocked
 - ◆ kernel is unaware there are other options
 - ◆ can not wait and run at the same time



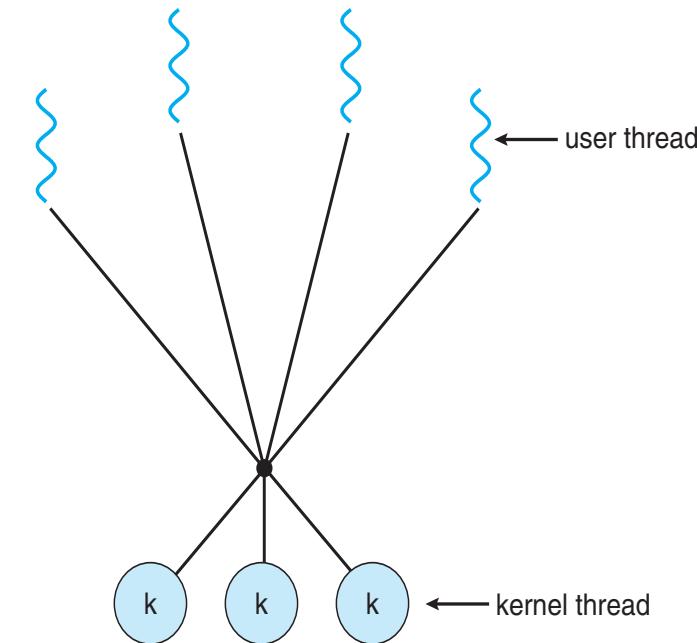
One-to-One Thread Model

- One user-level thread per kernel thread
 - A kernel thread is allocated for every user-level thread
 - Must get the kernel to allocate resources for each new user-level thread
- How does it work?
 - Create new thread
 - ◆ system call to kernel
 - Upon yield, switch to another kernel thread in system
 - ◆ kernel is aware
 - Upon wait, another thread in the process may run
 - ◆ only the single kernel thread is blocked
 - ◆ kernel is aware there are other options in this process



Many-to-Many Thread Model

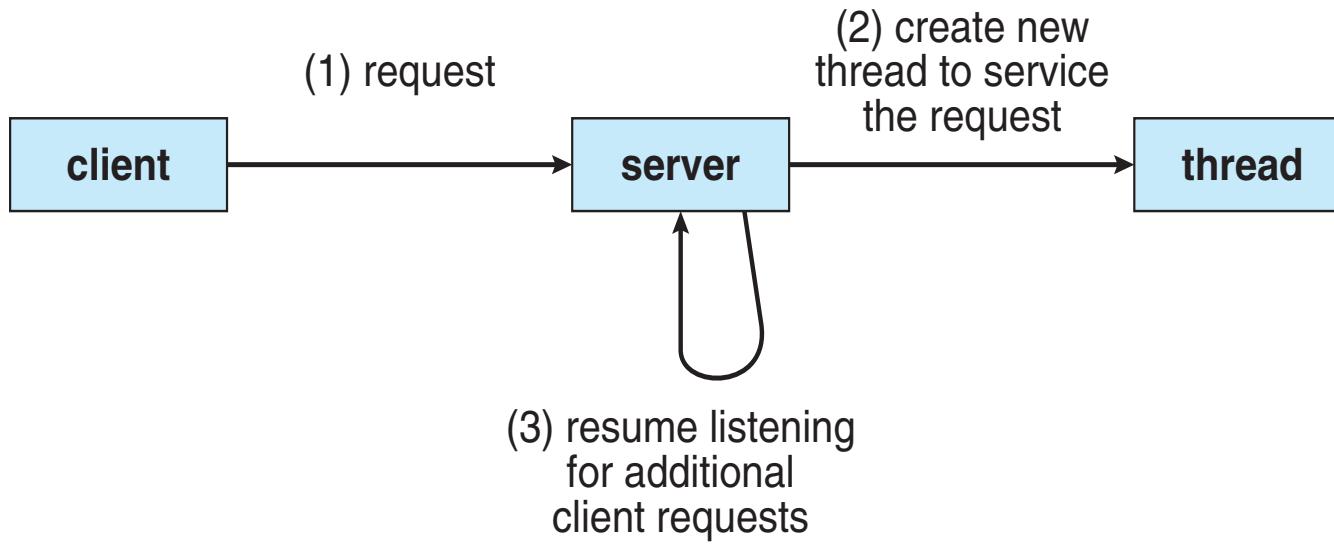
- A pool of user-level threads maps to a pool of kernel threads
 - Pool sizes can be different (kernel pool is no larger)
 - A kernel thread in the pool is allocated for every user-level thread
 - No need for the kernel to allocate resources for each new user-level thread
- How does it work?
 - Create new thread
 - ◆ may map to kernel thread dynamically
 - Upon yield, switch to another thread
 - ◆ kernel is aware
 - Upon wait, another thread in the process may run
 - ◆ if a kernel thread is available to be scheduled to that process
 - ◆ kernel is aware of the mapping between user and kernel threads



Problems Solved with Threads

- Imagine you are building a web server
 - You could allocate a pool of threads, one for each client
 - ◆ thread would wait for a request, get content file, return it
 - How would the different thread models impact this?
- Imagine you are building a web browser
 - You could allocate a pool of threads
 - ◆ some for user interface
 - ◆ some for retrieving content
 - ◆ some for rendering content
 - What happens if the user decided to stop the request?
 - ◆ mouse click on the stop button

Multithreaded Server Architecture



Linux Threads

- Linux uses a one-to-one thread model
 - Threads are called *tasks*
- Linux views threads as “contexts of execution”
 - Threads are defined separately from processes
 - There is flexibility in what is private and shared
- Linux system call
 - *clone(int (*fn)(), void **stack, int flags, int argc, ...)*
 - Create a new thread (Linux task)
- May be created in the same address space or not
 - Flags (on means “share”): clone VM, clone filesystem, clone files, clone signal handlers
 - If all these flags off, what system call is clone equal to?

POSIX Threads

- POSIX Threads is a thread API specification
 - Does not define directly the implementation!
 - Could be implemented differently
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
 - Specification, not implementation
 - Provided either as user-level or kernel-level (very interesting)
 - API specifies behavior of the thread library
 - Implementation is up to development of the library
- Common in UNIX operating systems
 - Solaris, Linux, Mac OS X
- POSIX Threads is also known as *Pthreads*



POSIX Threads

- *pthread_create()*
 - start the thread
- *pthread_self()*
 - return thread ID
- *pthread_equal()*
 - for comparisons of thread ID's
- *pthread_exit()*
 - or just return from the start function
- *thread_join()*
 - wait for another thread to terminate & retrieve value from *pthread_exit()*
- *pthread_cancel()*
 - terminate a thread, by TID
- *pthread_detach()*
 - thread is immune to join or cancel & runs independently until it terminates
- *pthread_attr_init()*
 - thread attribute modifiers

POSIX Threads FAQ

- How to pass multiple arguments to start a thread?
 - Build a struct and pass a pointer to it
- Is the pthreads ID unique to the system?
 - No, just process – Linux task ids are system-wide
- After *pthread_create()*, which thread is running?
 - It acts like fork in that both threads are running
- How many threads terminate when ...
 - *exit()* is called? – all in the process
 - *pthread_exit()* is called? – only the calling thread
- How are variables shared by threads?
 - Globals, local static, dynamic data (heap)

Pthreads Example

```
#include <pthread.h>
#include <stdio.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    if (argc != 2) {
        fprintf(stderr,"usage: a.out <integer value>\n");
        return -1;
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr,"%d must be >= 0\n",atoi(argv[1]));
        return -1;
    }
}
```

Pthreads Example (Cont.)

```
/* get the default attributes */
pthread_attr_init(&attr);
/* create the thread */
pthread_create(&tid,&attr,runner,argv[1]);
/* wait for the thread to exit */
pthread_join(tid,NULL);

printf("sum = %d\n",sum);
}

/* The thread will begin control in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```

Pthreads Code for Joining 10 Threads

```
#define NUM_THREADS 10

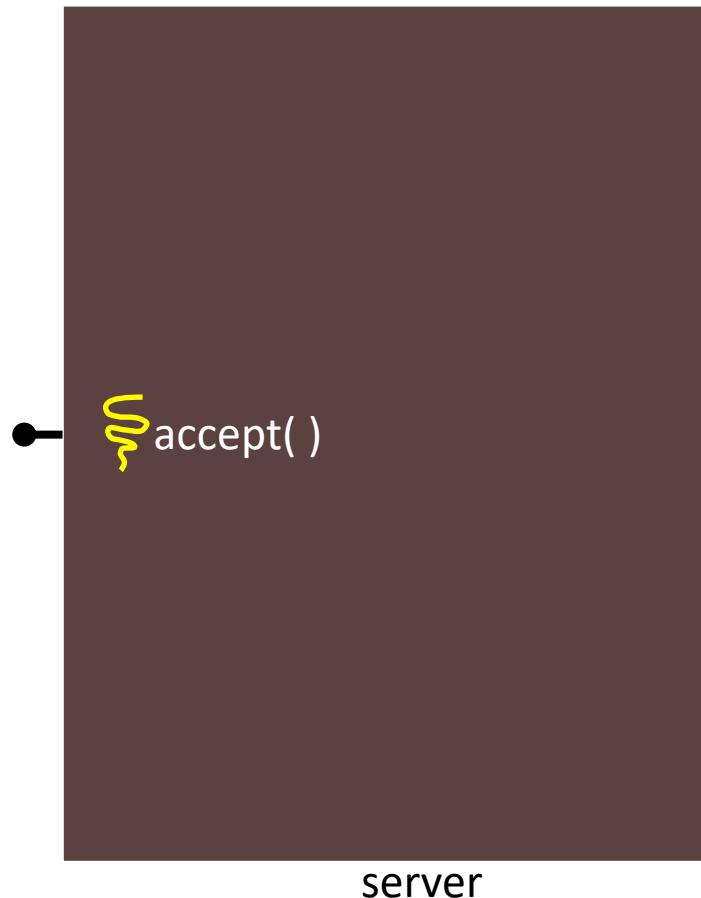
/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```

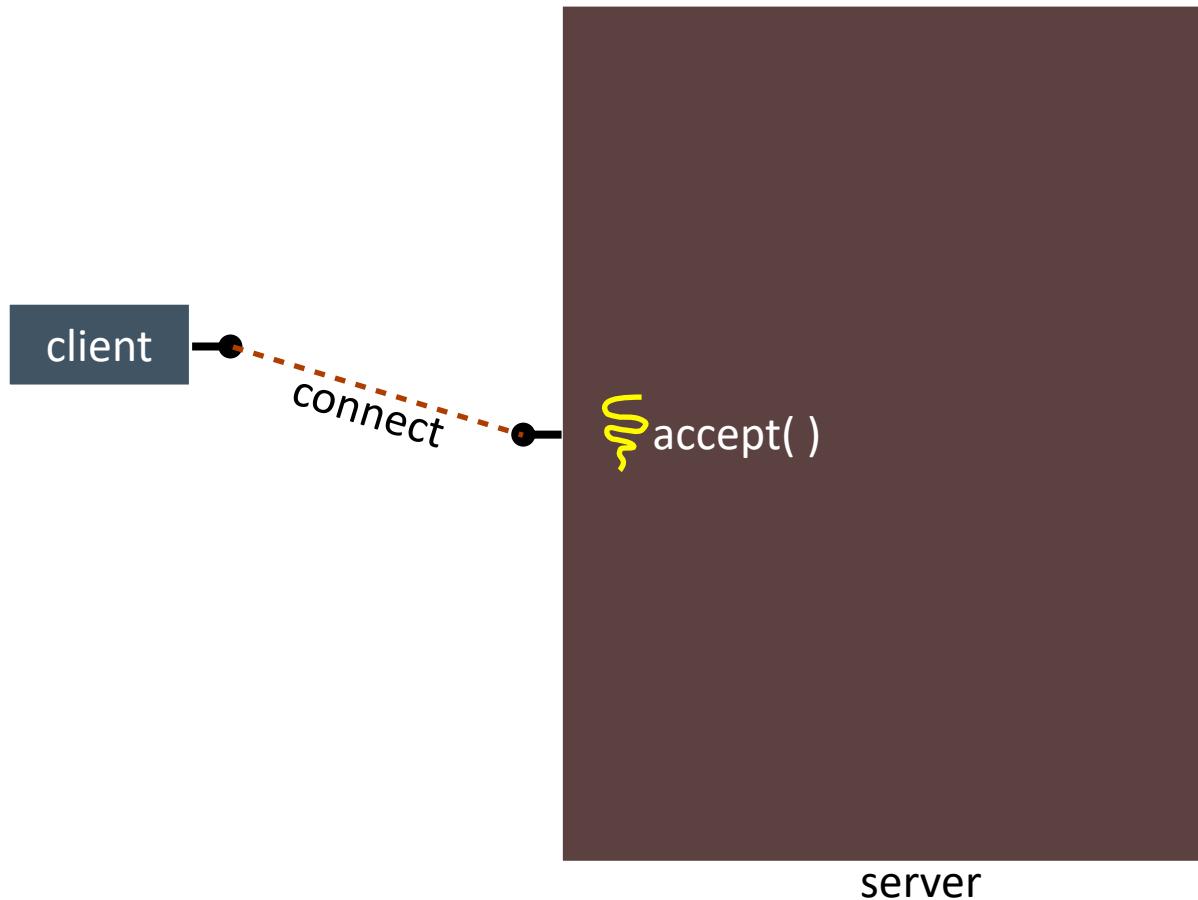
Concurrency with Threads

- Consider the client-server example again
- Now want to run with just a single process
 - Need to use threads to get concurrency
- Process “main” (parent) thread waits for clients
 - Parent thread forks (or dispatches) a new thread to handle each new client connection
 - Responsibilities of the child thread:
 - ◆ handles the new connection
 - ◆ exits when the connection terminates

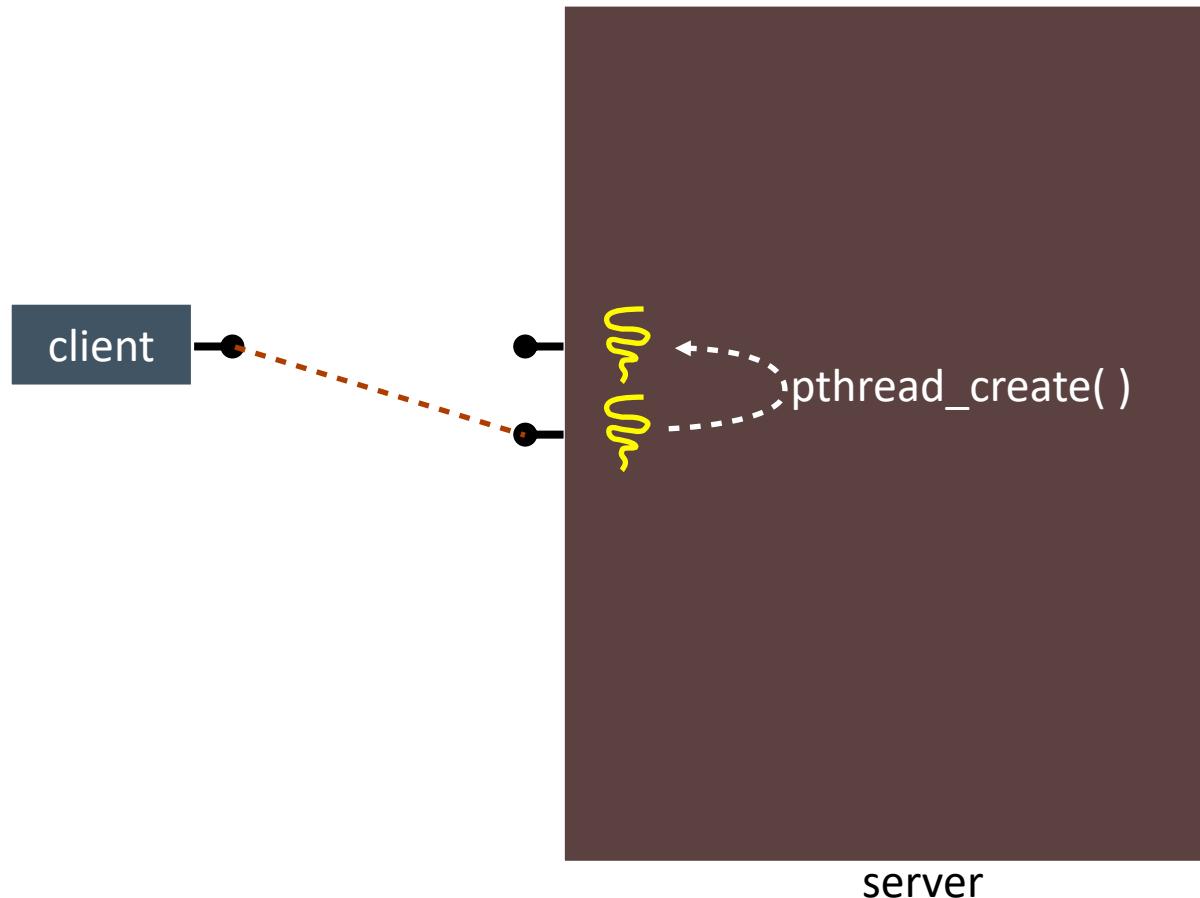
Client-Server with Pthreads (1)



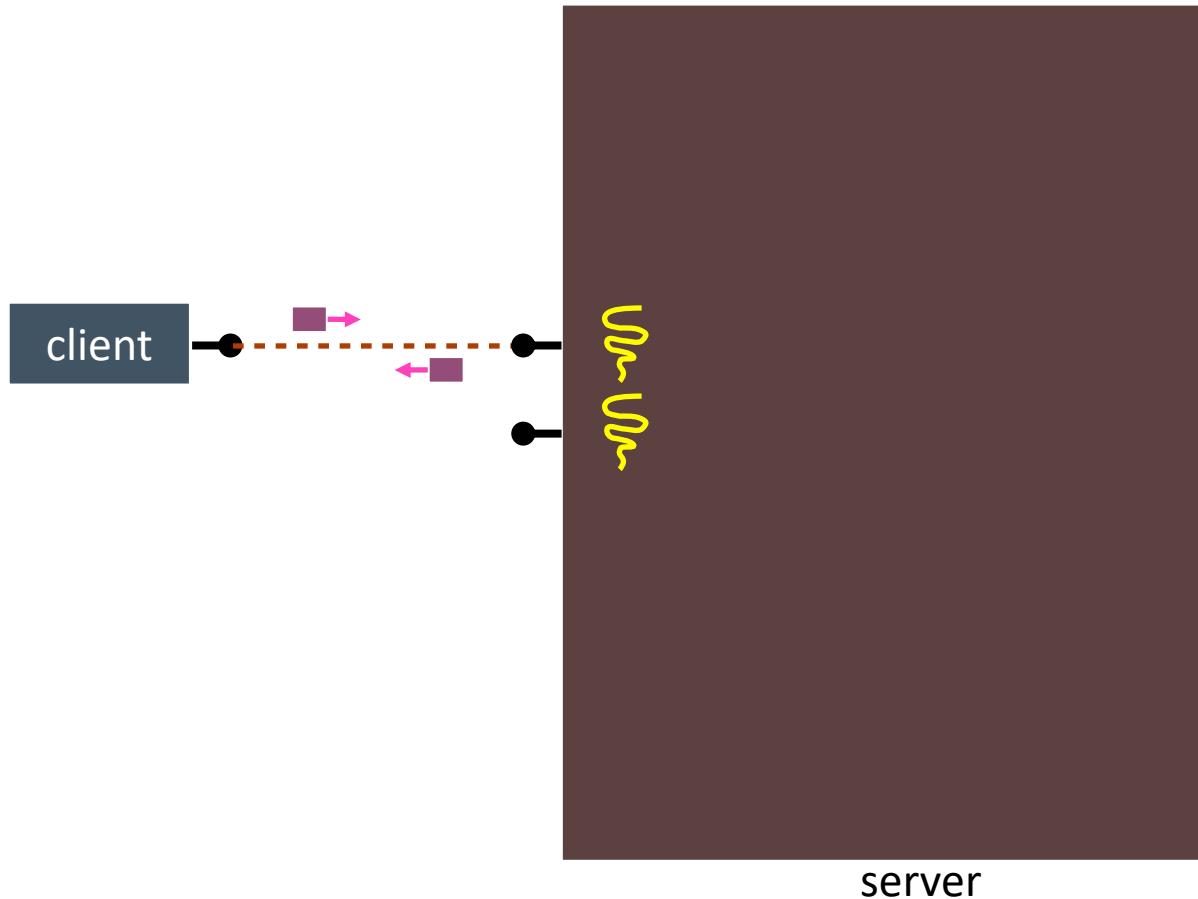
Client-Server with Pthreads (2)



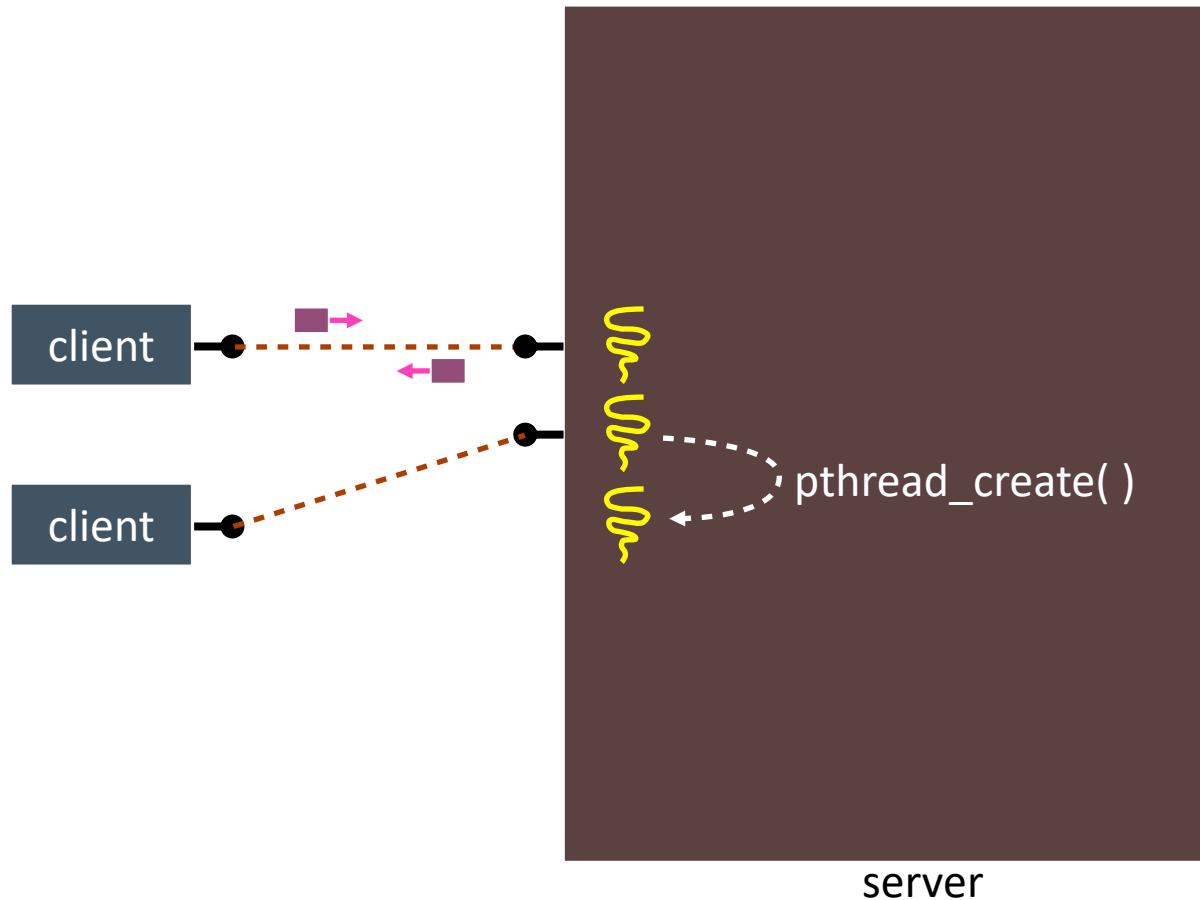
Client-Server with Pthreads (3)



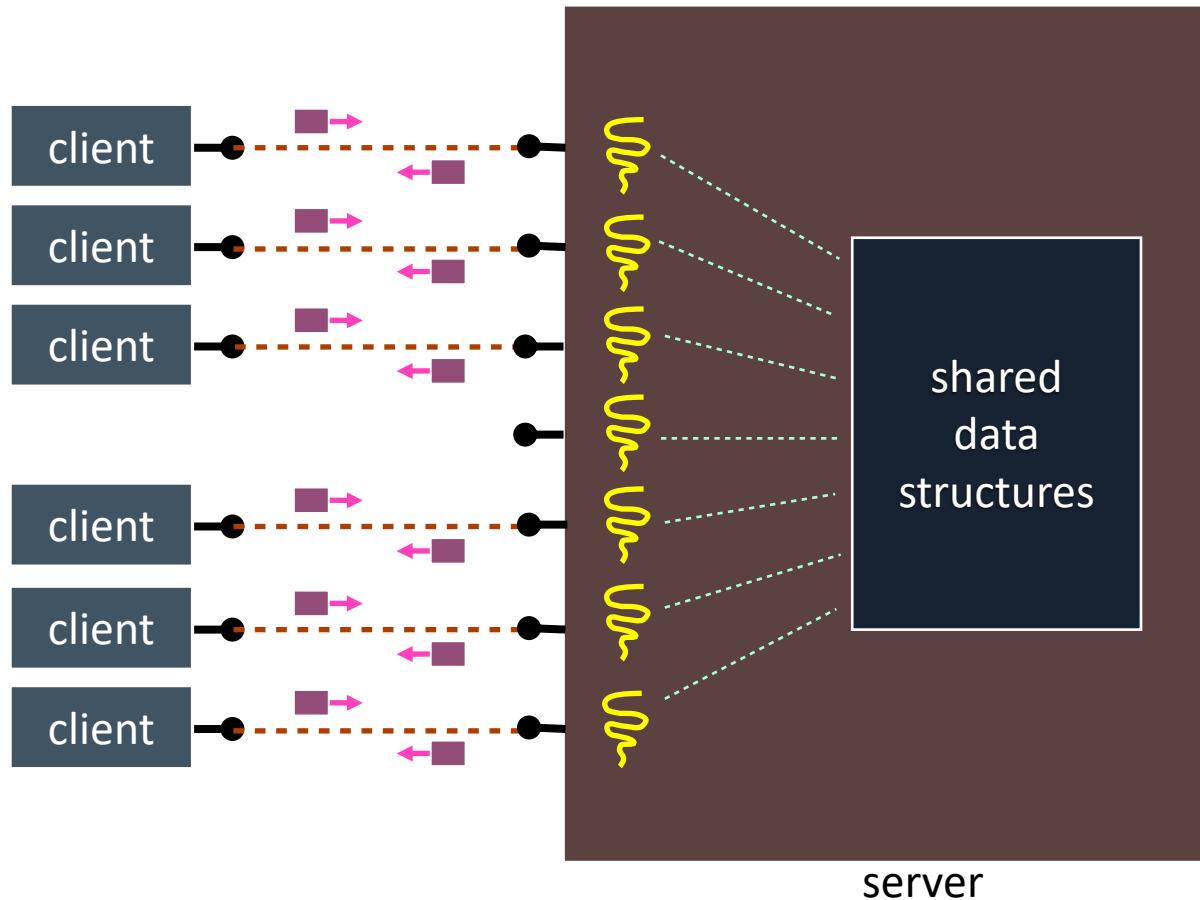
Client-Server with Pthreads (4)



Client-Server with Pthreads (5)



Client-Server with Pthreads (5)



Implications?

- Consider a web server on a Linux platform
- 0.0297 ms per thread create (time for *clone()*)
 - 10x faster than process forking
 - Maximum of $(1000 / 0.0297) = \sim 33,670$ connections/sec
 - 3 billion connections per day per machine
 - ◆ much, much better
- Facebook would need only 500 machines
- So, why do we need processes at all?
- Just write everything using threads
 - Writing safe multithreaded code can be complicated
 - Why? What are the issues?

Concurrent Threads

□ Benefits

- All threads are running the same code
 - ◆ still the case that much of the code is identical!
- Shared-memory communication is possible
- Good CPU and network utilization
 - ◆ lower overhead than processes

□ Disadvantages

- Synchronization is complicated
- Shared fate within a process
 - ◆ one rogue thread can hurt you badly

□ More to come on this topic ...

Inter-Thread Communication

- Can you use shared memory?
 - Sure, it is there, might as well use it
 - Just need to allocate memory in the address space
 - ◆ No need for fancy IPC shared memory
- Can you use message passing?
 - Of course
 - Would have to build infrastructure though
- Hmm, can threads utilize IPC mechanisms
 - That is a good question actually
 - Would need to make sure only 1 thread uses at a time

Fork/Exec Issues

- Semantics are ambiguous for multithreaded processes
- *fork()*
 - How does it interact with threads?
- *exec()*
 - What happens to the other threads?
- *fork*, then *exec*
 - Should all threads be copied?



Thread Cancellation

- So, you want to stop a thread from executing
 - Do not need it anymore
 - It is just hanging around and you want to get rid of it
- Two choices
 - Synchronous cancellation
 - ◆ wait for the thread to reach a point where cancellation is permitted
 - ◆ no such operation in Pthreads, but can create your own
 - Asynchronous cancellation
 - ◆ terminate it now
 - ◆ *pthread_cancel(thread_id)*

Signal Handling

- What's a signal?
 - A form of IPC
 - Send a particular signal to another process
- Receiver's signal handler processes signal on receipt
- Example
 - Tell the Internet daemon (*inetd*) to reread its config file
 - Send signal to *inetd*: *kill -SIGHUP <pid>*
 - *inetd*'s signal handler for the SIGHUP signal re-reads the config file
- Note: some signals cannot be handled by the receiving process, so they cause default action (kill the process)



Signal Handling

- Synchronous signals
 - Generated by the kernel for the process
 - Due to an exception -- divide by 0
 - ◆ events caused by the thread receiving the signal
- Asynchronous signals
 - Generated by another process
- Asynchronous signals are more difficult for multithreading

Signal Handling and Threads

- So, you send a signal to a process
 - Which thread should it be delivered to?
- Choices
 - Thread to which the signal applies
 - Every thread in the process
 - Certain threads in the process
 - A specific signal receiving thread
- It depends...

Signal Handling and Threads

- UNIX signal model created decades before Pthreads
 - Conflicts arise as a result
- Synchronous vs. asynchronous cases
- Synchronous
 - Signal is delivered to the same process that caused the signal
 - Which thread(s) would you deliver the signal to?
- Asynchronous
 - Signal generated by another process
 - Which thread(s) in this case?

Thread Pools

- Pool of threads
 - Create (all) at initialization time
 - Assign task to a waiting thread
 - ◆ it is already made so it should be fast
 - Use all available threads
- What about when that task is done?
 - Suppose another request is in the queue ...
 - ... should we use running thread or another thread?
- Concern for the setup time cost
- Faster than setting up a process, but what is necessary?
 - How do we improve performance?



Scheduling

- How many kernel threads to create for a process?
 - In M:N model
- If last kernel thread for an application is to be blocked
 - What happens?
 - Recall the relationship between kernel and user threads
- It would be nice if the kernel told the application and the application had a way to get more kernel threads
 - Scheduler activation
 - ◆ at thread block, the kernel tells the application
 - Application can then get a new thread created
 - ◆ see lightweight threads

Re-entrance and Thread-Safety

- Terms that you might hear:
- *Reentrant* code
 - Code that can be run by multiple threads concurrently
- *Thread-safe* libraries
 - Library code that permits multiple threads to invoke the safe function
- Requirements
 - Rely only on input data
 - ◆ Or thread-specific data
 - Must be careful about locking (later)

Why not threads?

- Threads can interfere with one another
 - Impact of more threads on caches
 - Impact of more threads on TLB
 - Bug in one thread...
- Executing multiple threads may slow them down
 - Impact of single thread vs. switching among threads
- Harder to program a multithreaded program
 - Multitasking hides context switching
 - Multithreading introduces concurrency issues

Summary of Threads

- Threads
 - A mechanism to improve performance and CPU utilization
- Kernel-space and user-space threads
 - Kernel threads are real, schedulable threads
 - User-space may define its own threads (but not real)
- Threading models and implications
 - Programming systems
 - Multi-threaded design issues
- Useful, but not a panacea
 - Slow down system in some cases
 - Can be difficult to program
- Multiprogramming and multithreading are important concepts in modern operating systems