



CIS 415

Operating Systems

Main Memory

Prof. Allen D. Malony

Department of Computer and Information Science
Spring 2020



UNIVERSITY OF OREGON

Logistics

- Congratulations on the midterm!
 - Hope everything went ok in taking the test
 - Hope you thought it was a fair evaluation
 - Questions were shuffled in each section
 - I will try to have them graded by next Tuesday
- Keep working hard on Project 2
 - Lab sessions will focus on the project

Outline

- Background
- Swapping
- Contiguous Memory Allocation
- Segmentation
- Paging
- Structure of the Page Table
- Examples
 - The Intel 32 and 64-bit Architectures
 - ARM Architecture

Objectives

- To provide a detailed description of various ways of organizing memory hardware
- To discuss various memory-management techniques, including paging and segmentation
- To give a detailed example pure segmentation and segmentation with paging

Background

- Program must be brought (from external storage) into memory and “loaded” in a process for it to be run
 - A program goes through several steps before producing an executable code image that the OS uses to create a process
- A program starts with “logical” addresses that are in the “logical” address space its process
 - Instructions produce addresses with respect to logical addresses (also called “virtual” addresses)
 - Need to bind (map between) “logical” addresses and “physical” addresses in memory
- All this depends on both software and hardware
 - Software includes compilers, loaders, and OS
 - Hardware include memory system, address translation, ...

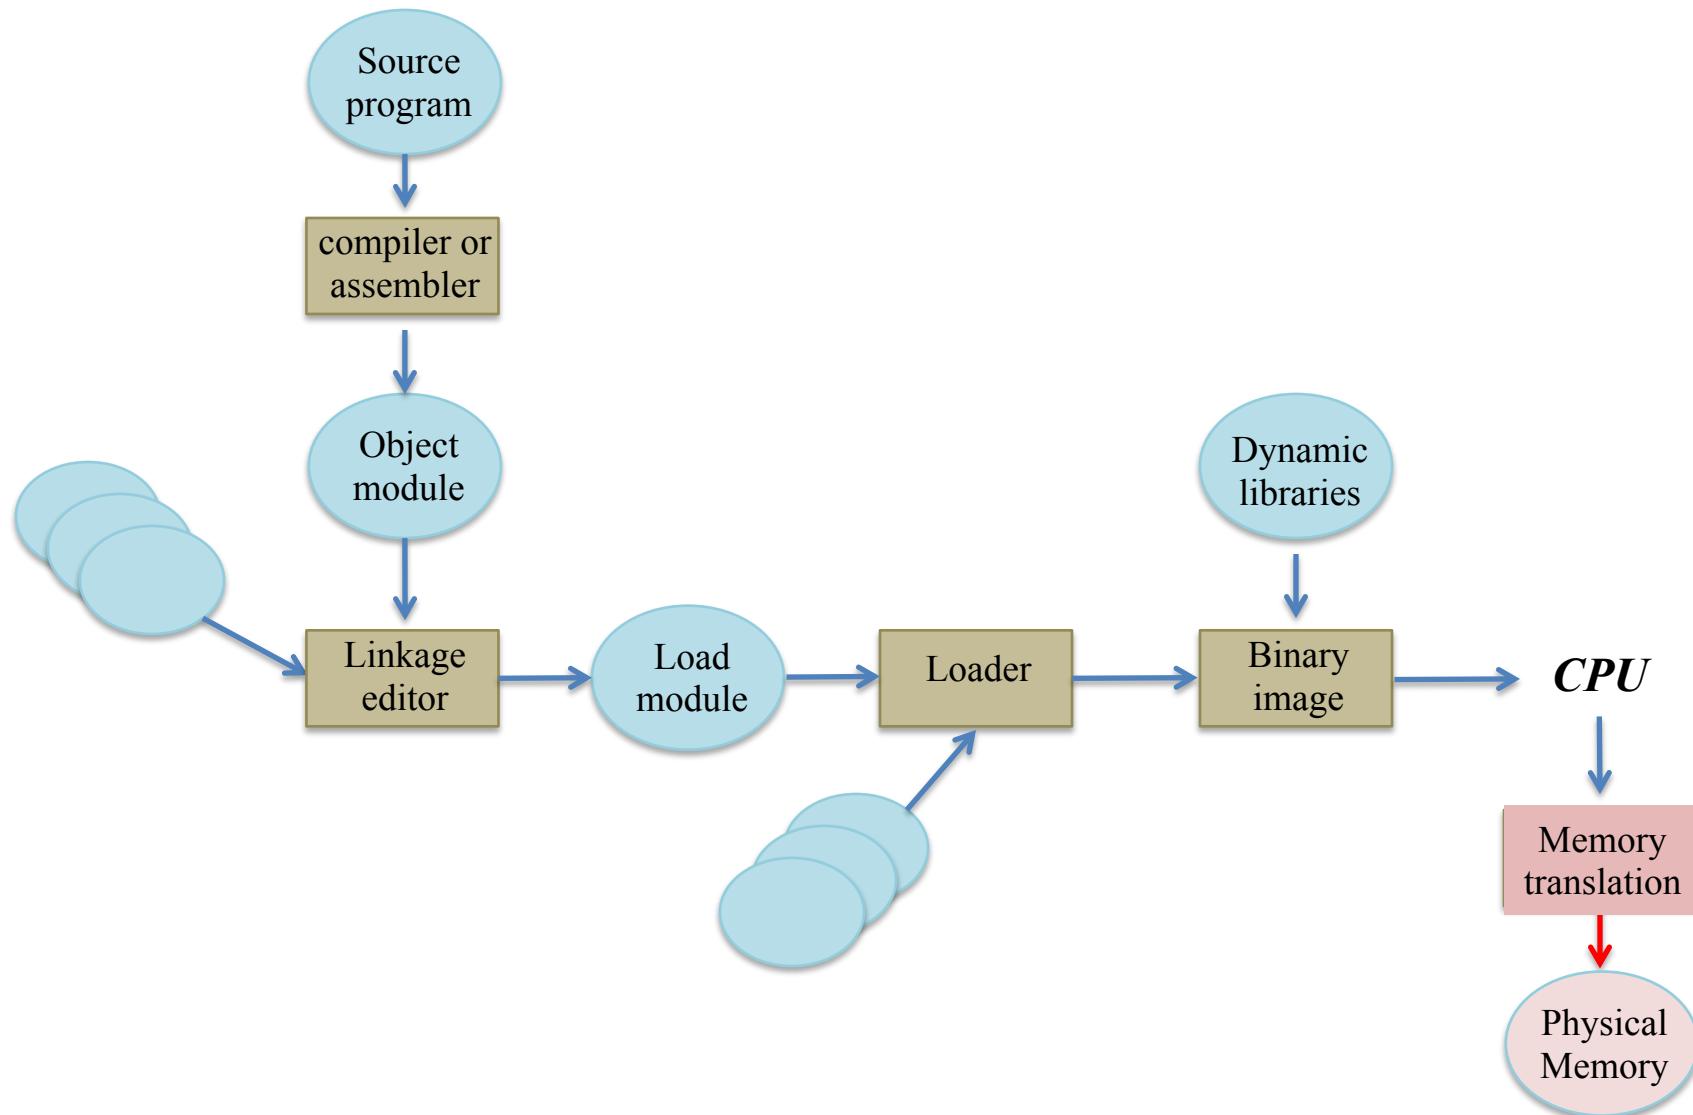
Address Binding

- Addresses represented in different ways at different stages of a program’s life
 - Need to bind one address to another
 - Source code addresses usually *symbolic* (i.e., names)
 - Compiled code has addresses that need to bind to *relocatable addresses*
 - Linker or loader will bind relocatable addresses to *absolute addresses*
 - Each binding maps one address space to another
- All of these addresses are “logical” addresses

Memory and Address Binding

- Main memory and CPU registers are the only storage directly addressable by a program
 - CPU registers are accessed in one CPU clock (or less)
 - Memory system sees addresses from read/write operations
- Addresses used in CPU instructions are logical addresses
 - Reference the logical address space of a process
 - These logical addresses must be “translated” to physical addresses
- What does address translation (address binding)?
 - *Compile time*: If a memory location is known *a priori*, code can be generated with absolute addresses, but needs to be recompiled if the location changes
 - *Load time*: If a memory location is known *relative* to an address, code can be generated with relocatable addresses
 - *Execution time*: addresses are mapped dynamically during execution using *memory management* hardware

Multistep Processing of a User Program

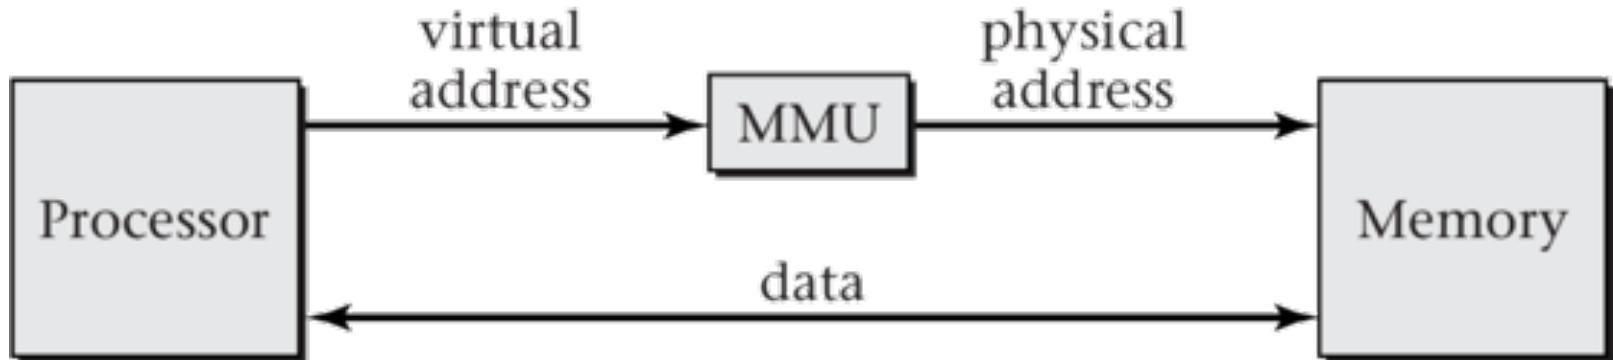


Logical vs. Physical Address Space

- The concept of a *logical address space* that is bound to a separate *physical address space* is central to proper memory management
 - *Logical address (virtual address)* – generated by the CPU
 - ◆ *logical address space* is the set of all logical addresses generated
 - *Physical address* – address seen by the memory unit
 - ◆ *physical address space* is the set of all physical addresses used
- Logical and physical addresses are the same in compile-time and load-time address-binding schemes
- Logical and physical addresses MUST be different at execution time ... Why?
 - Hardware-based address binding schemes here must do this
 - They must be fast!

Memory Management Unit

- Hardware device that maps a logical (virtual) address to a physical address
- How it does this is at the heart of all memory management schemes
- Memory management unit (*MMU*) changes the logical address into a physical address



Contiguous Allocation

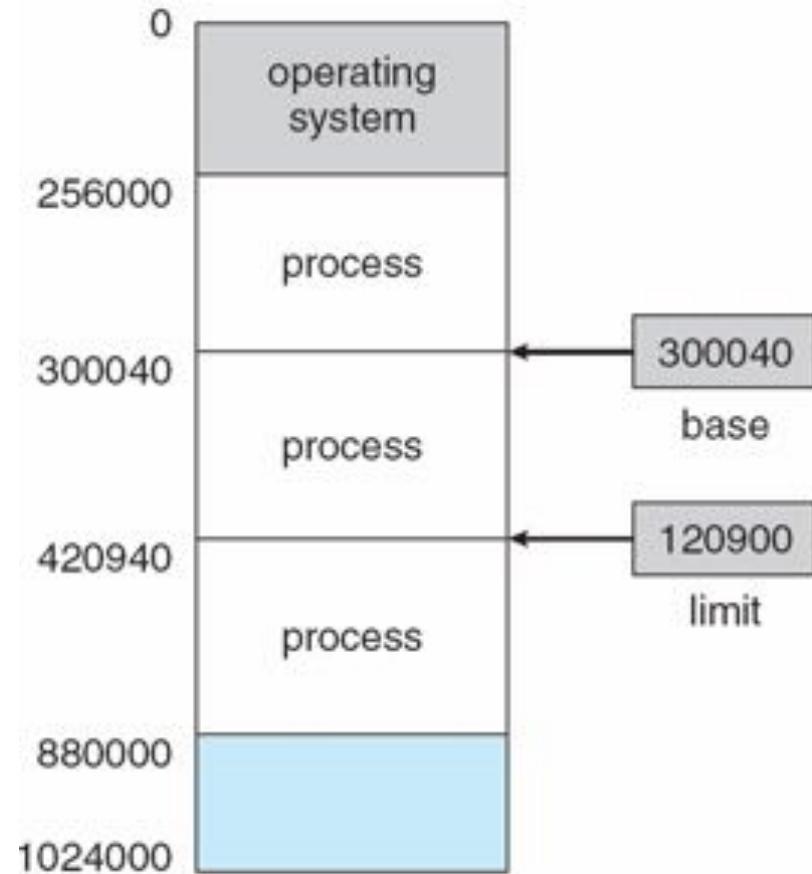
- Main (physical) memory must support both OS and user processes
 - It is a limited resource
 - Must allocated efficiently
- *Contiguous allocation* is one early method
- Main memory usually into two *partitions*:
 - Resident operating system
 - ◆ usually held in low memory with interrupt vector
 - User processes then held in high memory
 - Each process contained in single contiguous section (partition) of memory

Contiguous Allocation Address Translation

- One mechanism is to offset all logical address
 - A *base register* contains a value that is added to every logical address before it is sent to memory
 - A *limit register* is used to set a bound on the maximum address
- Another mechanism is to *relocate* an address
 - One or more relocation register(s) contains a value that is added to every logical address of a memory region before it is sent to memory
- User program's generate logical addresses
 - Their logical addresses change to physical addresses

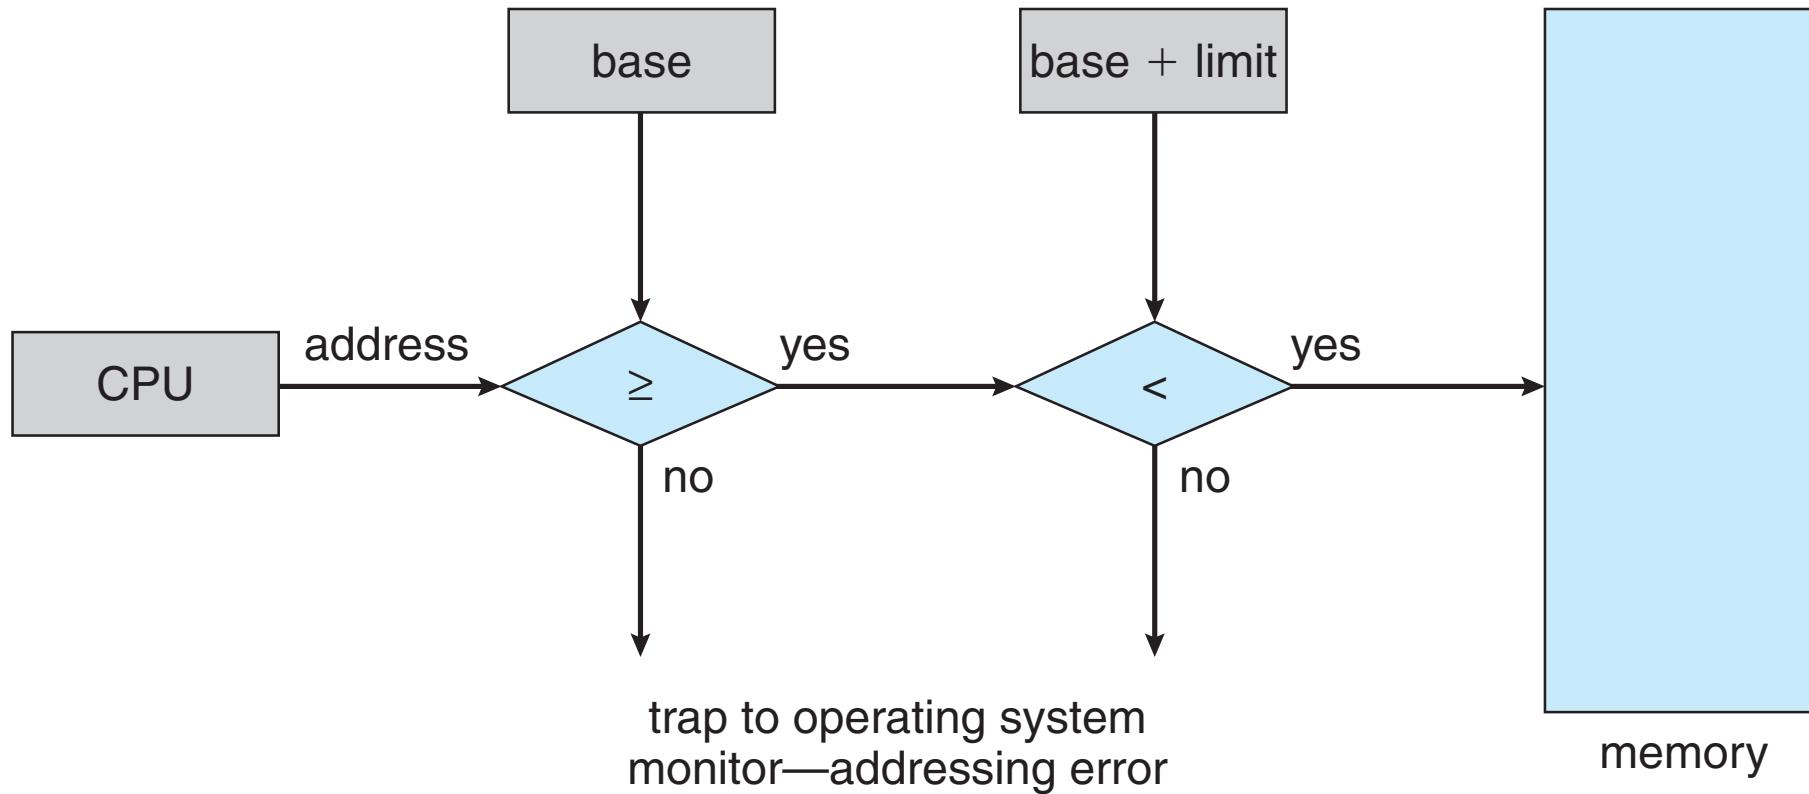
Base and Limit Registers with Partitioning

- How do we support the *logical address space*
 - A pair of *base* and *limit registers*
- CPU must check every memory access generated in user mode to be sure it is between base and limit for that user
 - If not, an exception occurs
- Need hardware support for doing this
- Useful for mapping entire process address space into a fixed address region
- OS sets up base/limit values



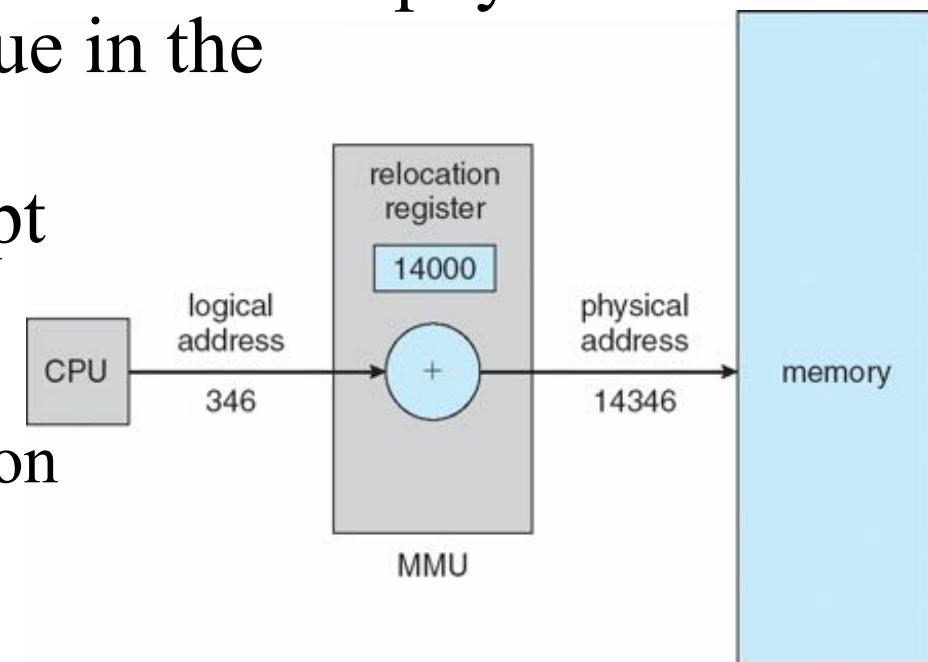
Hardware Address Protection

- If the OS did the check, it would be slow
- Hardware is used instead



Dynamic Relocation via Relocation Register

- CPU logical address is converted to a physical address by adding the value in the relocation register
- Like a base register, except you can have multiple relocation registers
 - One for each memory region
 - Code region
 - Heap region
 - Different routines (dynamic loading)



- Little support from the operating system is required
 - Implemented through program design
 - Used in coordination with base/limit registers

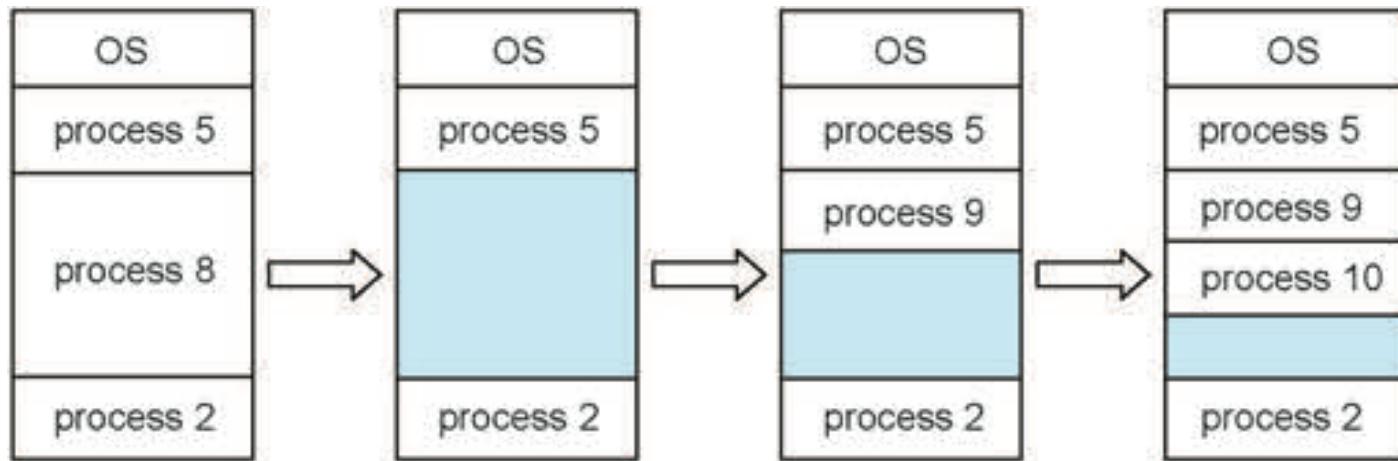
Dynamic Linking

- *Static linking* – system libraries and program code combined by the loader into the binary program image
- *Dynamic linking* –linking postponed until execution time
- Small piece of code (stub) used to locate the appropriate memory-resident library routine
- Stub replaces itself with the address of the routine and executes the routine
- OS checks if routine is in the processes memory address
 - If not in address space, add to address space
- Dynamic linking is particularly useful for libraries (Why?)
 - System libraries loaded when the program starts are also known as *shared libraries*
- Consider dynamic linking use to patching system libraries
 - Versioning may be needed

Multiple-partition Allocation

□ Multiple-partition allocation

- Degree of multiprogramming limited by number of partitions
- Variable-partition sizes for efficiency (sized to a given process' needs)
- Hole – block of available memory; holes of various size are scattered throughout memory
- When a process arrives, it is allocated memory from a hole large enough to accommodate it
- Process exiting frees its partition, adjacent free partitions combined
- Operating system maintains information about:
 - a) allocated partitions b) free partitions (*hole*)



Dynamic Storage Allocation Problem

- How to satisfy a request of size n ?
 - Assume there is a list of free holes
- *First-fit*
 - Allocate the first hole that is big enough
- *Best-fit*
 - Allocate the smallest hole that is big enough; must search entire list, unless ordered by size
 - Produces the smallest leftover hole
- *Worst-fit*
 - Allocate the largest hole; must also search entire list
 - Produces the largest leftover hole
- First-fit and best-fit better than worst-fit in terms of speed and storage utilization

Fragmentation

- *External fragmentation*
 - Total memory space exists to satisfy a request, but it is not contiguous
- *Internal fragmentation*
 - Allocated memory may be slightly larger than requested memory
 - This size difference is memory internal to a partition, but not being used
- First fit analysis reveals that given N blocks allocated, $0.5 N$ blocks lost to fragmentation
 - $1/3$ may be unusable -> *50-percent rule*

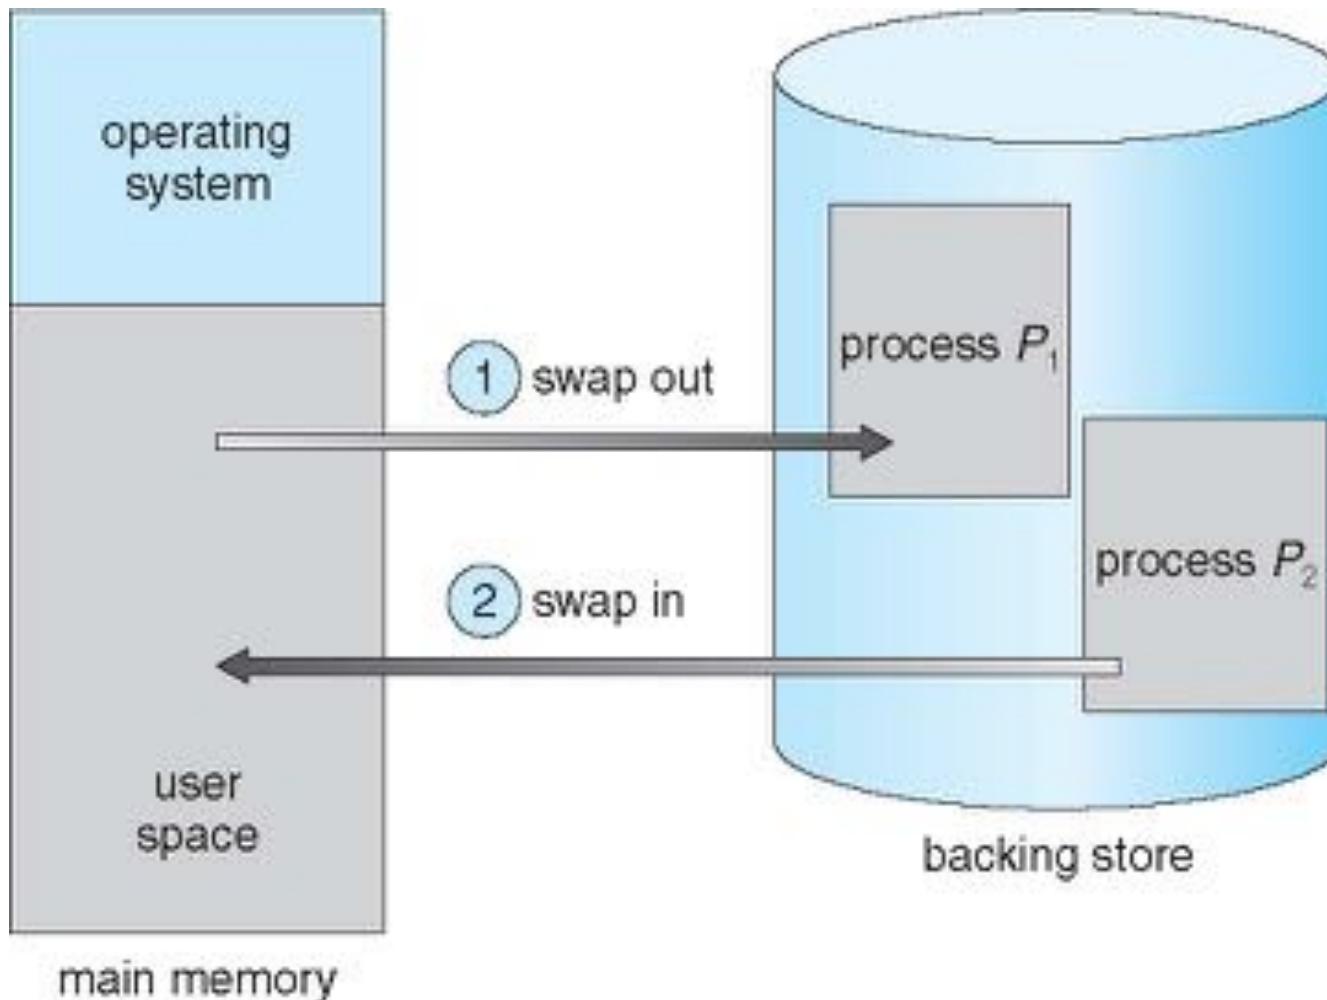
Fragmentation Improvement

- Reduce external fragmentation by *compaction*
 - Shuffle memory contents to place all free memory together in one large block
 - Compaction is possible only if relocation is dynamic, and is done at execution time
 - I/O problem
 - ◆ keep job in memory while it is involved in I/O
 - ◆ do I/O only into OS buffers

Swapping

- Physical memory space is finite
- OS takes up some portion of the physical memory
- Suppose we allocate physical memory to meet the needs of a process as discussed
 - What happens when the total physical memory space requested by processes exceeds physical memory?
 - Just consider 2 processes
- Swapping is the idea of sharing physical memory space by having only 1 process use the space at any time
- Still think about this with respect to contiguous memory allocation for a process

Schematic View of Swapping



Swapping Support

- A process can be *swapped* temporarily out of memory
 - Moved to a *backing store*— fast disk large enough to accommodate copies of all memory images for all users
 - Later brought back into memory for continued execution
- Swapping is used along with process scheduling
 - *Roll out, roll in* – swapping variant used for priority-based scheduling algorithms
 - Lower-priority process is swapped out so higher-priority process can be loaded and executed
- Major part of swap time is transfer time
 - Total transfer time is directly proportional to the amount of memory swapped
- System maintains a ready queue of ready-to-run processes which have memory images on disk

Swapping Methods

- Does the swapped out process need to swap back in to same physical addresses?
- Depends on address binding method
 - Also consider pending I/O since it might be associated with particular addresses
- Modified versions of swapping are found on many systems (i.e., UNIX, Linux, and Windows)
 - Swapping normally disabled
 - Started if greater than a memory allocation threshold
 - Disabled again once memory demand reduced

Context Switch Time including Swapping

- If next processes to be put on CPU is not in memory, need to swap out a process and swap in target process
- Context switch time can then be very high
 - 100MB process swapping to hard disk
 - Transfer rate of 50MB/sec
 - ◆ swap out time of 2000 ms
 - ◆ plus swap in of same sized process
 - ◆ total context switch swapping component time of 4000ms
- Can reduce if reduce size of memory swapped
 - Know how much memory really being used
 - System calls to inform OS of memory use via *request_memory()* and *release_memory()*

Swapping Constraints

- Other constraints as well on swapping
 - Can not swap out if process has pending I/O
 - ◆ otherwise I/O would occur to wrong process
 - Could transfer I/O to kernel space, then to I/O device
 - ◆ known as *double buffering*, adds overhead
- Standard swapping not really used in modern operating systems
 - But modified version common
 - ◆ swap only when free memory extremely low

Swapping on Mobile Systems

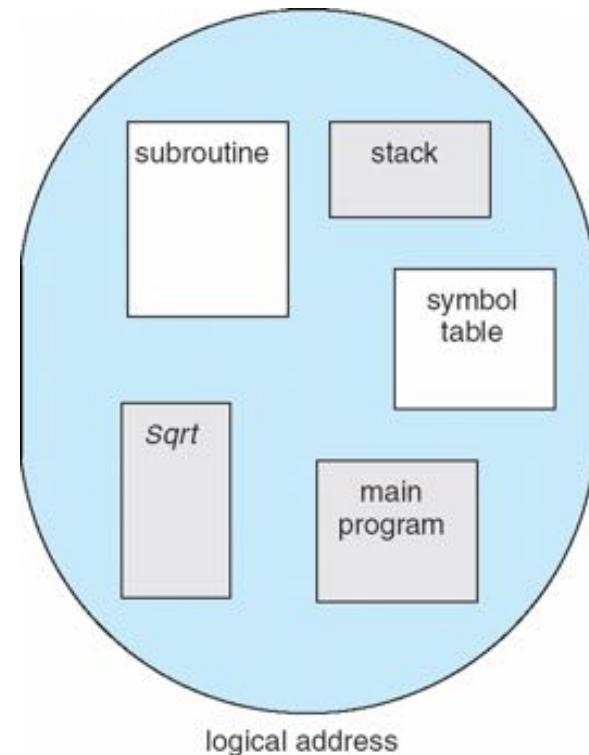
- Not typically supported
 - Flash memory based
 - ◆ small amount of space
 - ◆ limited number of write cycles
 - ◆ poor throughput between flash memory and CPU for mobile
- Instead use other methods to free memory if low
 - iOS asks apps to voluntarily relinquish allocated memory
 - ◆ Read-only data thrown out and reloaded from flash if needed
 - ◆ Failure to free can result in termination
 - Android terminates apps if low free memory, but first writes application state to flash for fast restart
 - Both OSes support paging as discussed below

Non-contiguous Memory Allocation

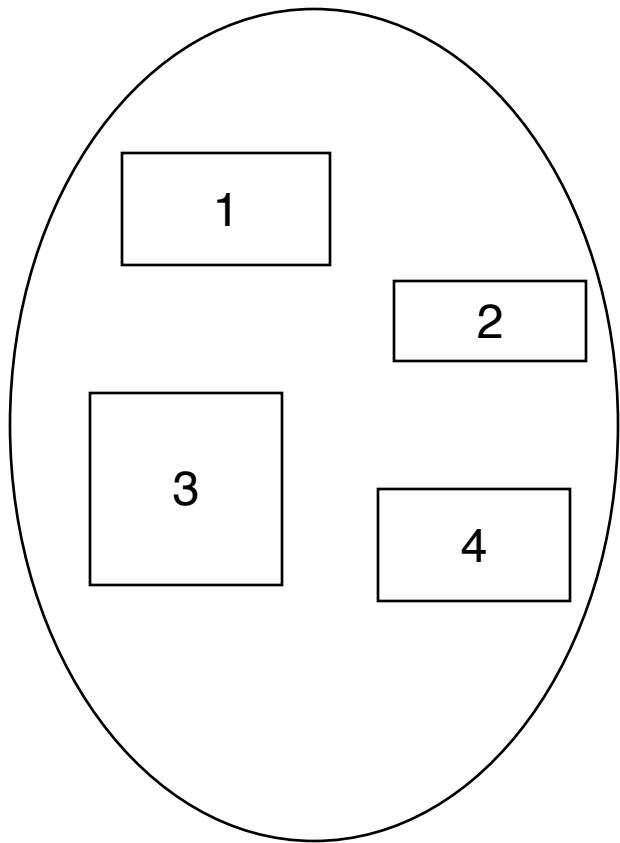
- Contiguous memory allocation is not all that great for managing memory
- Its advantages are that address translation is very straightforward and fast
 - Just need a base and limit register
- However, there are disadvantages
 - All logical addresses must be contiguous
 - Holes can develop requiring compaction to fix
- Suppose we loosen constraints and allow for non-contiguous memory allocation
 - Segments
 - Pages

Segmentation

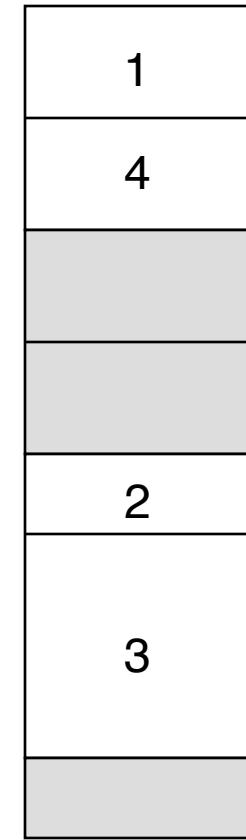
- Memory-management scheme that supports a user view of memory
 - Allows different parts to be allocated separately
- A program is a collection of segments
 - A segment is a logical unit
 - Instructions
 - ◆ main program
 - ◆ procedures and functions
 - ◆ libraries
 - Data
 - ◆ global variables
 - ◆ stack
 - ◆ arrays
- Associate (continuous) logical addresses with each segment



Logical View of Segmentation



user space

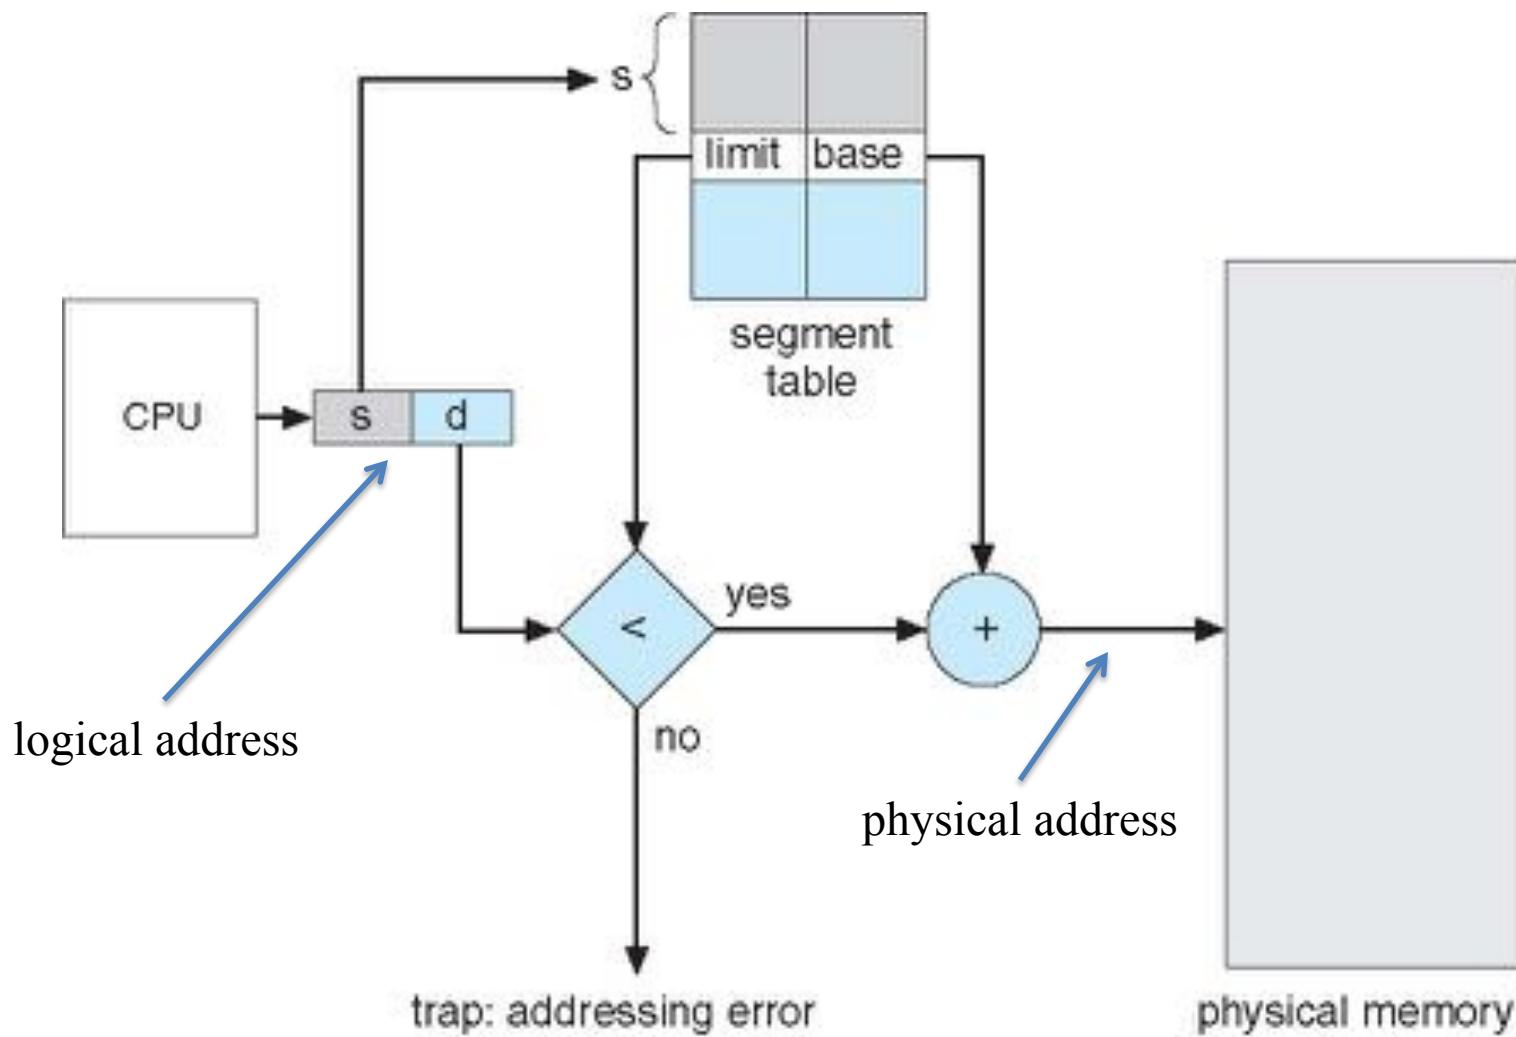


physical memory space

Segmentation Architecture

- Logical address consists of a two tuple:
 $\langle \text{segment-number}, \text{offset} \rangle$
- *Segment table*
 - Maps two-dimensional physical addresses
 - Each table entry has:
 - ◆ *base register* containing the starting physical address where the segments reside in memory
 - ◆ *Limit register* specifying the length of the segment
- *Segment-table base register (STBR)* points to the segment table's location in memory
- *Segment-table length register (STLR)* indicates number of segments used by a program
 - segment number s is legal if $s < \text{STLR}$

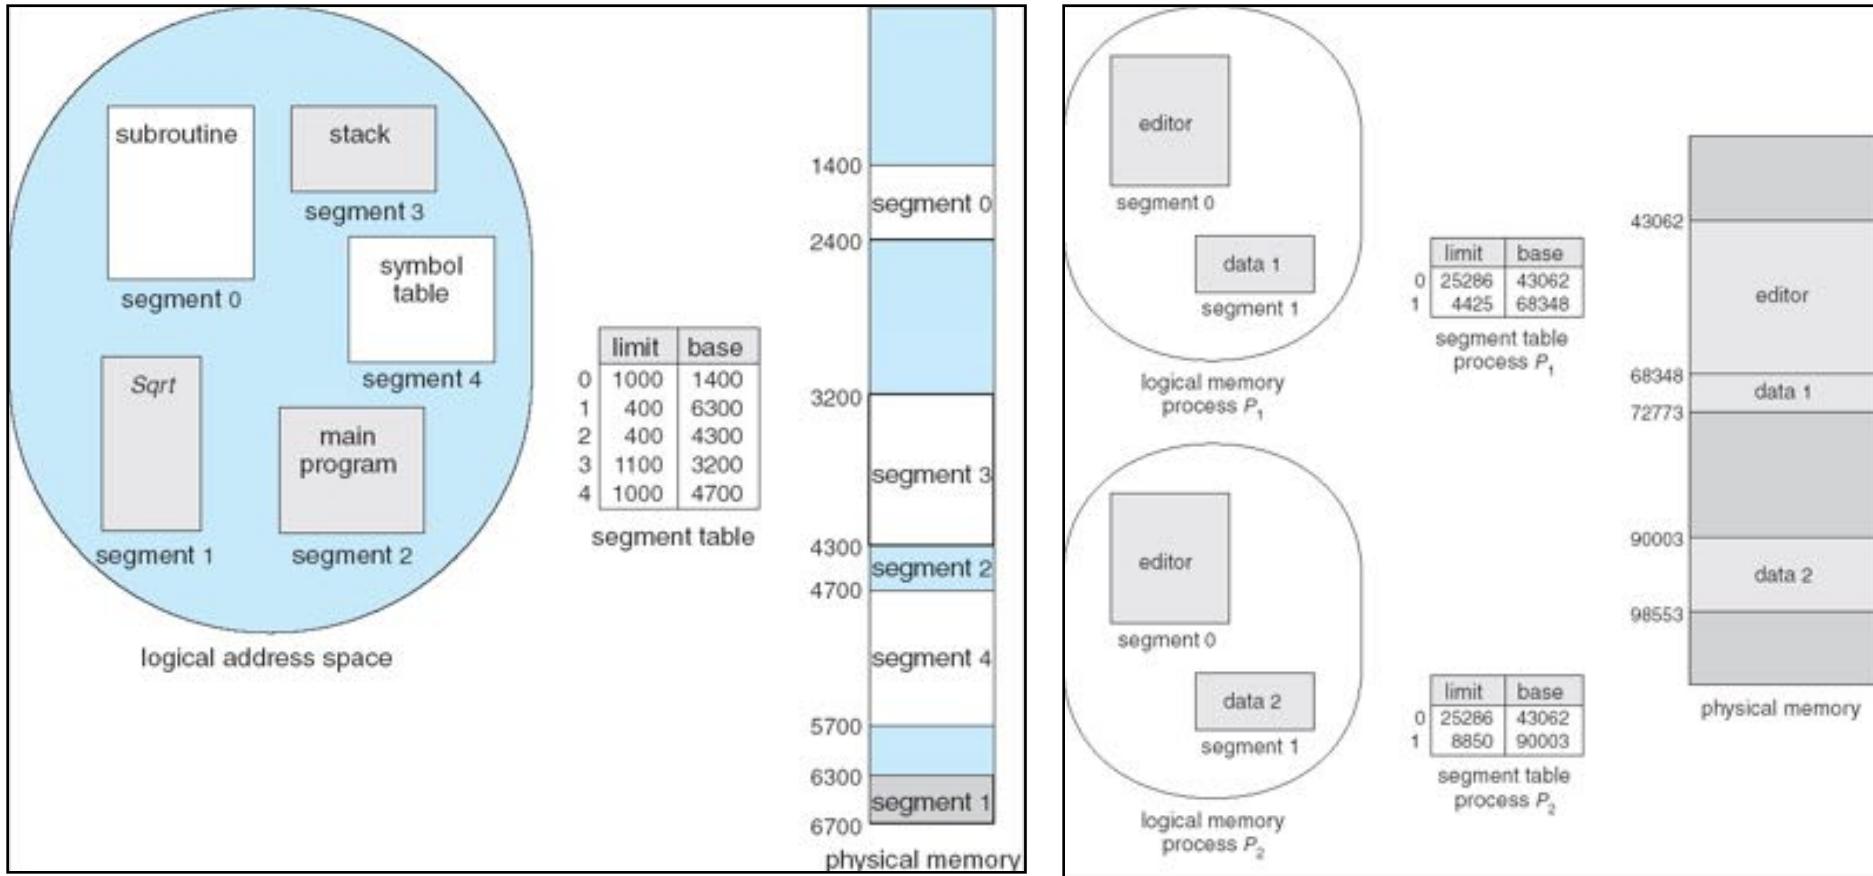
Segmentation Hardware



Segmentation Architecture Protection

- Protection
 - With each entry in segment table associate:
 - ◆ validation bit = 0 \Rightarrow illegal segment
 - ◆ read/write/execute privileges
- Protection bits associated with segments
 - Code sharing occurs at segment level
- Since segments vary in length, memory allocation is a dynamic storage-allocation problem
 - Could possibly still result in areas of physical memory that are too small to be allocated to a segment (holes)
- A segmentation example is shown next

Segmentation and Segment Sharing Example

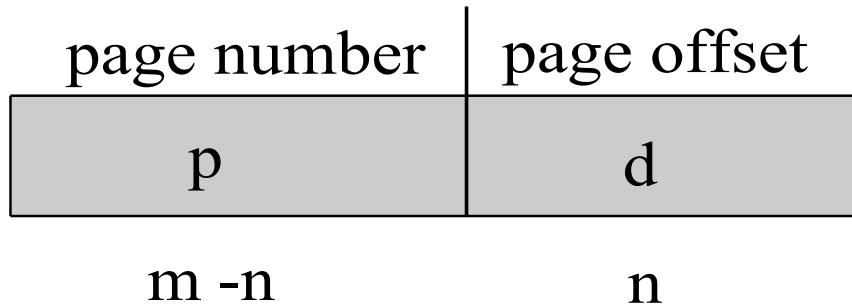


Paging

- Suppose we want to eliminate external fragmentation entirely
- Divide physical memory into fixed-sized blocks called *frames*
 - Size is power of 2, between 512 bytes and 16 Mbytes
- Divide logical memory into blocks of same size called *pages*
- Keep track of all free frames
- To run a program of size N pages, need to find N free frames and load program
- Set up a page table to translate logical to physical addresses
- Physical address space of a process is noncontiguous
- Memory management reduces to finding frames for pages
 - Process is allocated physical memory wherever pages are available
 - Avoids external fragmentation
 - Avoids problem of varying sized memory chunks
- Backing store likewise split into pages!
- Still have internal fragmentation ... Why?

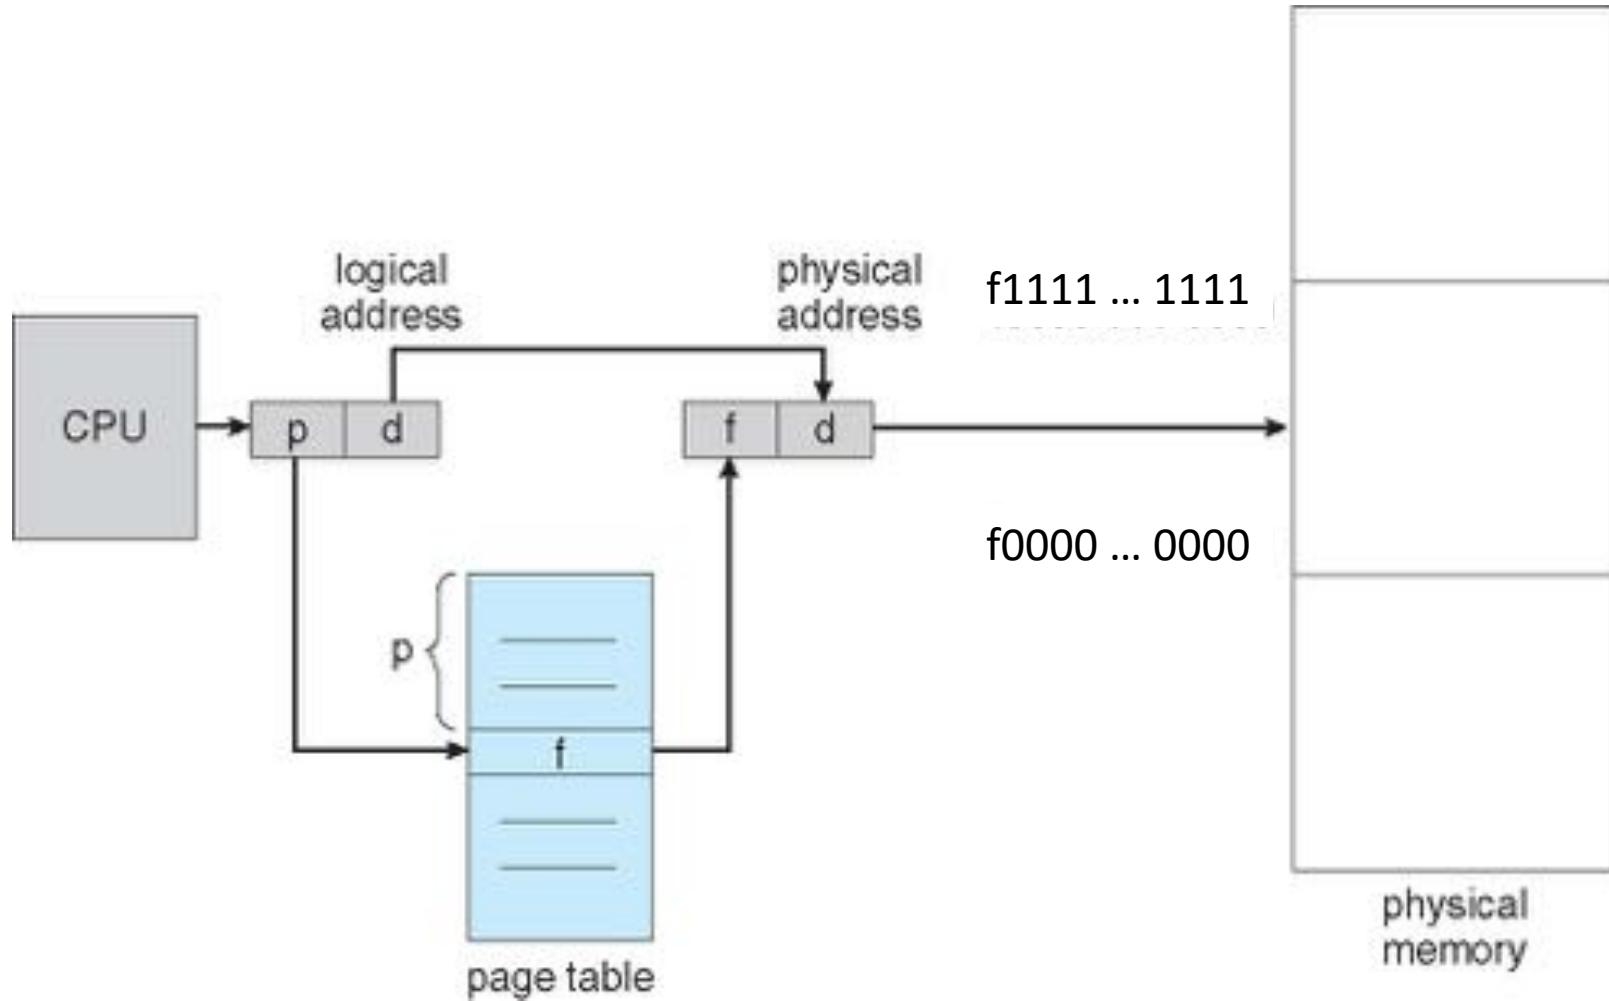
Address Translation Scheme

- Address generated by CPU is divided into:
 - *Page number (p)* – used as an index into a *page table* which contains base address of each page in physical memory
 - *Page offset (d)* – combined with base address to define the physical memory address that is sent to the memory unit

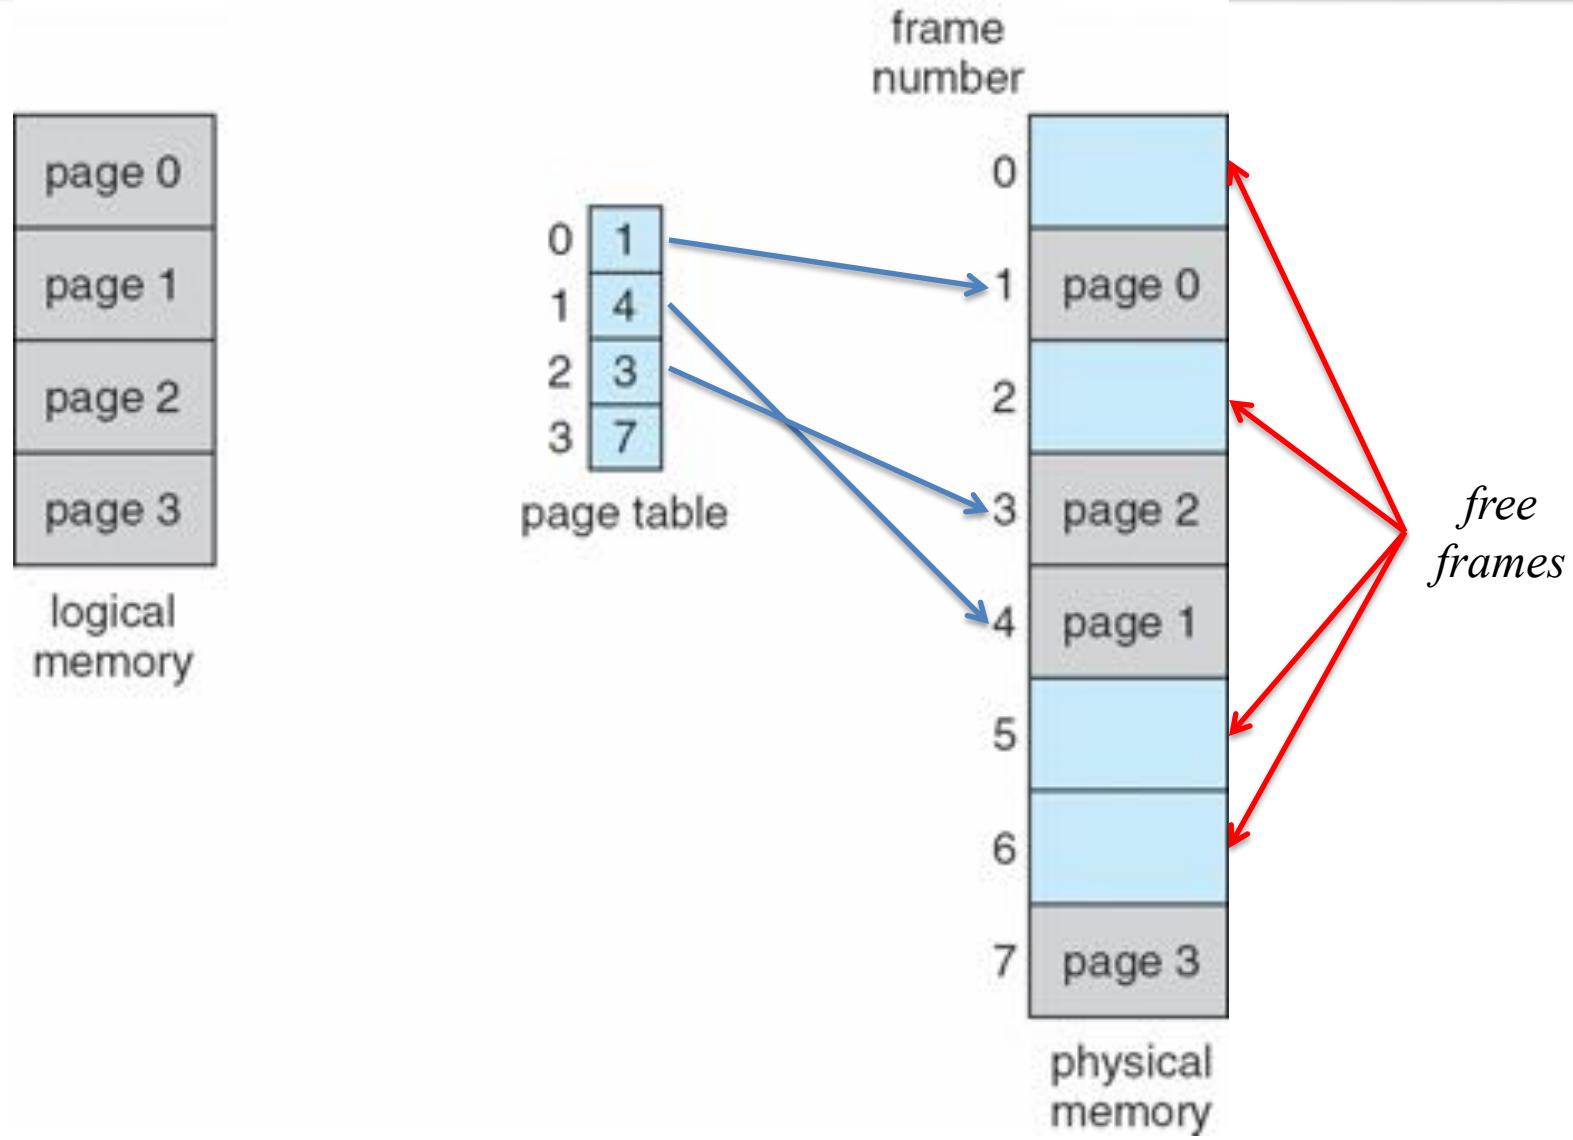


- For given logical address space 2^m and page size 2^n

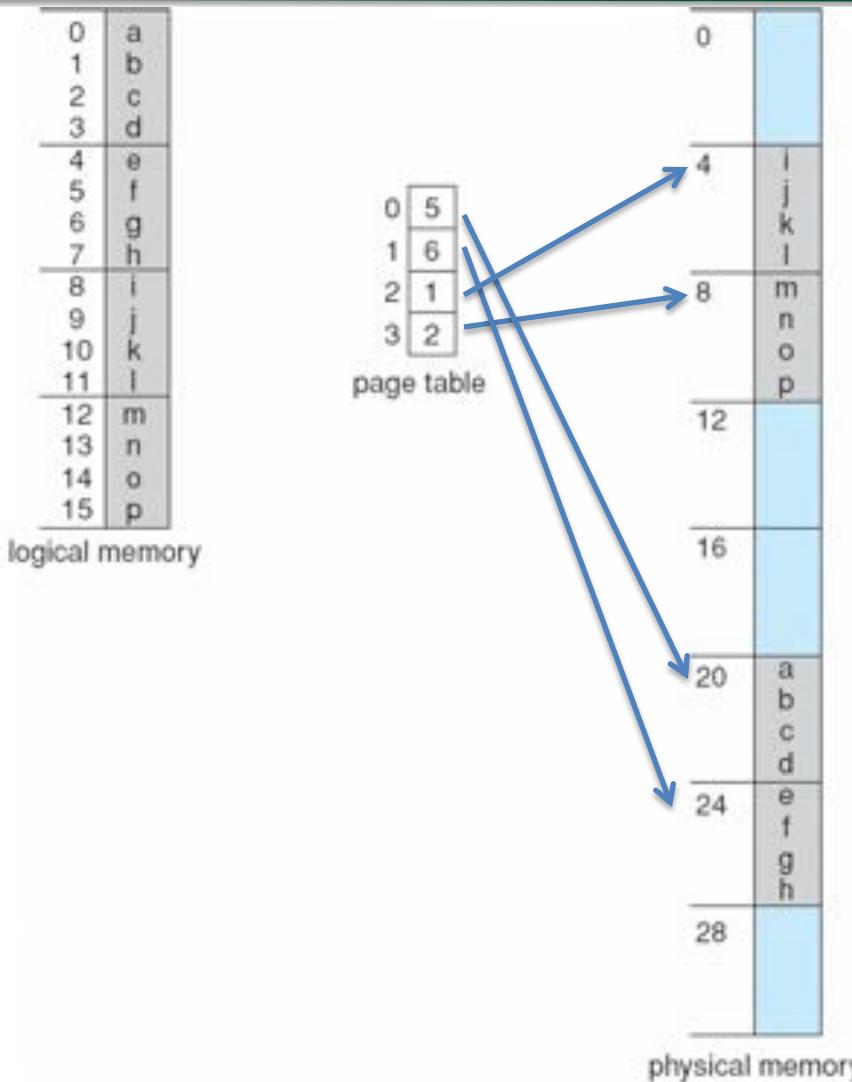
Paging Hardware



Paging Model of Logical / Physical Memory



Paging Example



$n=2$ and $m=4$ 32-byte memory and 4-byte pages

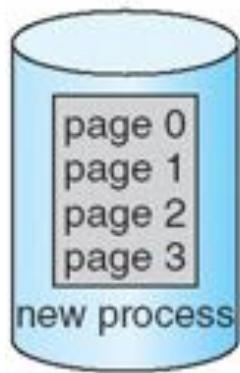
Paging and Fragmentation

- Calculating internal fragmentation
 - Page size = 2,048 bytes
 - Process size = 72,766 bytes
 - 35 pages + 1,086 bytes
 - Internal fragmentation of $2,048 - 1,086 = 962$ bytes
 - Worst case fragmentation = 1 frame – 1 byte
 - On average fragmentation = 1 / 2 frame size
 - So small frame sizes desirable?
 - But each page table entry takes memory to track
 - Page sizes growing over time
 - ◆ Solaris supports two page sizes – 8 KB and 4 MB
- Process view and physical memory now very different
- By implementation process can only access its own memory

Free Frames

free-frame list

14
13
18
20
15

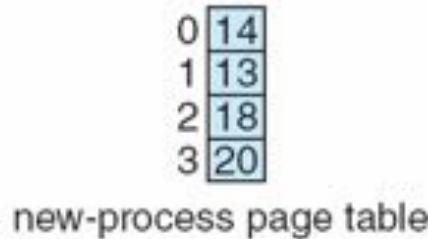
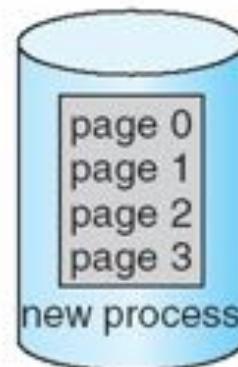


(a)

Before allocation

free-frame list

15



new-process page table



(b)

After allocation

Implementation of Page Table

- Page table is kept in main memory
- *Page-table base register (PTBR)* points to the page table
- *Page-table length register (PTLR)* indicates size of the page table
- In this scheme every data/instruction access requires two memory accesses
 - One for the page table and one for the data / instruction
- Two memory accesses can be improved by the use of a special fast-lookup hardware cache called *translation look-aside buffers (TLBs)*
 - Implemented with associative memory

Page Table TLB

- TLBs typically small (64 to 1,024 entries)
- On a TLB miss, value is loaded into the TLB for faster access next time
 - Replacement policies must be considered
 - Some entries can be wired down for permanent fast access
- Some TLBs store *address-space identifiers (ASIDs)* in each TLB entry – uniquely identifies each process to provide address-space protection for that process
 - Otherwise need to flush at every context switch

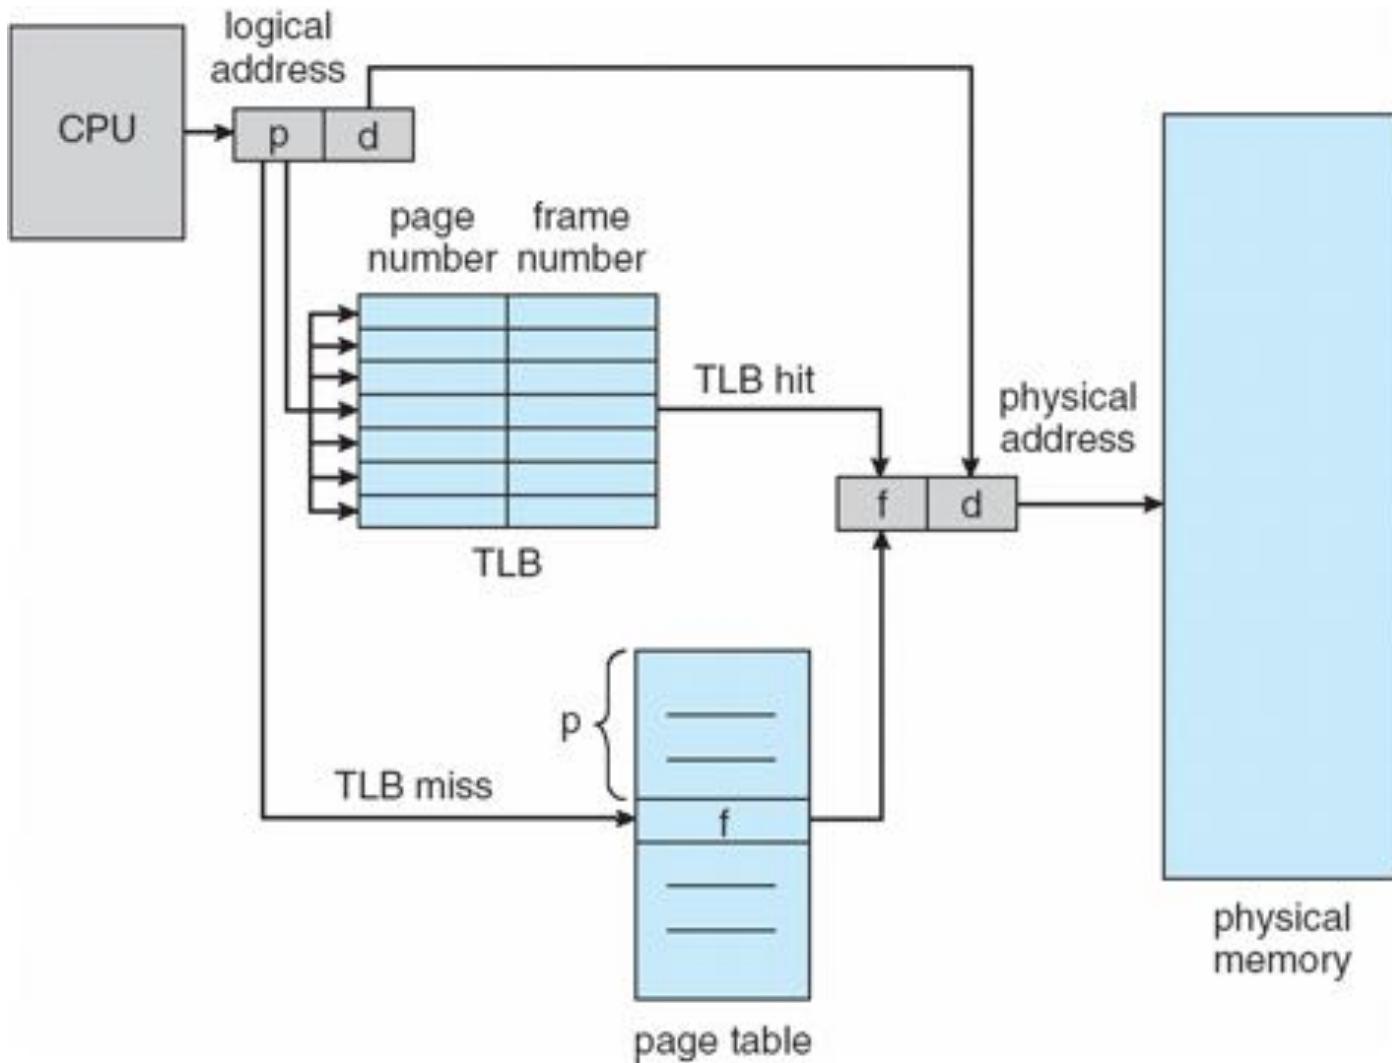
Associative Memory

- Associative memory – parallel search

Page #	Frame #

- Address translation (p, d)
 - If p is in associative register, get frame # out
 - Otherwise get frame # from page table in memory

Paging Hardware With TLB



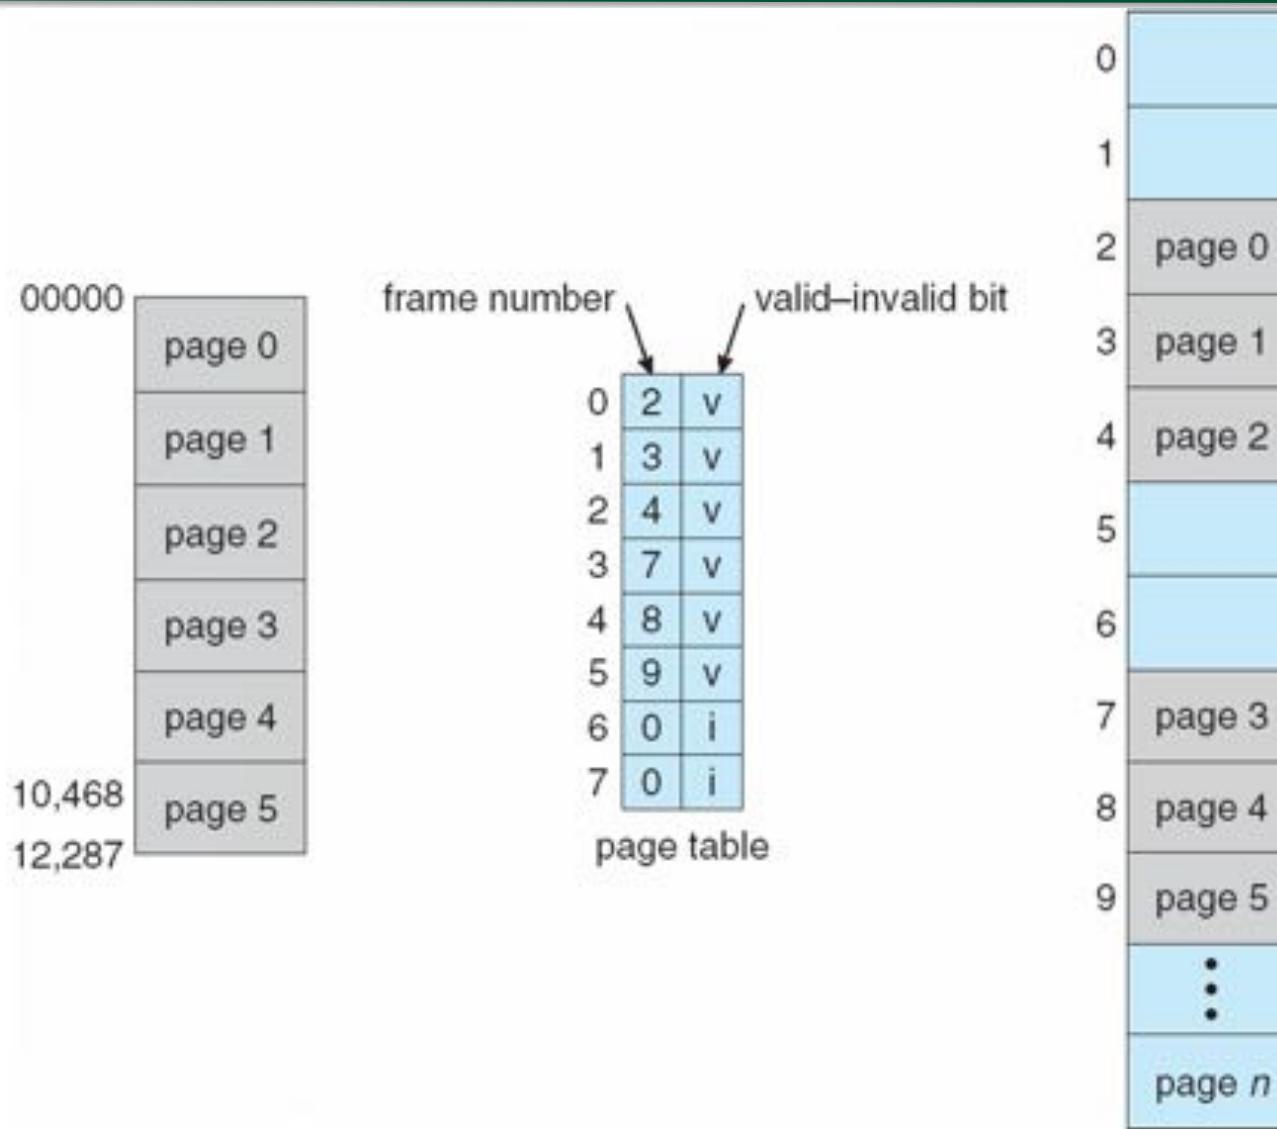
Effective Access Time

- Associative lookup = ε time unit
 - Can be < 10% of memory access time
- Hit ratio = α
 - *Hit ratio* – percentage of times that a page number is found in the associative registers; ratio related to number of associative registers
- Consider $\alpha = 80\%$, $\varepsilon = 20\text{ns}$ for TLB search, 100ns for memory access
- *Effective Access Time (EAT)*
$$\begin{aligned}\text{EAT} &= (1 + \varepsilon) \alpha + (2 + \varepsilon)(1 - \alpha) \\ &= 2 + \varepsilon - \alpha\end{aligned}$$
- Consider $\alpha = 80\%$, $\varepsilon = 20\text{ns}$ for TLB search, 100ns for memory access
 - $\text{EAT} = 0.80 \times 100 + 0.20 \times 200 = 120\text{ns}$
- Consider more realistic hit ratio -> $\alpha = 99\%$, $\varepsilon = 20\text{ns}$ for TLB search, 100ns for memory access
 - $\text{EAT} = 0.99 \times 100 + 0.01 \times 200 = 101\text{ns}$

Memory Protection

- Memory protection implemented by associating protection bit with each frame to indicate if read-only or read-write access is allowed
 - Can also add more bits to indicate page execute-only, and so on
- Use a *valid-invalid* bit attached to each entry in the page table:
 - “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal page
 - “invalid” indicates that the page is not in the process’ logical address space
 - Or use *page-table length register* (PTLR)
- Any violations result in a trap to the kernel

Valid (v) or Invalid (i) Bit In A Page Table



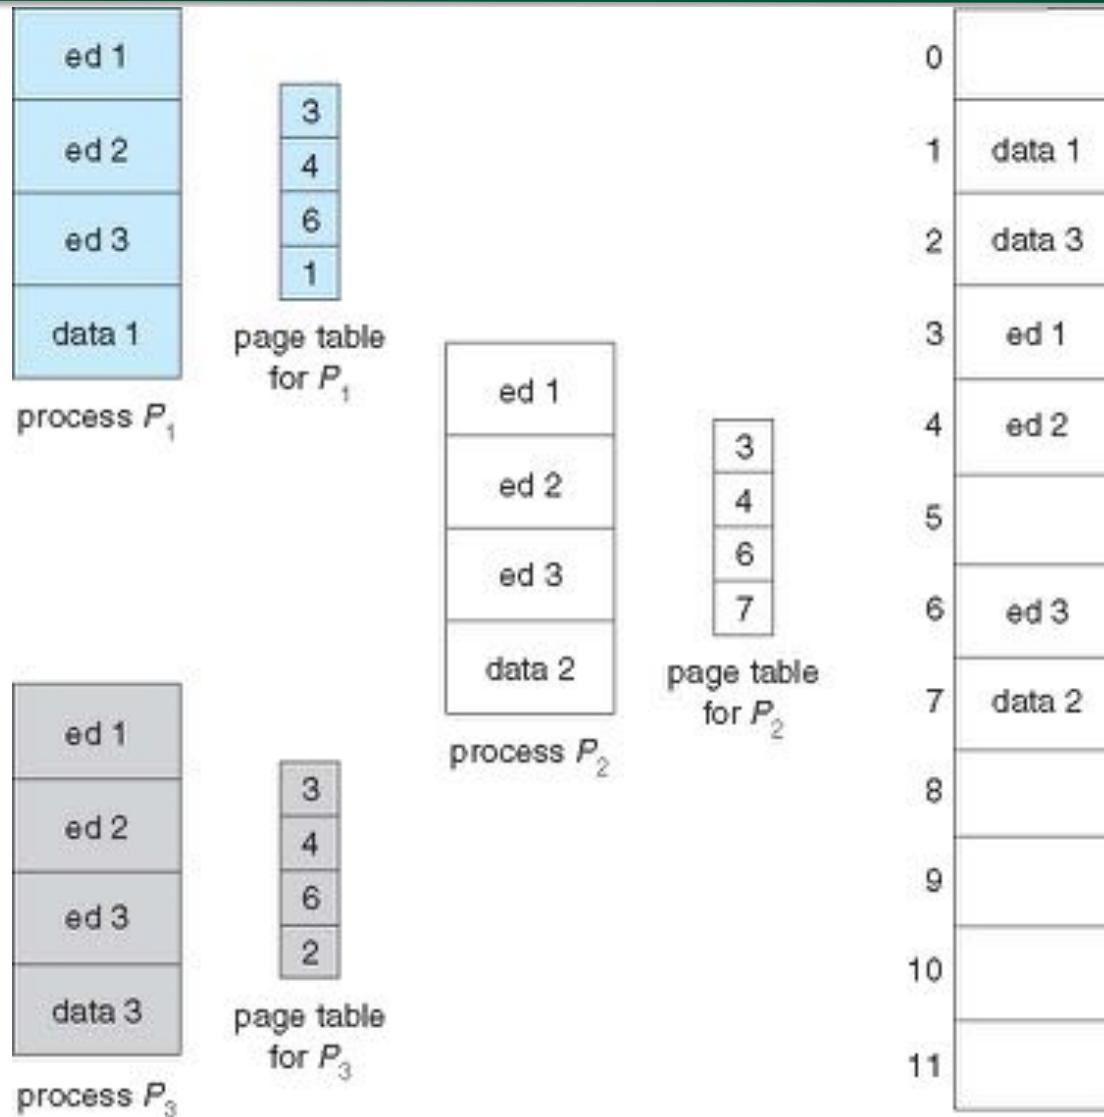
Shared Pages

- Shared code
 - One copy of read-only (reentrant) code shared among processes (i.e., text editors, compilers, window systems)
 - Similar to multiple threads sharing the same process space
 - Also useful for interprocess communication if sharing of read-write pages is allowed

□ Private code and data

- Each process keeps a separate copy of the code and data
- The pages for the private code and data can appear anywhere in the logical address space

Shared Pages Example



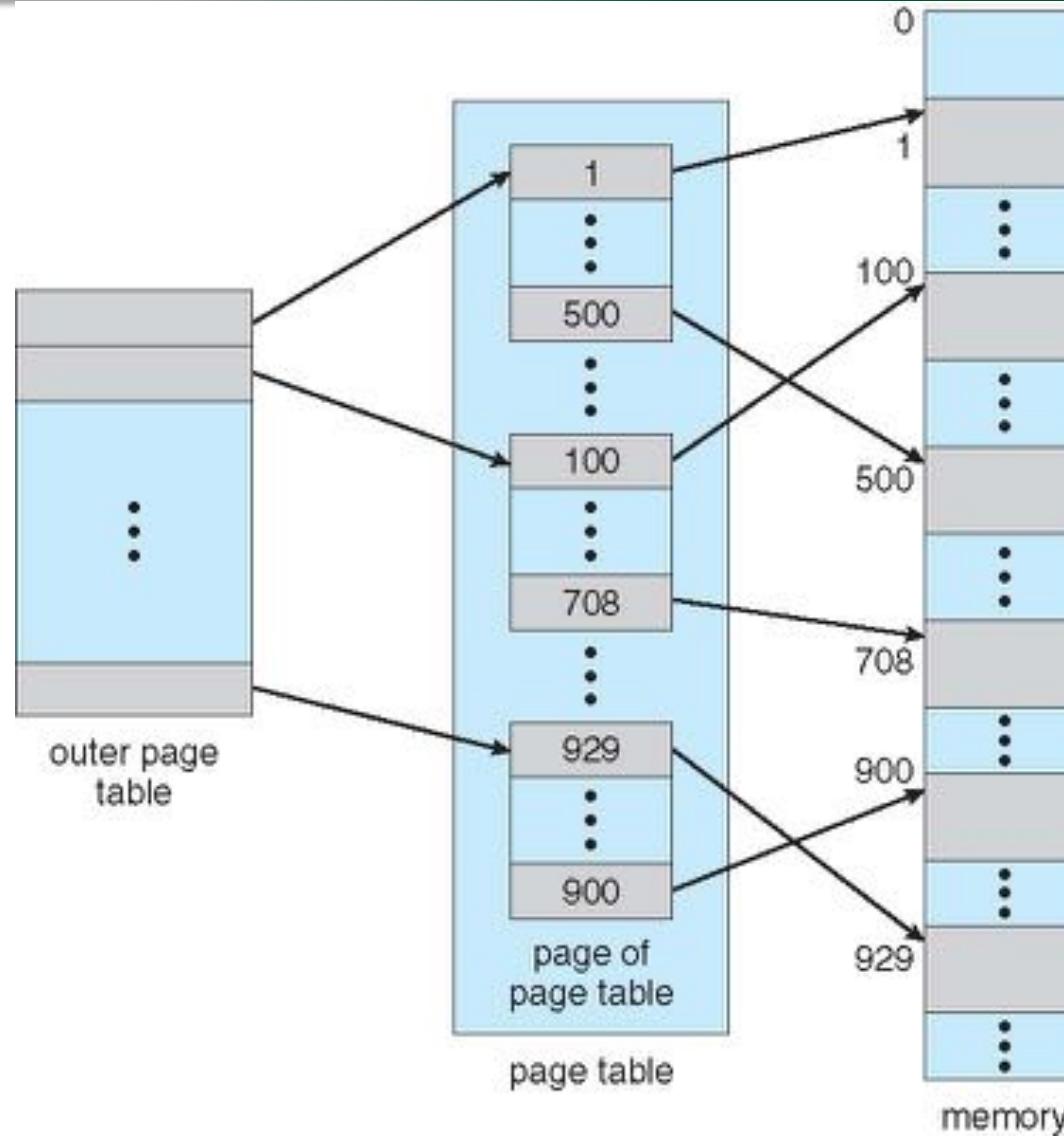
Structure of the Page Table

- Memory structures for paging can get huge using straight-forward methods
 - Consider a 32-bit logical address space as on modern computers
 - Page size of 4 KB (2^{12})
 - Page table would have 1 million entries ($2^{32} / 2^{12}$)
 - If each entry is 4 bytes -> 4 MB of physical address space / memory for page table alone
 - ◆ that amount of memory used to cost a lot
 - ◆ do not want to allocate that contiguously in main memory
- Hierarchical paging
- Hashed page tables
- Inverted page tables

Hierarchical Page Tables

- Break up the logical address space into multiple page tables
- A simple technique is a two-level page table
- We then page the page table

Two-Level Page-Table Scheme



Two-Level Paging Example

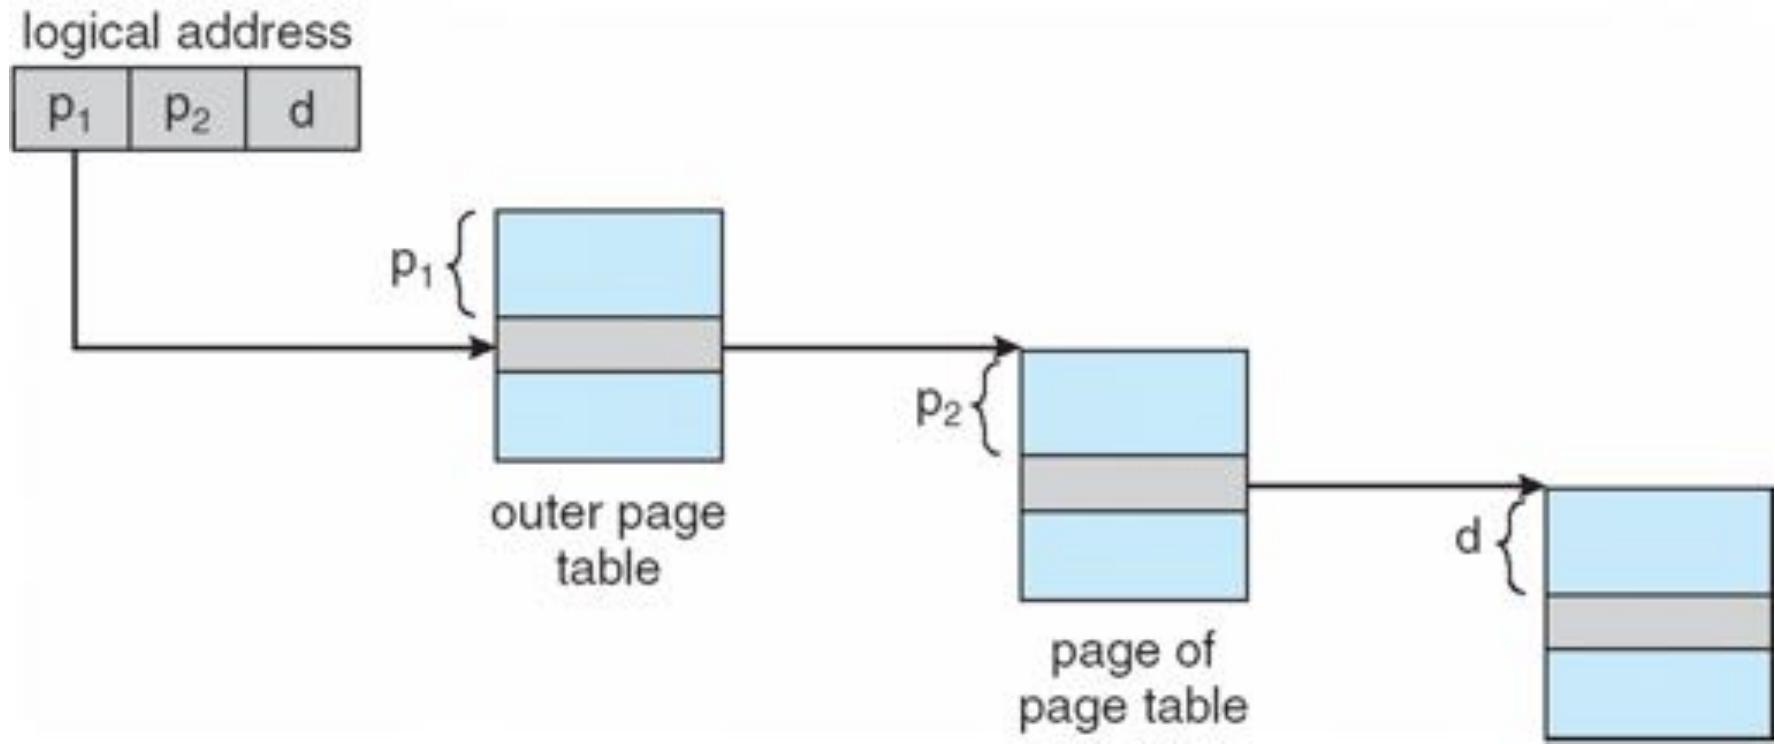
- A logical address (on 32-bit machine with 1K page size) is divided into:
 - A page number consisting of 22 bits
 - A page offset consisting of 10 bits
- Since the page table is paged, the page number is divided into:
 - 12-bit page number
 - 10-bit page offset
- Thus, a logical address is as follows:

page number	page offset
p_1	p_2
12	10

where p_1 is an index into the outer page table, and p_2 is the displacement within the page of the inner page table

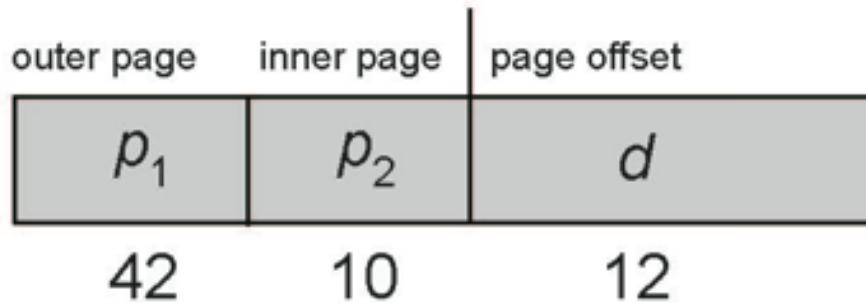
- Known as *forward-mapped page table*

Address-Translation Scheme



64-bit Logical Address Space

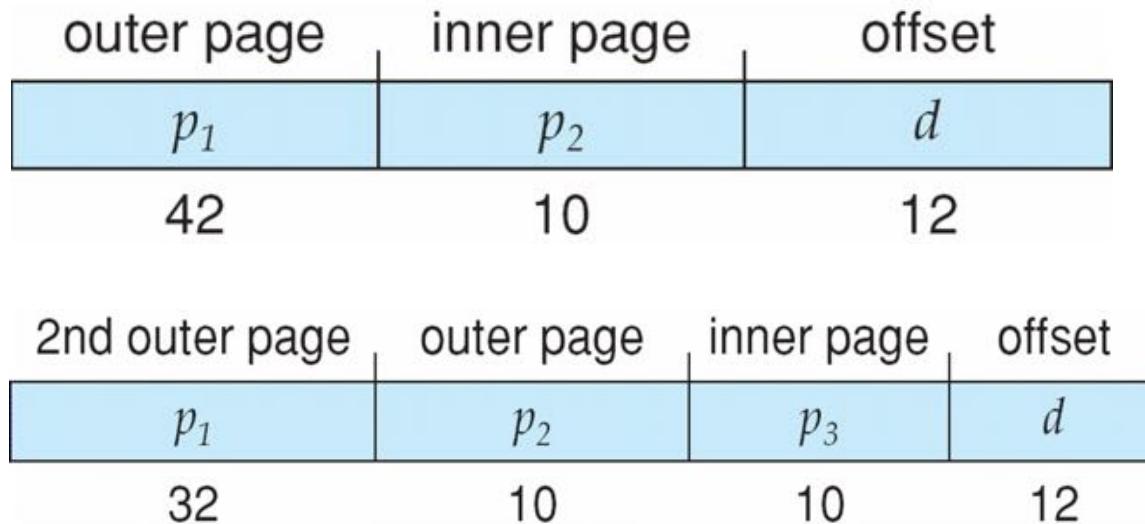
- Even two-level paging scheme not sufficient
- If page size is 4 KB (2^{12})
 - Then page table has 2^{52} entries
 - If two level scheme, inner page tables could be 2^{10} 4-byte entries
 - Address would look like



- Outer page table has 2^{42} entries or 2⁴⁴ bytes

2nd Outer Page Table

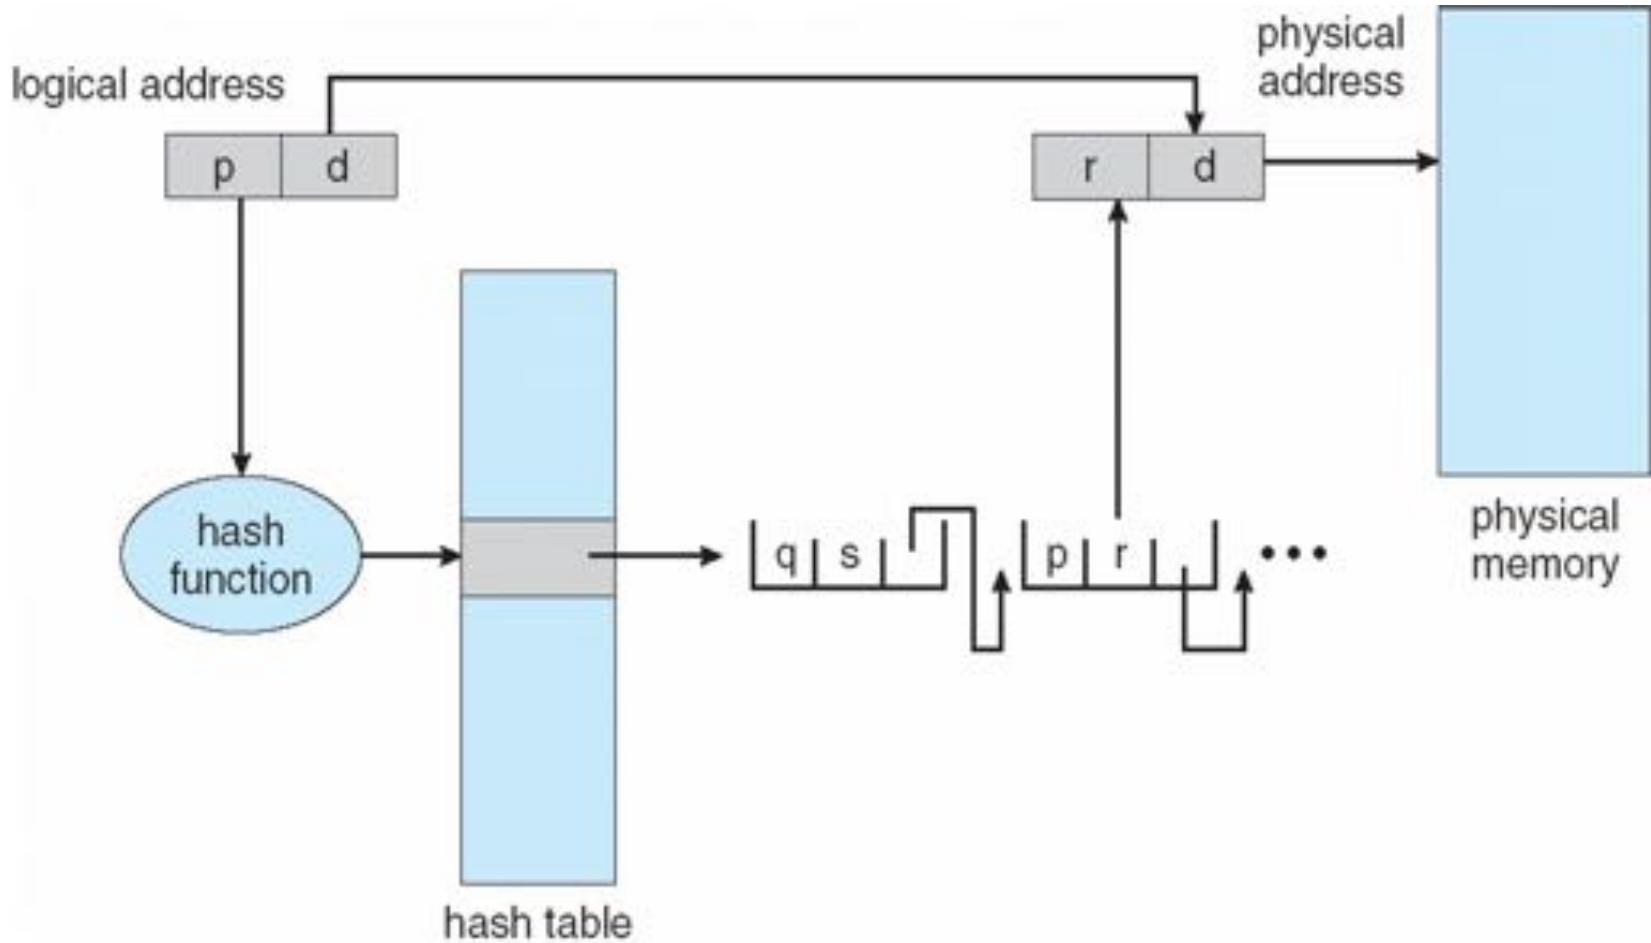
- One solution is to add a 2nd outer page table
- In this example, the 2nd outer page table is still 234 bytes in size
 - Possibly 4 memory access to get to one physical memory location



Hashed Page Tables

- Common in address spaces > 32 bits
- The virtual page number is hashed into a page table
 - This page table contains a chain of elements hashing to the same location
- Each element contains (1) the virtual page number (2) the value of the mapped page frame (3) a pointer to the next element
- Virtual page numbers are compared in this chain searching for a match
 - If a match is found, the corresponding physical frame is extracted
- Variation for 64-bit addresses is *clustered page tables*
 - Similar to hashed but each entry refers to several pages (such as 16) rather than 1
 - Especially useful for sparse address spaces (where memory references are non-contiguous and scattered)
- Hashing has to be done by the MMU ... Why?

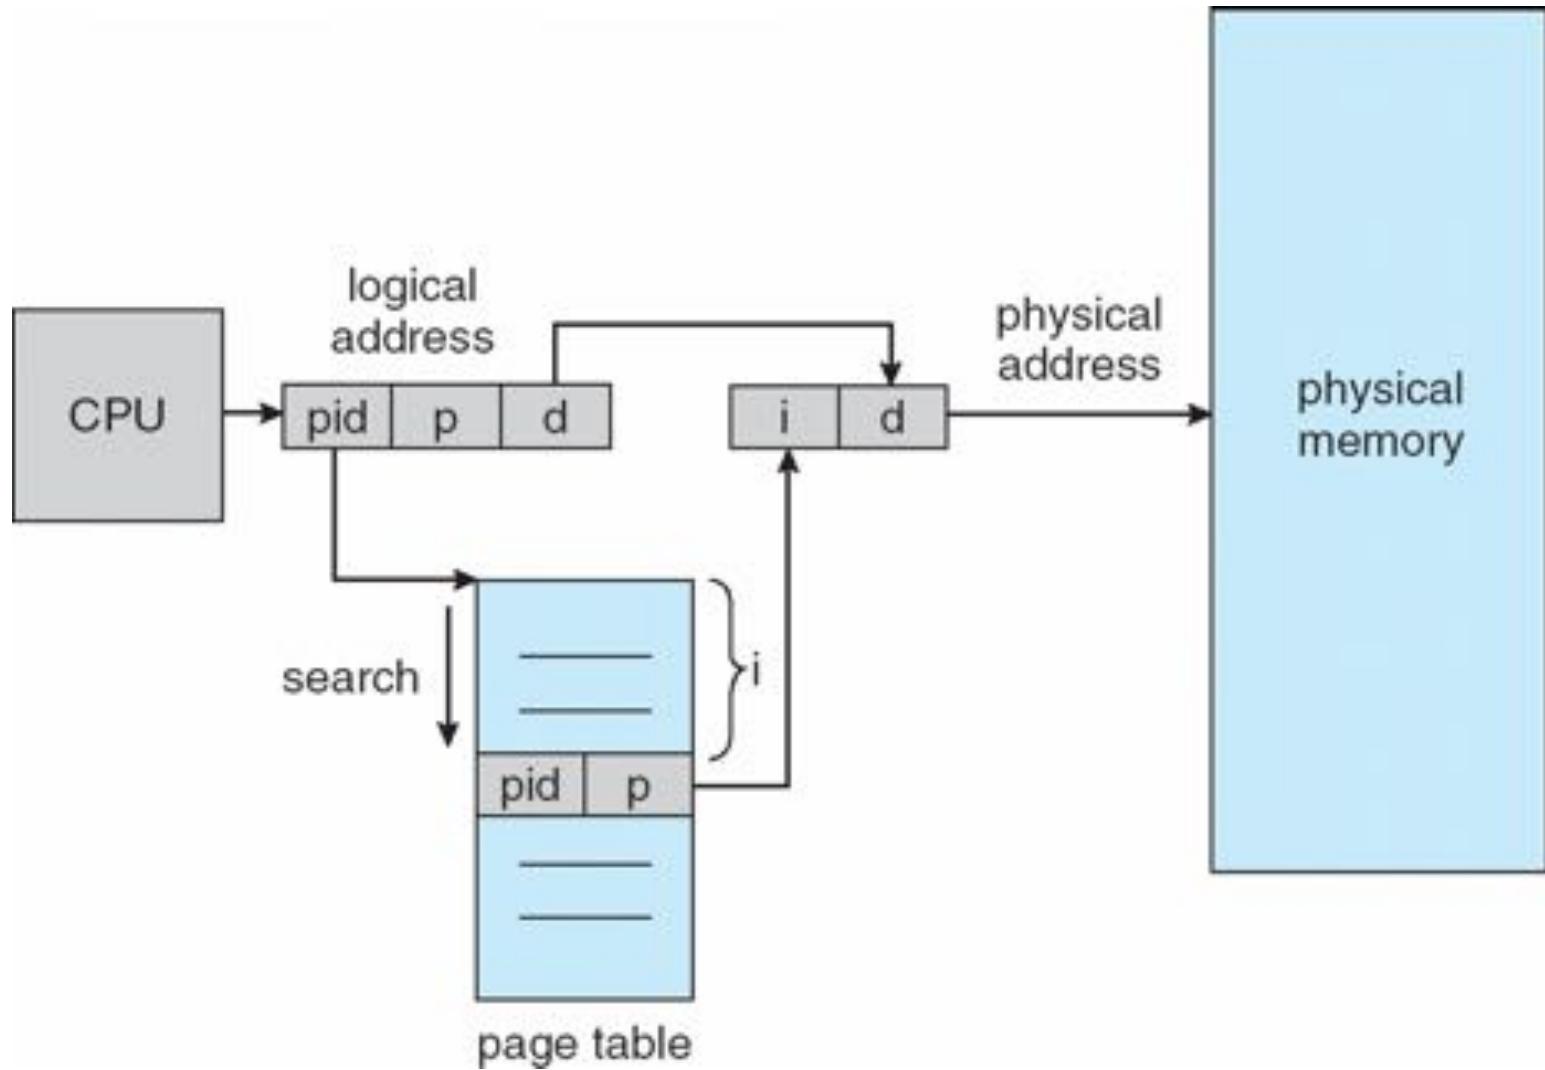
Hashed Page Table



Inverted Page Table

- Rather than each process having a page table and keeping track of all possible logical pages, track all physical pages
- One entry for each real page of memory
- Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page
- Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs
- Use hash table to limit the search to one or at most a few page-table entries
 - TLB can accelerate access
- But how to implement shared memory?
 - One mapping of a virtual address to the shared physical address

Inverted Page Table Architecture



The Intel 32 and 64-bit Architectures

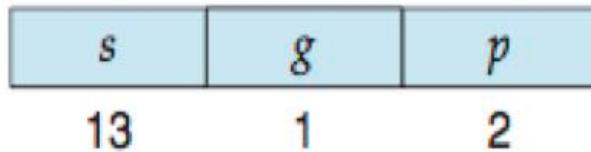
- Dominant industry chips
- Pentium CPUs are 32-bit
 - Called *IA-32* architecture
- Current Intel CPUs are 64-bit
 - Called *IA-64* architecture
- Many variations in the chips, cover the main ideas here

The Intel IA-32 Architecture

- Supports both segmentation and segmentation with paging
 - Each segment can be 4 GB
 - Up to 16 K segments per process
 - Divided into two partitions
 - ◆ first partition of up to 8 K segments are private to process (kept in *local descriptor table (LDT)*)
 - ◆ second partition of up to 8K segments shared among all processes (kept in *global descriptor table (GDT)*)

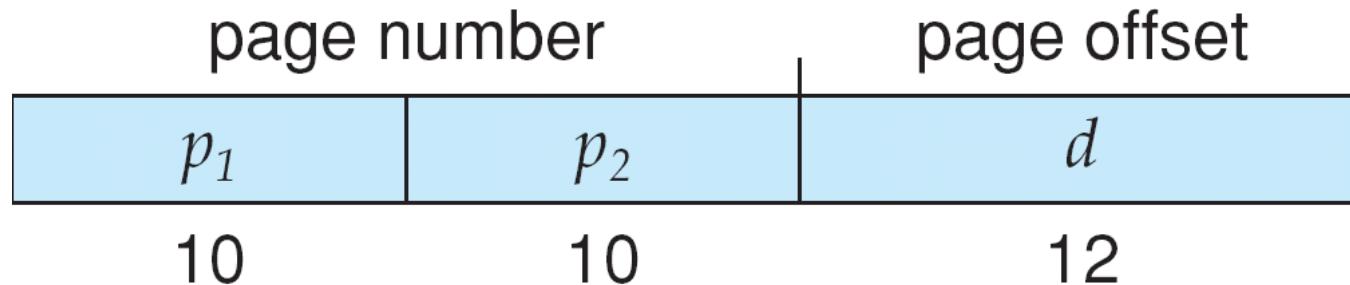
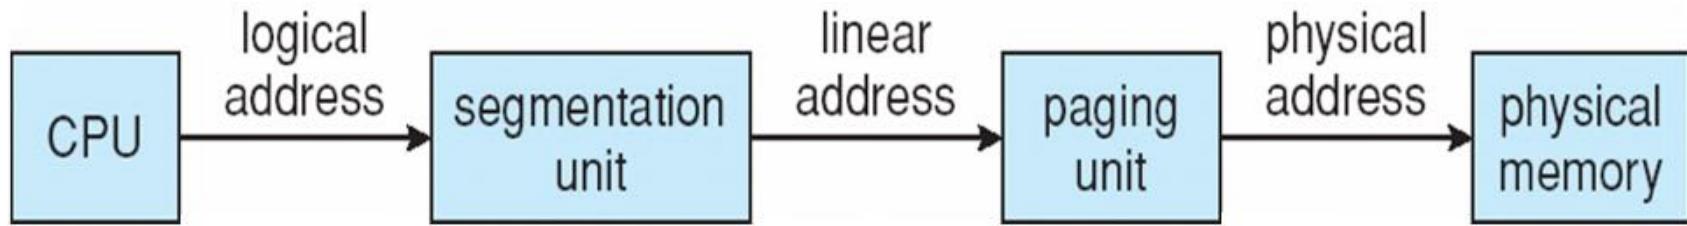
The Intel IA-32 Architecture Addresses

- CPU generates logical address
 - Selector given to segmentation unit
 - ◆ which produces linear addresses

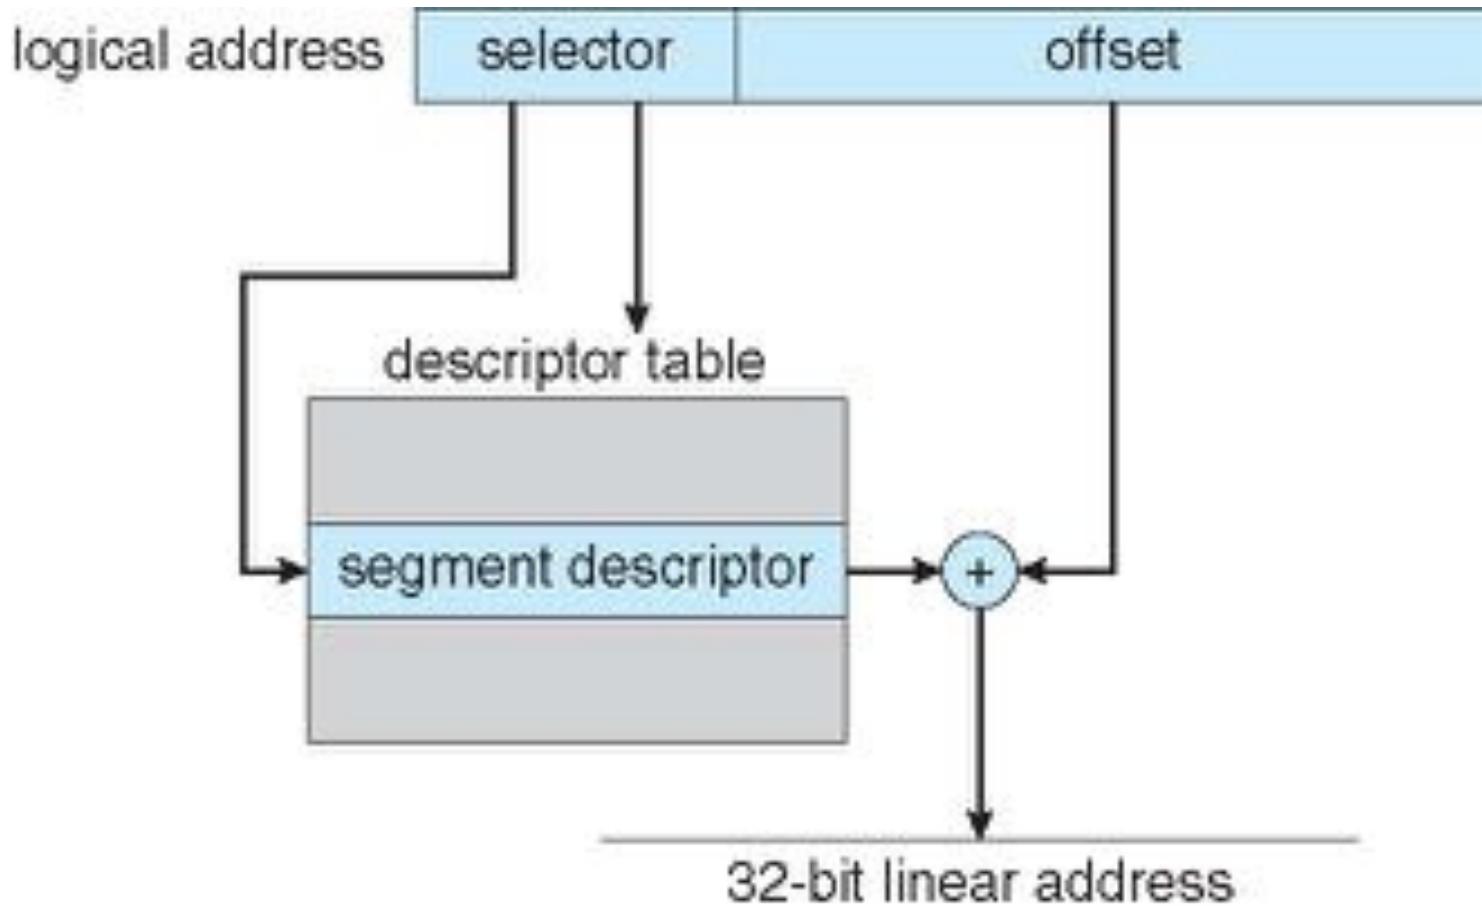


- Linear address given to paging unit
 - ◆ which generates physical address in main memory
 - ◆ paging units form equivalent of MMU
 - ◆ pages sizes can be 4 KB or 4 MB

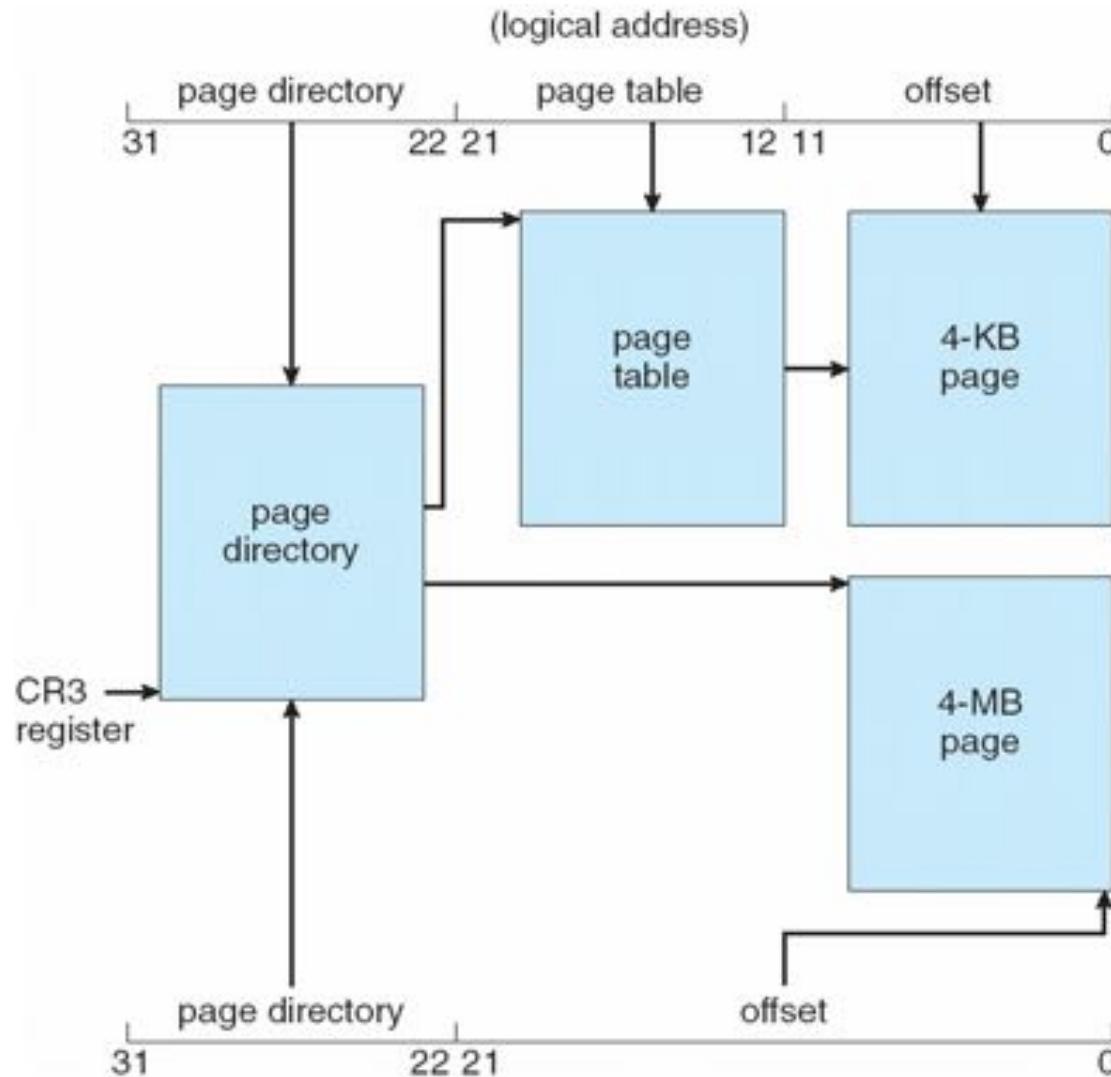
IA-32 Logical to Physical Address Translation



Intel IA-32 Segmentation

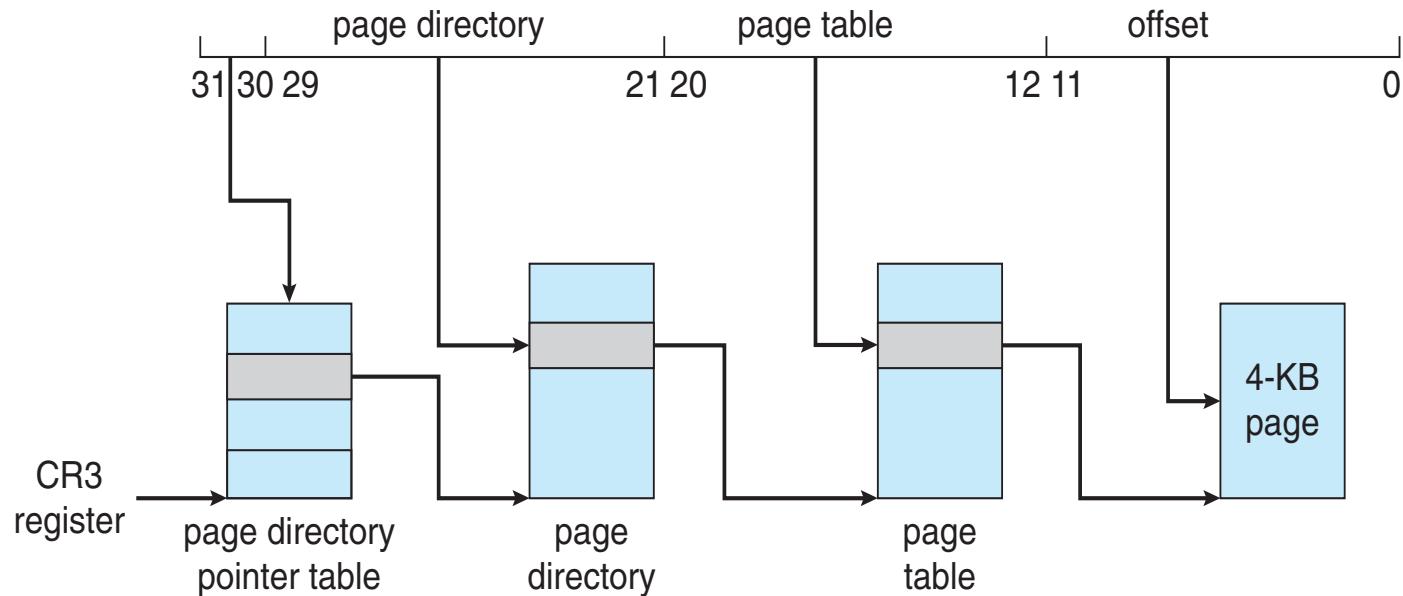


Intel IA-32 Paging Architecture



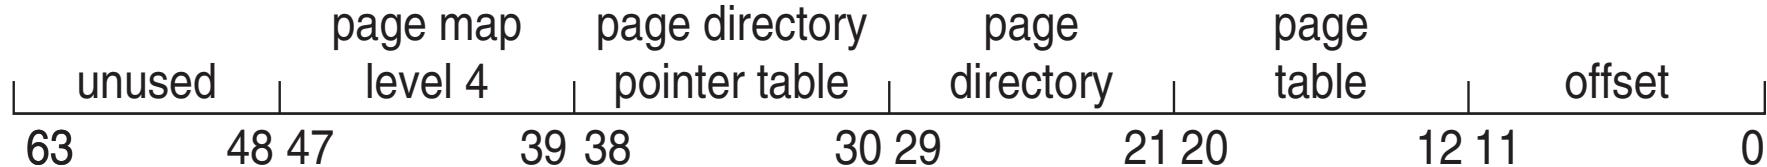
Intel IA-32 Page Address Extensions

- 32-bit address limits led Intel to *create page address extension (PAE)*, allowing 32-bit apps access to more than 4GB of memory space
 - Paging went to a 3-level scheme
 - Top two bits refer to a *page directory pointer table*
 - Page-directory and page-table entries moved to 64-bits in size
 - Increases address space to 36 bits (64GB physical memory)



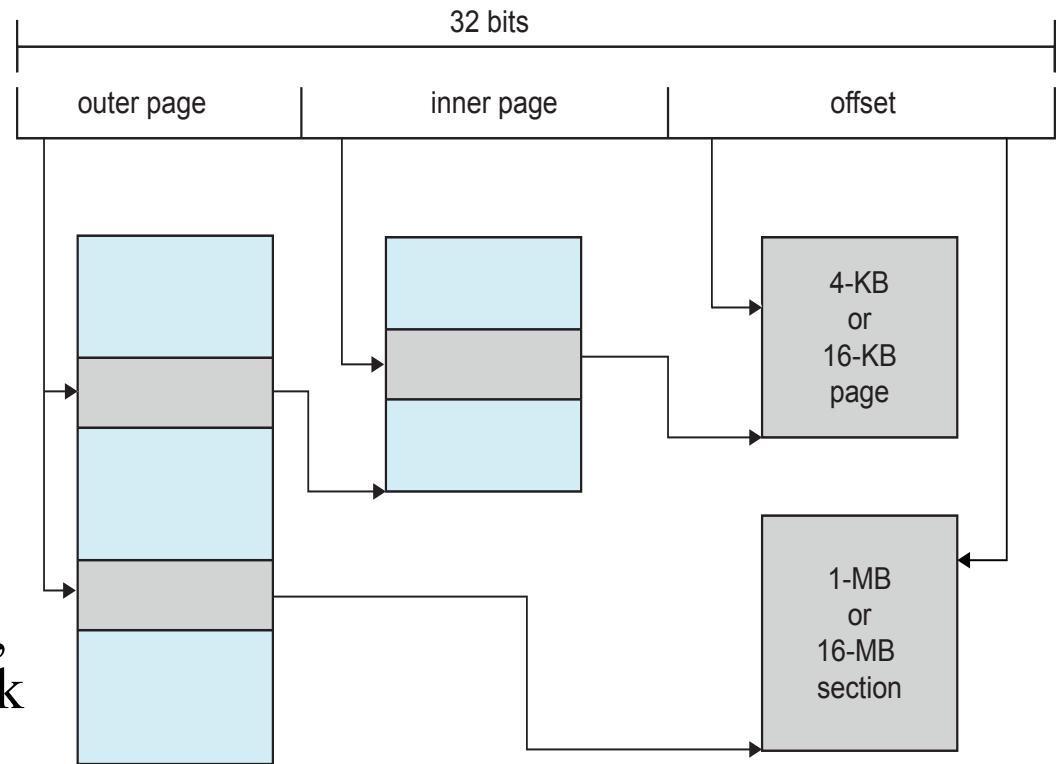
Intel x86-64

- Current generation Intel x86 architecture
- 64 bits is ginormous (>16 exabytes)
- In practice only implement 48 bit addressing
 - Page sizes of 4 KB, 2 MB, 1 GB
 - Four levels of paging hierarchy
- Can also use PAE so virtual addresses are 48 bits and physical addresses are 52 bits



Example: ARM Architecture

- Dominant mobile platform chip
- Modern, energy efficient, 32-bit CPU
- 4 KB and 16 KB pages
- 1 MB and 16 MB pages (termed sections)
- One-level paging for sections, two-level for smaller pages
- Two levels of TLBs
 - Outer level has two micro TLBs (one data, one instruction)
 - Inner is single main TLB
 - First inner is checked, on miss outers are checked, and on miss page table walk performed by CPU



Next Class

- Virtual memory